

content-boundary-class-a-judgement

Content boundary — class A INWARD rows, judged (2026-09-08)

Status: JUDGED. NOTHING WAS REDACTED, NOTHING WAS ALLOW-LISTED, NO THRESHOLD OR BASELINE WAS MOVED, AND `scripts/verify-content-boundary.sh` WAS NOT EDITED. The gate still exits 1, and it should.

This closes the option-4 half of the decision packet carried in the four root governance carriers under “*Content-boundary re-baseline — the decision packet*”, and answers **R11** of [../specs/005-clone-and-clear-red/spec.md](#). It does **not** make the detector green — spec 005 puts that explicitly out of scope — and it must not be read as clearing the population the gate reports.

Machine-readable artefact: [findings/content-boundary-class-a-judgement.jsonl](#) — one object per judged row. **Paths, classes, digests, epochs and counts only. No matched text appears in it, in this document, or in any report derived from either.**

What was judged, and what “class A” means

Class A is the content-boundary population whose PUBLIC destination is one of this umbrella’s two spec trees, `specs/001-*` and `specs/002-*`. It is the largest structural class in the packet, and the packet’s recommended action was to judge its **INWARD** rows — the rows whose text was committed on the PRIVATE side first.

Two full gate runs were taken, tied together with --expect-corpus, and BOTH are reported — not the more convenient one. The population moved sharply between them; §“Why the total fell” below explains why, measured rather than guessed.

	run 1 · 2026-09-08 12:14	run 2 · 2026-09-08 12:53
gate exit	1	1
gate total	16,155 (prose 15,107, short 831, name 217)	9,945 (prose 9,456, short 373, name 116)
class A	9,513 (prose 9,168, short 295, name 50)	7,086 (prose 6,902, short 139, name 45)
class A OUTWARD — public text is older	7,631 (80.2 %)	5,815 (82.1 %)
class A INWARD — private text is older	1,832 (19.3 %)	1,226 (17.3 %)
class A direction UNDETERMINED	50 (0.5 %)	45 (0.6 %)
corpus during the run	MOVED , 66 of 14,891 files	MOVED , 4 of 14,837 files
corpus vs run 1's manifest	(wrote it)	DIFFERS

Judged here: the UNION of both runs' INWARD and direction-UNDETERMINED rows — 1,933. 1,882 came from run 1, and 51 further rows appeared only in run 2 (all G2, all into specs/001-.../contracts/*); judging the union rather than either run alone is what stops a moving tree from quietly dropping a row out of the reading queue.

The direction split has now been measured **five times over five populations whose totals differ by more than 3,000 rows**, and OUTWARD has landed between **79.1 % and 82.1 %** every single time (79.7 / 79.3 / 79.1 / 80.2 / 82.1). **The totals move; the direction does not.** That is the finding to act on.

Why the total fell — measured, not guessed

An unexplained fall is worse than a stable red, so this one is attributed. The gate total fell 16,155 → 9,945 in 39 minutes. **Nothing was silenced to achieve it**, and each of the four candidate mechanisms was tested rather than assumed:

Candidate mechanism	Measured
the instrument was weakened	No. scripts/verify-content-boundary.sh is unmodified; the path-reference mask is identical in both runs (1,425 long + 5,613 short shingles)
an exemption was added	No. .content-boundary-allow is unmodified; pardons rose 4,077 → 4,083, moving with the corpus
private files were deleted or untracked	No. Of the 139 private files behind the lost rows, 0 are gone from disk and 0 have left git ls-files
a new public repository exonerated them	No. The derived fleet is identical in both runs (15 rows, same roles); a 40-row sample of the lost rows found 0 whose text is present in the newest public submodule
the private corpus itself shrank	YES. Private prose tokens fell 3,597,903 → 3,129,293 (-13.0 %) , name candidates 527 → 353, already-public keys 2,306 → 869

The mechanism is concurrent work, not remediation. The private workshop submodule had **217 modified files** in its working tree while both runs were taken, and the gate reads the working tree. A 13 % smaller private corpus yields a smaller private key space and therefore fewer matches, in every class at once — which is exactly the shape observed. **This is a measurement of a moving tree, and neither total may be quoted as the count of any single tree.** Both runs said so themselves.

Method — stated so it can be attacked

Direction is measured **at the level of the text, not the file**, and **at corpus level, not pairwise**:

1. Every class-A row's matched string is normalised exactly as the gate normalises (`[^A-Za-z0-9]+` → space, lowercased — the same rule at `scripts/verify-content-boundary.sh:1938`).
2. Every commit of every repository is walked **chronologically**, and every blob it adds or modifies is normalised the same way and scanned for every still-undated string. A string is dated at the **earliest commit whose blob content actually contains it**, anywhere in that repository — so a string that moved between files is still caught, and a file merely renamed does not restart its clock.
3. Public = the umbrella. Private = workshop, ai_interviewing, monetization. Blobs scanned: **6,449 public, 2,881 private** (workshop 2,072, ai_interviewing 794, monetization 15). **Nothing sampled, nothing extrapolated.**
4. A string committed on one side only is **UNDETERMINED**, never a verdict.

The probe's normalisation was validated against the gate rather than assumed to agree with it. On a seeded random sample of **60** class-A rows, the row's matched string was located, under this probe's own normalisation, in the current content of both the named public file and the named private file: **60 of 60 on each side, 0 misses**. Two corpora normalised differently is the one failure the gate's own header says must never happen, and this is the check that it did not.

The name class cannot be dated at all: the gate **withholds** the matched text by design, so there is nothing to normalise. All 50 class-A name rows are therefore direction-UNDETERMINED **by construction**, and defeating that protection to date them would be the disclosure the protection exists to prevent.

What “private committed first” does NOT mean

It establishes an **order**, nothing more. A symbol name, a task line, an error string, or a design decision made in code and written up afterwards all produce INWARD rows with nothing having crossed the boundary. **1,832 rows are not 1,832 leaks and must never be reported as such.** What the number does establish is that the

reassuring story — “*the specs are outward propagation*” — is true for four rows in five and **false for the fifth**, which is why the fifth was read.

The 68 rows against the recorded teaching material — read first

These are the only judged rows whose private side sits under workshop/chapters/, the directory holding a recording of a private teaching session with an identifiable third party. They carry the highest cost if wrong, so they were judged first and in the most detail.

Rows	68, all prose
Private source	one file: workshop/chapters/01/transcript/accuracy-plan.json
Public destinations	specs/001-.../tasks.md (45, at 4 distinct lines) and specs/002-.../tasks.md (23, at 3 distinct lines)
Distinct matched strings	30
Direction	INWARD, lead 141 seconds — private 2026-09-03T20:15:14, public 2026-09-03T20:17:35

Verdict: NOT A DISCLOSURE. Three independent measurements, each re-derivable:

1. **The private file contains no recorded material of any kind.** It is a stratified **sampling plan** for measuring transcript accuracy. Its fields are a chapter id, a filename, a duration, a seed, thirty {t0,t1} window bounds, and five descriptive strings. It carries **no transcript text, no speaker turn, no utterance, no participant identifier and no personal data** — there is no field of that kind in it to carry any.
2. **Every one of the 30 matched strings resolves to a methodology field**, and the split was measured rather than asserted: estimator_unit **23**, the plan’s own how-to-use instruction field **6**, and reference_schema.windows[].onset_s **1**. **Zero** resolve to any window, timing, content or identity field.
3. **The public side is the instruction to use that plan.** The four specs/001-.../tasks.md lines and the three specs/002-.../tasks.md lines are task prose telling a human what to produce, in what form,

and why the estimator unit matters. Naming the plan's path is a **reference**, which this boundary explicitly permits.

The 141-second lead is the whole mechanism: one authoring session in which the plan was committed privately and the public task text describing it was written two and a half minutes later. **The direction is real and the crossing is not.**

Separately, the 5 class-A name-class rows whose private side is the private notes PDF are covered under G8 below. Its **filename carries the third party's given name**, so that path is **withheld** from the machine-readable artefact — a path is normally quotable, this one is not.

Every judged group, with its verdict

Group	Rows	Digests	Private files	Public files	Median lead	Verdict · reason
G1 chapter sampling plan	68	30	1	2	141 s	NOT A DISCLOSURE · methodology field of a sampling plan
G2 platform source (platform/, pipeline/, scripts/, curriculum/)	1,222	885+	102	15	29 min	NOT A DISCLOSURE · spec and implementation co-authored
G3 session-evidence briefs	184	121	16	5	141 s	NOT A DISCLOSURE · build-session brief about the specified platform
G4 platform documentation (limits, manual, faq, README, runbooks, prompts)	163	55	20	6	11 h	NOT A DISCLOSURE · documentation of the same platform
	126	71	14	5	194 s	

Group	Rows	Digests	Private files	Public files	Median lead	Verdict · reason
G5 training areas						NOT A DISCLOSURE · shared engineering argument about the public contract
G6 chapter-01 redaction review	6	3	2	2	21 h	NOT A DISCLOSURE · by-design counts-only public summary
G7 evidence files and gate registries	114	102	18	9	7 min	NOT A DISCLOSURE · gate and evidence vocabulary
G8 name class (direction undetermined)	50	1	10	1	—	NOT A DISCLOSURE · digest previously adjudicated as a product brand
TOTAL	1,933					0 DISCLOSURE · 0 unjudged

Row counts are over the **union** of both runs; the per-group digest, file and lead-time columns are run 1's, which is the larger population.

Every row landed in a group; none is ungrouped. The artefact carries a G9-UNGROUPED verdict of UNDETERMINED precisely so that a future run over a changed corpus cannot silently absorb a row that no group describes.

G2 — 1,222 rows, and why they are the least surprising half

Private **source code** (Go 479, Python 377, shell 323, TypeScript 13 by row) against the public **contract and task documents that specify it**. The private side is the implementation of the platform specs/001-** and specs/002-** define; the public side is its specification. The overlapping text is comment, docstring and status prose about that platform's own behaviour — the public document's own subject matter. Median lead 29 minutes; **74.7 %** of all INWARD rows are inside one hour and **90.1 %** inside a day, which is the signature of one author writing spec and code in one sitting, in both orders.

G5 — the one group where a second opinion is reasonable

126 rows, dominated by docs/training/areas/05-evidence-gates-and-anti-bluff.md (64) reaching specs/001-.../contracts/pipeline-cli.md and .../http-api.md. Both sides argue the same point about the platform's own status.sh exit-code contract — a three-numbered refusal of the “benign state” reading — and the private side is a training Q&A restating it.

Judged NOT A DISCLOSURE: the subject is the public contract's own subject, the argument concerns an exit code the public document normatively defines, and **no pedagogy, no participant, no session content and no personal data appears on either side of any matched string**. The remaining 62 rows of the group are short-class generic phrases and multi-format duplicates (.md / .html / .pdf renders of one document each counted separately) across 13 digests.

Recorded honestly: this is authored curriculum material, it is private-first, and the wording is reused. If the operator regards the *wording* of a training area as a commercial asset distinct from the contract it describes, that is a licensing judgement, not a boundary one — and it is the operator's to make. It does not change the boundary verdict.

G6 — the pair that looks worst and is designed

workshop/docs/redaction-review-chapter-01.{md,pdf} reaching specs/001-.../redaction-review-summary.md and .../tasks.md, 6 rows over 3 digests. The public file **states its own boundary in its first paragraph:** counts and verdict only, no chapter content, no personal name, no pid,

no transcript line reference, with the detailed review named by private path. The overlap is the vocabulary of the category headings and the decision counts. That is the design working, not failing.

G8 — 50 rows, one digest, and a limit that is the protection

All 50 class-A name-class rows carry the **single** digest 04bc71a2, all in specs/001-.../research/transcription.md, from 10 private files across workshop and ai_interviewing. That digest was resolved in content-boundary-incident-2026-09-01.md §12.2 to a **machine-learning model-hub brand, two tokens — not a person**, by re-implementing the gate's own h() and rebuilding the reported line's capitalised runs.

Honest boundary (§11.4.6): that resolution was NOT independently re-derived in this session. The identity evidence carried here is narrower and is stated as such: same gate, same digest function, same 8-hex digest, same public file, same class. The verdict rests on the earlier wave's adjudication, and a reader who wants it re-derived should re-run §12.2's method rather than trust this sentence.

Direction for these 50 rows remains UNDETERMINED and cannot be resolved with this instrument, because the gate redacts the matched text by design. That is the same mechanism which keeps a real person's name out of an archived artefact: **the protection and the blind spot are one thing**, and trading the first for the second would be the wrong trade.

Two checks that do not rely on the gate's own exonerations

Both bypass the detector entirely. Both were re-run in this session.

1. **The third party's given name is measured absent.** The candidate tokens are **derived programmatically from the private artefact's own path** — never typed, never written down — and searched as whole words across every tracked file of all ten in-scope public repositories, and separately across every untracked file of this public working tree. Of 5 derived tokens, 4 are ordinary words with 182–1,269 public files each; **exactly 1 has zero occurrences anywhere public**, tracked or untracked. §11.3's redaction holds.
2. **No judged row's private path was published if the path itself discloses.** The one private path whose basename carries that token

is masked in the machine-readable artefact by the same derived rule, so the artefact cannot publish what the redaction removed.

What this did NOT do

- **No public file was edited to remove a match.** No redaction was performed.
- **.content-boundary-allow was not touched.** No pair was added, and none is proposed. An allow-list entry buys a green exit at the cost of hiding the row from the next reader, and adding one here would assert boundary flows that do not exist.
- **scripts/verify-content-boundary.sh was not edited.** No window, floor, rank, bucket or filter was moved. The counts this document reports are the unmodified gate's.
- **No baseline was moved and no count was re-based.** The gate exits 1.
- **Nothing about the incident's \$3, \$8B or \$11.4 changed.** History still carries the wave-1 and wave-2 material; the provider purge is still unsent; the conversation with the third party is still the substantive remedy and is still the operator's to have.

Honest boundary — read this before quoting any number above

1. **The corpus MOVED during both measuring runs, and the gate said so.** Run 1 reported corpus MOVED – 66 of 14,891 enumerated file(s) changed; run 2 reported 4 of 14,837 **and** corpus DIFFERS from --expect-corpus. So **neither 16,155 nor 9,945 may be quoted as the count of any single tree**, and this document does not quote either as one. The movement is concurrent work by other agents in this one checkout — a 13 % contraction of the private corpus, attributed above. **This is the reason the union of both runs is judged rather than either alone.**
2. **What the instability does not touch is the direction split**, which is the figure this document acts on. It has now been measured on five populations whose totals differ by more than 3,000 rows, and OUTWARD has landed between **79.1 % and 82.1 %** every time.
3. **Judgement is by GROUP, evidenced by inspection of the largest and the riskiest members — not by reading all 1,138 distinct digests one at a time.** G1 was resolved string-by-string (30 of 30). G5, G6 and the 182 rows whose lead exceeds 24 hours — the rows where the same-session mechanism is weakest — were opened

individually. The remainder was judged on its group's measured shape. **The 51 rows seen only in run 2 were direction-measured but were not separately opened**; all 51 are G2, into specs/001-.../contracts/*, and they inherit that group's verdict rather than carrying an independent one. A reader who disagrees with a group can re-derive it from the artefact, which names every row's private path, public path and line.

4. **This judges class A only.** The other ~6,600 rows of the population — public reusables extracted from the private tree, the governance cascade, and a genuinely unassessed remainder — are **not** judged here and are **not** cleared by it.
5. **This says nothing about git history.** Every pass reads the working tree.

What the operator is asked to decide next

Nothing here requires a decision to stand; the judgement is recorded and the gate is unchanged. Three things are worth deciding, in this order:

1. **Nothing, for class A.** The recommended action of the decision packet has now been carried out and found **0 disclosures over 1,882 rows**. Option 2 of that packet — allow-listing classes B and C — is still **not recommended**: it would hide rows from the next reader, and its direction evidence remains file-level and sampled, which is weaker than what class A now has twice over.
2. **Option 3 — teach the gate direction.** This is now the only remaining option that raises precision without hiding a row, and its feasibility is no longer in question: the probe behind this document dated 16,599 rows across two runs, against 9,330 blobs, in about four minutes per run. Subtracting **OUTWARD** rows with the recall cost printed, exactly as the detector already prints its already-public subtractions, would remove ~80 % of class A from the reading queue **without asserting a single boundary flow that does not exist**.
3. **The G5 wording question**, if the operator regards a training area's wording as a commercial asset. That is a licensing decision and is not a boundary one.

Re-derive rather than trusting this document

```
bash scripts/verify-content-boundary.sh --json --expect-corpus /  
    tmp/cb.manifest  
bash scripts/verify-content-boundary-judgement.sh  
    # 0 / 1 / 2  
bash scripts/verify-content-boundary-judgement.sh --prove-failure
```