

OPERATOR-DECISIONS-2026-09-08

Operator decisions — 2026-09-08

Recorded per §11.4.208 so a decision is not re-litigated and not silently reinterpreted. Each row states the question, the decision, and what it obliges.

CONTENT BOUNDARY. This is a PUBLIC repository. Every row below refers to private workshop material BY IDENTIFIER OR PATH ONLY. No transcript text, no note text, no area title, no question text and no person's name appears here.

#	Question	Decision	Status
1	Semantic-search leg needs an embedder; local models now off	Keep semantic OFF. Measured identical quality (top-1 18/22, top-5 19/22) at 9.4x the speed. Reversible with one env var	DECID
2	meeting-notes route deliberately NOT_BUILT on one undecided row	Brainstorm fully, then implement — including Meeting Notes	DECID
3	3 documents unpublished on W2 violations	Cite them, then publish	DECID
4	12 served areas carry no assessment	Fill / generate what is missing. “Stick to common sense”	DECID
5	The one flagged span (id 01M1H9DKFTVEJMW5YFZ1GC7PYN)	Serve it. Operator inspected the span and judged it carries nothing private. Add a reason code for “flagged withheld-only by corpus statistics, judged not private” and record the decision against the row	DECID
6	The word-run rule that flagged it		DECID

#	Question	Decision	Status
		Keep the rule as-is. It caught 4 genuine rows because it is wide. Future flags route to the operator as judgements	
7	meeting-notes contract shape	Match the three siblings exactly — same cumulative count VS own_count split, same status set	DECID
8	66 plan-coverage rows, 0 decided	Show all in one pass	DONE (sheet produ
9	Sequencing against the QA session	One batch, then QA. Land everything, restart once, run the second adversarial review, then hand over	DECID
10	20-passage evidence floor sorts on the wrong quantity (r = -0.066 against actual authorability)	Re-baseline the floor onto anchored usable citable passages at 8, and author the banks. REVISES decision #3 of 2026-09-07 on measured grounds	DECID
11	Filling every area empties the M4 proof arm's population	Fill everything, give M4 a synthetic no-bank fixture so the arm keeps detecting	DECID
12	8 rows read as likely covered	All 7 distinct points recorded COVERED (one appears in two chapters). 5 carry evidence locators; 1 rests on operator judgement with no locator and says so in its own reason field	DONE
13	11 plan rows are ASR fragments naming no deliverable	Fix the extractor, retire the 11 as extraction defects — NOT as dropped, which would put them in the commitment ledger	DECID
14	MIN_SEGMENTS_FOR_INFORMATIVE_ABSENCE = 500 makes not_covered unreachable for short chapters	Scale the guard to chapter length , with its own validation and a paired mutation	DECID
15	Document 04's last 8 blocks cite evidence that has no passage id	Ingest pipeline/ {README,CALIBRATION,AUDIT}.md.	DECID

#	Question	Decision	Status
16	264 pre-existing citations of unverified relevance	Operator rejected deleting the 8 measured claims and rejected the editorial exemption Audit them all , and the same pattern in the already-published areas 01 and 03. Depublication is an acceptable outcome	DECID
17	publication-reviews.jsonl / area-materials.jsonl gitignored — the files gating what is served are not in version control	Track them	DECID

18 | Every graded question answerable by shape alone: 138/138 keys are the longest option (100%), 122/138 at position 1 (88%), pass mark 70%, no shuffle anywhere in the serving path | **Fix properly: deterministic shuffle AND distractor rebalancing**, plus a gate that fails if either heuristic beats chance, with a paired mutation and a control | DECIDED |

19 | The producer/verifier split is process-scoped only; two invocations 2s apart publish on same-minted citations, and the prover's key arm is inverted so it passes with the protection removed | **Make it durable, fix the prover**, register it in the check registry, and close the promote_curriculum_areas.py --apply path that wrote 37 of 41 live review rows with no citation checking | DECIDED |

20 | QA session timing given the above | **Fix the blockers first, then QA**. The shape leak and the split are fixable in hours; everything else is measured and sound | DECIDED |

21 | 32 answer keys appear verbatim in the lesson section the question links to (31 of 290) — all against area-document prose, never a transcript | **Separate study mode from assessment mode**. Keep the keys and the study links; suppress lesson_sections in assessment/interview rendering. Rewording 31 correct keys was rejected as degrading good prose to defend a property only one mode needs | DECIDED |

22 | Two questions carry answers the evidence contradicts (an asserted absence the corpus refutes; "appears in neither transcript" where a redacted passage names it) | **Rewrite both to what the evidence supports**, neither withdrawn. The second preserves the ABSENT vs WITHHELD distinction the platform is built on | DONE |

23 | 52 flagged rows in the 30 older committed banks, unread | **Read and repair all 52**. Result: 51 repaired, 1 confirmed already supported

— 98% PRECISION, which says nothing about recall | DONE |

24 | Areas 26 (5 anchored) and 23 (6) serve tests while 19 (5), 27 (6) and 28 (6) are refused — authorship date, not evidence, is the selector |

Author banks for 19, 27 and 28 to match. Level up rather than withdraw two real banks. An agent may still decline any of the three if the evidence genuinely cannot support one | DECIDED |

25 | Four kg_* kinds (226 rows of TODOs, meeting notes, open questions, next-meeting points) exist in the index but were never in the advertised vocabulary | **Make them first-class searchable.** This is the session-record content built today — reachable through chapter routes but not through search | DECIDED |

26 | ROOT CAUSE: a bank obeying its own “cite nothing outside the evidence pack” constraint could not cite its own evidence |

SUPERSEDED — see 26a. Originally “re-cut the packs on sentence boundaries”; the premise was measured false | WITHDRAWN |

26a | Re-measured: there is NO pack producer and NO artefact — the pack is derived by reading a Markdown file, and NO gate ever enforced membership (B1–B8 check resolution, redaction and lesson-section path, never pack membership). Only 30.6% of outside-pack citations sit at distance 1; the rest reach up to 644 segments away. Best closure admits 41.3% at 3.75× pack growth. Corpus re-cut separately proven unsafe: one merged pair moves root_hash 62e62eea... → 960db1d2..., orphans an anchor, and models are off so nothing can be re-embedded |

Withdraw the constraint entirely. Replacement rule: cite the passage that carries the fact; admissible iff it RESOLVES UNREDACTED and AN AUTHOR HAS READ IT against the prompt. The constraint is wrong in KIND, not granularity | DECIDED |

26b | The withdrawn constraint’s prose still stands verbatim in 6 banks, instructing future authors to obey the rule that caused the defect | **The agent already reading those files corrects the 6 \$comment blocks in-pass** — recording what the rule said, that it is withdrawn, and why, rather than deleting it silently | DECIDED |

27 | 182 older-bank questions the screen never flagged, unread. Precision is 98%; recall is known poor (reading found 73% where the screen flagged 37.5%) | **Read all 182.** The measured defect rate in the UNFLAGGED population is the first honest estimate of the screen’s recall | DECIDED |

The defect behind the defect — recorded because the pattern repeated six times

Every stream today found that the reported symptom was real and the stated mechanism was wrong. Fixing the stated mechanism would have produced a green gate over an unfixed product in at least three cases.

- Shuffling choices would NOT have fixed the shape leak: the choice **id** was itself the key (-c1 on 88.4%), so the key would have moved with its label on.
- The graded path was not BYPASSING citation withholding — the type it uses has no citations field, so the protection is **inapplicable**, and wiring it would have shipped an inoperative gate reading green forever.
- proposer was empty because a downstream check inferred independence from it; the field was lying to protect a **proxy**, and fixing the inference let the field tell the truth.
- The anchor defect was **fleet-wide**, not introduced this session — 34.2% in the older banks against 37.5% in the new ones.
- The evidence-pack constraint **guaranteed** the defect it existed to prevent.
- CK021 counted “no bank authored” and was read as “below the floor” — a correct number over the wrong set.

Claims WITHDRAWN on 2026-09-08 after an adversarial review

Recorded because a withdrawn claim is worth more than a quiet correction.

- **“01 and 03 serve on hand-verified citations.”** OVERSTATED. The citations were hand-read, but the review artifact records only a MECHANICAL check and says so in its own stamped text: it establishes that a citation RESOLVES, not that the cited passage SUPPORTS the claim. reviewer is the pipeline, not a person.
- **“dropped 0, vanished 0.”** dropped=0 is a real measurement. vanished=0 is arithmetically FORCED — it is a residual, every branch increments exactly one counter, and the gate asserting it is therefore a false-green control (§11.4.201). The two must never be reported together as if both were measured.

- **“5 areas below the evidence floor.”** WRONG POPULATION. CK021 fires on “no bank was authored”, not “below the floor”. Seven areas are below 8; two areas serve 9-question banks on 6 and 5 anchored passages because they predate the re-baseline and were grandfathered without a decision.
- **“261 questions across 34 banks”**, reported as an achievement without ever checking whether the questions test anything. See decision 18.
- **“The producer has been split from the verifier.”** REFUTED. See decision 19.

Still not put to the operator

- Document 05’s four umbrella-facts have no citation path from inside this private repo. Options: rewrite to what this corpus supports (losing true content), ingest public umbrella governance output into the private corpus, or leave 05 unpublished. None taken.
- Document 04’s nine calibration blocks stay uncitable while the ingest is stopped. The ingest is blocked by a DETERMINED NEGATIVE: with models off there is no embedding endpoint, so new passages would enter the index with no vectors and be silently absent from semantic retrieval.
- The served redactions table holds 0 rows while the log records 162 effective suppressions. No text is exposed; the artifact simply cannot state WHY a row is withheld. Materialising it is an index rebuild.
- -question-verify none is not a recognised value (focus|llm), so L5 is unwired. The server reports it as unwired rather than silently selecting one.

Decisions that overruled a recommendation

Recorded because a recommendation overruled is worth more than one accepted.

- **#12, the “three-hours web build” point.** The coordinating session advised AGAINST recording it, because the session demonstrates a web surface but no instrument established it is the SAME artefact the planned point refers to. **The operator recorded it anyway.** Obligation, and it is discharged: the disposition row carries NO evidence locator and its reason field states in full that it rests on

operator judgement alone and must not be cited as evidence-backed. A covered decision is the only irreversible direction here — an undecided point merely carries forward — so this row is the one to revisit first if the plan later looks short.

- **#11 vs #10.** These two answers pull against each other: the re-baseline at 8 clears 8 of the 12 areas, while “fill everything” reaches all 12 including the 4 a measurement calls genuinely thin. Resolved as: fill every area the evidence honestly supports, fabricate nothing, and REPORT any area that cannot be authored truthfully. An honest empty bank beats a padded one, and the M4 fixture removes the proof-population objection either way.

Not yet put to the operator

- The served database’s redactions table holds **0** rows while the log records **162** effective suppressions. No text is exposed — measured by S1 — but the served artifact cannot state WHY any row is withheld. Materialising it is an index rebuild.
- `-question-verify none` is not a recognised value (`focus|llm`), so **L5 is unwired**. The server reports it as unwired rather than silently selecting one, which is correct; but “none” arguably ought to be a legal value meaning deliberately-off.
- The 8 weakest evidence rows still await confirmation.