

QWEN

vasic — QWEN.md

INHERITED FROM submodules/constitution/ QWEN.md

All rules in submodules/constitution/QWEN.md (and the submodules/constitution/Constitution.md it references) apply unconditionally to this project. Project-specific rules below extend them — they do NOT weaken or override any universal clause.

When this file disagrees with the constitution submodule, the constitution wins.

Base agent rules: submodules/constitution/QWEN.md — READ IT FIRST. The base file is authoritative for any topic not covered here. Locate it from any nested depth with submodules/constitution/find_constitution.sh.

This carrier is read by the Qwen Code CLI agent (qwen-code). It is one of the four repository-root governance context carriers (CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md) that anchor 11.4.157 requires to be maintained in lockstep. Everything below this paragraph is byte-identical across all four; only this opening section differs.

Where the canonical rules actually live

Nothing is copied here. This file is a pointer; the authority is the submodule.

What	Canonical path in this repository
The universal constitution (11,787 lines, 255 ### § anchors,	

What	Canonical path in this repository
1,801,862 bytes, re-measured on disk 2026-09-03 BEFORE AND AFTER the fourth fast-forward, at 3be10826f3d2 and again at 2887b42e9349 — every figure reproduced identically, sha256 973e473c823b4a13... on both sides; the SAME Constitution.md blob ad32fc94... as all four superseded pins 902979027a90, f16ea779b82a, f5876a3b700e and 3be10826f3d2, so these figures have not moved across any of the four bumps)	submodules/constitution/ Constitution.md
Claude Code carrier	submodules/constitution/ CLAUDE.md
Codex / Cursor / Aider / OpenCode / Crush / Kimi CLI carrier	submodules/constitution/ AGENTS.md
Qwen Code carrier	submodules/constitution/QWEN.md
Gemini CLI carrier	submodules/constitution/ GEMINI.md
Full anchor corpus companion	submodules/constitution/ CLAUDE_ANCHORS_FULL.md
Parent-walk resolver (works from any nested depth)	submodules/constitution/ find_constitution.sh
Post-pull governance hook	submodules/constitution/ scripts/post_update_hook.sh
Forbidden-command PreToolUse guard	submodules/constitution/ scripts/hooks/guard-forbidden-commands.sh
Propagation / covenant gates	submodules/constitution/ scripts/gates/

The pin has now been FAST-FORWARDED FOUR TIMES — 2026-09-01, 2026-09-02, and again on 2026-09-03, each on explicit operator authorization. Every bump closed one INSTANCE of the gap; the HOLE that let the first one open unnoticed is now WATCHED, and the watcher has gone GREEN, RED and GREEN again inside three days. submodules/constitution is currently checked

out at 2887b42e9349, on branch main, and that **EQUALS** `git ls-remote git@github.com:HelixDevelopment/HelixConstitution.git HEAD` as re-measured **2026-09-03 after the fourth move**. Every earlier “checked out at X” in this lineage is WITHDRAWN, and the sequence is kept because it is the whole lesson: **a pin bumped to equal its upstream does not stay equal** — this one has now gone stale within a day on FOUR consecutive occasions, and the green below will not survive the night either.

```
902979027a90 -> f16ea779b82a    2026-09-01, operator-authorized,
0 divergent / 3 behind
f16ea779b82a -> f5876a3b700e    0 divergent / 4 behind
f5876a3b700e -> 3be10826f3d2    2026-09-02, operator-authorized,
0 divergent / 2 behind
remote HEAD 2026-09-02          = 3be10826f3d2    -- EQUAL, no
drift (WITHDRAWN)
remote HEAD 2026-09-03 (morning) = 2887b42e9349    -- DIFFERS,
direction UNDETERMINED (WITHDRAWN)
3be10826f3d2 -> 2887b42e9349    2026-09-03, operator-authorized,
0 divergent / 1 behind
remote HEAD 2026-09-03 (after)  = 2887b42e9349    -- EQUAL, no
drift
```

The fourth move was operator-authorized and CLASSIFIED BEFORE it was made, and the earlier UNDETERMINED reading was resolved by the authorized --fetch rather than by assumption. The claim this block carried — “*the direction is UNDETERMINED and this session did not resolve it*”, resting on `cat-file -t 2887b42e9349` returning “*could not get object info*” — was TRUE when written and is now **SUPERSEDED**: the fetch brought the object in, and `git merge-base --is-ancestor 3be10826f3d2 2887b42e9349` returned TRUE with `git rev-list --left-right --count 3be10826f3d2...2887b42e9349` returning **0 / 1 — 1 behind, 0 divergent**. `git merge --ff-only` performed it, which would have refused anything that was not a true fast-forward.

The fourth move touches NO governance document, and that was checked by name rather than eyeballed. `diff --stat across 3be10826f3d2..2887b42e9349` is **five files, 142 insertions, 20 deletions**, from the single commit “*fix(helix_code): hc_status reported UP services as DOWN — wrong port + wrong scheme*”:

```
docs/scripts/helix_code_services.docx    Bin 12354 -> 12805
docs/scripts/helix_code_services.html    47 ++++++--
docs/scripts/helix_code_services.md      39 ++++++--
docs/scripts/helix_code_services.pdf     Bin 48214 -> 41698
scripts/helix_code/helix_code_services.sh 76 ++++++-----
```

Filtering that name list against every governance path — Constitution.md, the four carriers, CLAUDE_ANCHORS_FULLL.md, find_constitution.sh, helix-deps.yaml, templates/, scripts/gates/, scripts/hooks/ and scripts/post_update_hook.sh — returns **NONE**. A service-status script and its four rendered exports; nothing this repository inherits changed.

Re-measured 2026-09-03 after the move. The index and the submodule working tree both sit at 2887b42e9349; the umbrella's HEAD still records f16ea779b82a, so the whole four-step move is **STAGED and not yet committed**. helix-deps.yaml records 2887b42e9349 staged in the same change, so gitlink and manifest moved together exactly as C9 requires — bash scripts/verify-manifest-pins.sh exits **0** at **12 MATCH / 0 DRIFT / 0 UNDETERMINED**, and bash scripts/verify-submodule-remote-sync.sh is back to exit **0** at **12 CURRENT / 0 DRIFT / 0 UNDETERMINED**.

The third move, kept because its lesson is the same one. It too was classified before it was made: merge-base --is-ancestor f5876a3b 3be10826 TRUE, rev-list --left-right --count 0 / 2 — **2 behind, 0 divergent** — and git merge --ff-only performed it. Its diff --stat touched **two files, both inside the constitution's own scripts/gates/**; no governance document, no carrier, not Constitution.md.

C9 went FAIL for exactly as long as the manifest was ahead of the index, and that is the gate working. Moving helix-deps.yaml first left the recorded ref at 3be10826 while git ls-files -s still reported f5876a3b; verify-manifest-pins.sh read 11 MATCH, 1 DRIFT until the gitlink was staged, then returned to **12 MATCH / 0 DRIFT**. **Stage the gitlink and the manifest together, or the gate will — correctly — call you out.**

Do not read any SHA in this file as current; run the gate. Two consecutive sessions have now bumped this pin and found it behind upstream again within the day. Treat the constitution pin as a standing operator decision, not a task that completes.

Nothing this repository records about the corpus has moved in any of the four bumps, and the fourth was re-measured on both sides of the move rather than after it. Constitution.md is the SAME git blob ad32fc94c3ce80bad38b4d97ff3b8d83a236e860 at all five pins. Read on disk at 3be10826f3d2 and again at 2887b42e9349, every figure reproduced identically:

	at 3be10826f3d2	at 2887b42e9349
Constitution.md blob	ad32fc94...	ad32fc94...
lines	11,787	11,787
### § anchors	252	252
bytes	1,801,862	1,801,862
sha256	973e473c823b4a13...	973e473c823b4a13...
du -h CLAUDE.md	784K	784K
du -h Constitution.md	1.7M	1.7M

No figure in this file needed correcting for the fourth move — which is a measurement, not a rule. The second bump’s diff --stat likewise touched two paths, neither a governance document: design-toolkit (1 -) and docs/codegraph/Status.md (18 +). Four corpus-neutral moves in a row are a measured coincidence of what upstream happened to change, **not** a guarantee that a pin move is corpus-neutral — re-measure after the next one, on both sides.

Four findings from the FIRST bump, each measured rather than inferred, kept because each still describes how a pin move must be done here:

- 1. The move was a true fast-forward, not a rebase.** git merge-base --is-ancestor 902979027a90 f16ea779b82a returned TRUE and git rev-list --left-right --count <pin>...<remote> returned 0 / 3: **0** commits existed on the pin that upstream lacked, **3** existed upstream that the pin lacked. The move itself was performed with git merge --ff-only, which would have refused anything else. The local branch carried no unpushed work — 902979027a90 was a local commit made in this checkout, and it is an ancestor of the remote head, so nothing was lost. **Nothing was pushed to the constitution repository**; this was a local pin move only.
- 2. No recorded anchor count or line count in this repository went stale.** Constitution.md is literally the same git blob at both commits — ad32fc94c3ce80bad38b4d97ff3b8d83a236e860 — re-measured ON DISK AFTER the checkout at 11,787 lines and 255 ### § anchors, sha256 973e473c823b4a13... identical to the old pin’s. git diff --stat across the three commits touches three paths and none is a governance document: design-toolkit (1 +), docs/codegraph/Status.md (12 +), and the submodules/design-toolkit gitlink. Every figure recorded in this tree still describes the pinned text.
- 3. helix-deps.yaml moved in the SAME change — and did so again at the second bump.** Its constitution entry records 3be10826f3d2

today, staged together with the gitlink; a bumped gitlink with a stale ref: is precisely what C9 exists to catch, and it has now had two chances to catch one and had nothing to report either time. Re-measured 2026-09-02: `bash scripts/verify-manifest-pins.sh` exit **0** at 12 MATCH / 0 DRIFT / 0 UNDETERMINED, and `bash scripts/verify-governance-cascade.sh` exit **0** at 12 PASS / 0 FAIL / 0 ENV / 8 NOTE.

4. **No gate in this tree detected local-vs-remote submodule drift, and the bump did not change that. One was written later the same day, and it is RED.** Cascade check C9 and `scripts/verify-manifest-pins.sh` compare `helix-deps.yaml` to the **local gitlink**, never to the remote. Both exited **0** while the pin sat 3 commits behind, and both exit **0** now. They cannot tell the two states apart, so neither exit code is, or ever was, evidence about the remote — which is precisely why a second instrument was needed. See “Submodule-vs-remote drift” below.

Upstream defect carried in by the FIRST fast-forward — now FIXED UPSTREAM by the second, and the fix is measured, not assumed.

f16ea779b82a added a SECOND design-toolkit gitlink at the constitution’s own ROOT (`create mode 160000 design-toolkit`) alongside the existing submodules/design-toolkit, while its `.gitmodules` mapped the entry NAMED design-toolkit to path submodules/design-toolkit only — leaving the root gitlink unregistered. This repository correctly declined to fix it, because fixing it means committing to the constitution submodule, which this repository does not do. **Waiting was the right call: upstream removed it.** Re-measured 2026-09-02:

```
git -C submodules/constitution ls-tree f16ea779b82a --format='%
(objectmode) %(path)' | grep design-toolkit
# 160000 design-toolkit      <- the unregistered root
    gitlink
git -C submodules/constitution ls-tree HEAD --format='%
(objectmode) %(path)' | grep design-toolkit
# (none)                  <- gone at f5876a3b700e
```

That deletion is the design-toolkit | 1 – line in the second bump’s diff --stat. **Record the precedent:** a defect in a CONSUMED submodule was resolved by waiting for upstream, not by committing into a repository this project does not own.

Separately, the design-toolkit ref CONFLICT recorded under `conflict_resolution` in `helix-deps.yaml` is a DIFFERENT thing and is **UNCHANGED**: the constitution’s own `helix-deps.yaml` says ref:

16e4e76 at all five pins — 902979027a90, f16ea779b82a, f5876a3b700e, 3be10826f3d2 and 2887b42e9349, each read directly with `git show <pin>:helix-deps.yaml`. Do not read the gitlink fix as closing this.

The residual risk this block used to report as “not started” is now WATCHED, and the watching instrument is RED. The risk was real and is restated because the reason it went unseen is the lesson: any submodule in this tree could sit arbitrarily far behind its upstream — corpus changes included — with every gate in this repository green, because C9 compares the manifest to the local gitlink and **nothing compared the gitlink to the remote**. The bump closed one INSTANCE, not the CLASS. The class is now covered by `scripts/verify-submodule-remote-sync.sh`, registered in `scripts/check-registry.tsv` as `submodule-remote-sync` with a `--prove-failure` paired proof — which is exactly what R5 demanded. **Watched became FIXED, then RED, then FIXED again — inside three days.** Re-measured **2026-09-03 after the fourth fast-forward**, the gate exits **0 at 12 CURRENT / 0 DRIFT / 0 UNDETERMINED** of 12 owned gitlinks probed. The reading this block carried earlier the same day — *“exits 1 at 11 CURRENT / 1 DRIFT, and the single DRIFT row is submodules/constitution itself”* — is **WITHDRAWN as current** and was true when written. Every reading in the lineage is now historical: “6 CURRENT / 6 DRIFT” (2026-09-01), “11 CURRENT / 1 DRIFT” (earlier 2026-09-02), “12 CURRENT / 0 DRIFT” (later 2026-09-02), “11 CURRENT / 1 DRIFT” (2026-09-03 morning), “12 CURRENT / 0 DRIFT” (2026-09-03, after the fourth move). **Green is a measurement of today, not a property of the tree** — the gate says so itself, and TWO cycles of green-then-red-then-green are the proof. **Do not bank this green; the pin has broken it four times.** See “Submodule-vs-remote drift” below.

FIFTH FAST-FORWARD, 2026-09-08, operator-authorized — AND IT IS THE FIRST THAT IS NOT CORPUS-NEUTRAL. EVERY FIGURE IN THE TABLE BELOW IS WITHDRAWN. The pin moved 1f672725c823 -> 573c01905564, classified BEFORE the move as **0 divergent / 5 behind** and performed with `git merge --ff-only`, which would have refused anything else. `git submodule update --init --recursive` inside the submodule returned rc 0.

Constitution.md CHANGED, so the four bumps’ “same blob” lineage ENDS HERE. Re-measured on disk at 573c01905564:

	at 1f672725c823 (WITHDRAWN)	at 573c01905564 (measured)
lines	11,787	11,794

	at 1f672725c823 (WITHDRAWN)	at 573c01905564 (measured)
### § anchors	255	255
bytes	1,801,862	1,804,638
sha256	973e473c823b4a13...	b9e438c95e7a7fd3...

diff --stat across the five commits is **7 files, 323 insertions, 11 deletions**: Constitution.md (11 ±), all four carriers (2 ± each), scripts/hooks/guard-forbidden-commands.sh (177 ±) and a new scripts/hooks/test_guard_forbidden_commands.sh (138 +). **This is the first bump in this lineage to touch a governance document**, so the standing sentence that four consecutive corpus-neutral moves were “a measured coincidence of what upstream happened to change, not a guarantee” is now demonstrated rather than merely asserted. Re-measure after the next one, on both sides.

STANDING AUTHORIZATION for fast-forward-only constitution-pin bumps (2026-09-09)

The operator has converted the per-bump approval into a STANDING one. Five of the six moves in the lineage above were authorized one at a time, and the pin went stale again within the day on four consecutive occasions; the approval is now given once, with the terms fixed. **The terms are the authorization — a bump that does not satisfy every one of them is NOT authorized:**

1. An agent MAY `git fetch` inside submodules/constitution and fast-forward it **without asking again**. The sentence that stood here — “*a fetch is a mutating command and an operator decision*” — is **SUPERSEDED for this one submodule**, and for no other.
2. ONLY when `git merge-base --is-ancestor <pin> <remote>` returns TRUE **and** `git rev-list --left-right --count <pin>...<remote>` shows **0 divergent**. **ANY divergence: STOP and report. Do not merge.**
3. ONLY via `git merge --ff-only`. Never rebase, never `--force`, never `--force-with-lease`, never a `+ref` push. §11.4.113 is absolute and this authorization does not touch it.
4. The `diff --stat` across the moved range **MUST** be reported, and Constitution.md **MUST** be re-measured on **BOTH** sides of the move — blob sha, line count, ### § anchor count, byte count, sha256. **The 2026-09-08 bump was the first in this lineage to change Constitut**

ion.md: corpus-neutrality is NOT a property of a pin move, it is a measurement that must be taken each time.

5. The gitlink and the helix-deps.yaml deps[].ref entry **MUST** be staged **together in one change**, or cascade check C9 will correctly report DRIFT — as it has, twice.
6. **Nothing is ever pushed to the constitution repository.** This is a local pin move only.

Scope, stated so it is not over-read: this authorization covers submodules/constitution and NOTHING else. Every other owned gitlink still requires a per-bump operator decision, even when verify-submodule-remote-sync.sh reports it BEHIND and even when the move would be a clean fast-forward.

Re-derive. The pin equals the remote head as re-measured 2026-09-09, and that sentence has already been true and then false three times. The probes below are read-only and tell you THAT it differs; under the standing authorization above, a --fetch may then be run to classify WHICH WAY:

```
git -C submodules/constitution rev-parse HEAD #
    the local pin
git ls-remote git@github.com:HelixDevelopment/
    HelixConstitution.git HEAD
git -C submodules/constitution rev-parse HEAD:Constitution.md #
    blob identity
grep -c '^### §' submodules/constitution/Constitution.md #
    255
wc -l < submodules/constitution/Constitution.md #
    11794
bash scripts/verify-manifest-pins.sh #
    0 = ref == gitlink
```

Re-measured 2026-09-09, and the whole fleet is level — no move was made and none was needed. git -C submodules/constitution rev-parse HEAD and git ls-remote ... HEAD both return 573c019055646e72aeb1ee22dabb9005e94e855e, so the standing authorization had nothing to act on the day it was granted. bash scripts/verify-submodule-remote-sync.sh exits **0** at **13 CURRENT / 0 DRIFT / 0 UNDETERMINED** of **13 owned gitlinks probed, 14 declared** — the first reading in this file’s history with no BEHIND row anywhere in the fleet. bash scripts/verify-manifest-pins.sh exits **0** at **13 MATCH / 0 DRIFT / 0 UNDETERMINED**. Constitution.md re-measured on disk at that pin: blob d95bf0c4f47c..., **11,794** lines, **255** ### § anchors, **1,804,638** bytes, sha256 b9e438c95e7a7fd3... — every figure identical to the post-09-08 table above, which is what “re-measure, do not quote” buys you.

The fleet grew by one and this file had not recorded it. `submodules/curriculum-kit` (`vasic-digital/curriculum-kit`, at `07af6fb79e20`) joined the owned set; it is probed by both gates and is CURRENT. Three counts moved with it, each quoted from its own gate rather than derived by hand: `verify-submodule-remote-sync.sh` probes 13 owned gitlinks of 14 declared (was 12 of 14); cascade C1 classifies 12 owned + 1 governance source + 1 third-party (was 11 owned); C2/C3 now find $12 \times 4 = 48$ owned-submodule carriers (was $11 \times 4 = 44$). Note the two vocabularies disagree on purpose — the remote-sync gate probes `submodules/constitution` and counts it among its 13, while C1 classifies it as the cascade's SOURCE and excludes it from its 12. **Every “11 owned”, “12 owned gitlinks”, “44 of 44” and “12 CURRENT” figure elsewhere in this carrier is SUPERSEDED.** Re-derive the roster from the gates; never quote it from here.

Do not read the 13/13 green as durable. This carrier has recorded five green-then-red cycles on this one pin; the standing authorization exists precisely because the green does not survive the week.

Note the path: this repository places the submodule at `submodules/constitution/`, not at the top-level `constitution/` that the constitution's own prose and templates use as their example. Both layouts are supported by `find_constitution.sh` (its candidate list is `constitution` then `submodules/constitution`). Whenever a quoted constitution snippet says `constitution/<file>`, the file in THIS repository is `submodules/constitution/<file>`.

If the submodule directory is empty, it has not been initialised. Initialise it read-only-safely with `git submodule update --init submodules/constitution` before relying on any rule below. Do not run that, or any other mutating git command, without operator authorization.

Why a pointer block and not an @import

Anchor 11.4.35 invariant 6 declares the two inheritance forms **equivalent**: a consumer carrier MUST start with either the native `@constitution/CLAUDE.md import` OR the portable `## INHERITED FROM ...` pointer block. This repository uses the pointer block. The reason is stated openly rather than left implicit: `submodules/constitution/CLAUDE.md` is 784 KB and `Constitution.md` is 1.7 MB (re-measured 2026-09-03 at pin `2887b42e9349` with `du -h`, unchanged by the third or fourth bump; both grew from the 654 KB / 1.53 MB figures this file used to carry), so a native import would load roughly 165k tokens of context into every session before any work begins. Re-derive with

du -h rather than trusting these numbers — they are a dated observation, and the pin moves. The pointer form carries the same authority at no standing cost, and it is the form the constitution's own propagation gates recognise (submodules/constitution/scripts/gates/lib/pointer_carrier.sh classifies any carrier whose first non-fenced line-anchored heading is ## INHERITED FROM as a legitimate pointer-inheritance consumer).

Read the canonical text on demand, not eagerly. Useful entry points:

```
# Locate the submodule from any nested depth
bash submodules/constitution/find_constitution.sh

# Read one anchor instead of the whole corpus
grep -n '^### §11.4.35 ' submodules/constitution/Constitution.md
awk '/^### §11.4.35 /{f=1} f{print} f&&/^### §11.4.36 /{exit}' \
    submodules/constitution/Constitution.md
```

Critical base rules restated (for agents that do not resolve pointers)

The nine items below are reproduced verbatim from the constitution's own consumer scaffold submodules/constitution/templates/AGENTS.project.md.template. That bounded restatement is authored by the constitution for exactly this purpose, so reproducing it is the prescribed mechanism — it is not a copy of the corpus. The corpus itself is never duplicated here.

- **No bluffing.** Every PASS carries positive evidence. Constitution §11.4.
- **Mutation-paired gates.** Every new gate has a paired mutation proving it catches regressions. Constitution §1.1.
- **No guessing language.** likely, probably, maybe, seems, appears etc. are forbidden when reporting causes. Constitution §11.4.6.
- **Credentials never tracked.** .env patterns git-ignored; runtime-load only; per-service file separation. Constitution §11.4.10.
- **Use the project's commit wrapper.** No direct `git add / git commit / git push` on main repo.
- **Never force-push.** Force-push requires explicit per-session authorization AND a green §9.1.5 post-op gate.
- **Hardlinked backup before any destructive op.** Constitution §9.
- **CONTINUATION.md kept in sync** in every non-trivial commit. Constitution §12.10.

- **60% RAM cap.** Heavy work wrapped in bounded execution scope. Constitution §12.6.

Anything not covered by those nine items is covered by the canonical files. Go read them; do not guess.

Content boundary — do NOT quote private material into this public repository

This is a **PUBLIC** repository (private=false, verified against the provider 2026-09-01). **Three** of its submodules are **PRIVATE**: workshop (milos85vasic/workshop_curriculum), ai_interviewing, monetization. A gitlink exposes only a commit SHA, so their content stays private — until somebody copies it out.

design-toolkit is no longer one of them. It was flipped **public** on 2026-09-01 after a clean full-history privacy audit, so the four-private roster this paragraph used to carry is **WITHDRAWN**, not restated. Re-derive the whole roster rather than trusting this sentence — visibility is a provider setting and changes outside this tree:

```
git config -f .gitmodules --get-regexp 'submodule\..*\.'url'
# 13 URLs
gh api repos/<owner>/<name> --jq '.visibility' #
# per repository
```

Measured 2026-09-01 with exactly that pair: milos85vasic/workshop_curriculum, milos85vasic/ai_interviewing and milos85vasic/monetization are **private**; vasic-digital/design-toolkit, vasic-digital/containers, vasic-digital/LLMProvider, vasic-digital/RAG, vasic-digital/verdict, vasic-digital/passage, vasic-digital/vasic-digital.github.io and milos85vasic/milosvasic.ru are **public**.

The rule: naming a private path is fine; copying what is inside it is not.

```
docs/x.md: see workshop/chapters/01/<recording>.mp4      <-
reference, allowed
docs/x.md: the notes say "<sentence from that recording>" <-
LEAK, forbidden
```

One pasted paragraph is public permanently and irreversibly: it is then in a public repository's history, and history is not editable after a push. The pressure that causes this is diligence, not carelessness — an agent asked to document the work comprehensively reaches for an illustrative quote, and the quote is the disclosure. workshop/chapters/**

holds a recording of a private teaching session with an identifiable third party; treat it, and the notes derived from it, as quotable by path only.

Enforced by `scripts/verify-content-boundary.sh` (registered in `scripts/check-registry.tsv` as `content-boundary`; paired proof --prove-failure; three-valued exit where a repository whose visibility cannot be established is 2, never a pass). Legitimate cross-boundary flows are declared as private/public **pairs** with a mandatory reason in `.content-boundary-allow`. Full design, the measured false-positive reasoning and the fleet map: [docs/content-boundary.md](#).

This gate is RED on this tree today, and that is the designed state — do not make it green. Re-measured 2026-09-03, `bash scripts/verify-content-boundary.sh` exits **1** and prints:

```
LEAK — 15308 surviving match(es) (prose 14552, short 642, name
114); 19 row(s) also could not be determined
corpus MOVED — 18 of 12535 enumerated file(s) changed between the
pre- and
      post-analysis fingerprints
```

THE “NOT REPRODUCIBLE RUN-TO-RUN” FINDING IS RESOLVED, AND THE EARLIER VERDICT — “*the exact mechanism is UNDETERMINED and was not established*” — IS WITHDRAWN AS SUPERSEDED, measured 2026-09-04. It was true when written.

The algorithm is DETERMINISTIC. Four runs against a provably frozen copy of the fleet produced **byte-identical output** — 12939 every time (prose 12368, short 486, name 85), with every `CB_TRACE` stage count identical at every stage. The snapshot was a `cp -al` outside the repository with tracked files de-hardlinked, verified frozen for 60 s before use and bracketed by fingerprints around every run.

The mechanism was CONCURRENCY, not the gate. The live tree was never quiescent. A fingerprint of exactly what the gate reads, sampled every 20 s, showed the corpus changing on nearly every sample: two tracked files inside the private `workshop` submodule were being rewritten continuously by another agent, and **both are private-side key sources**, so each edit changes the private key space and therefore the total. Same command, no gate or allow-list edit, live tree: **12939 / 12939 / 13058 / 12968**. See [\[\[measure-on-a-quiet-tree\]\]](#).

The three-run table below is HISTORICAL and is kept only so a stale reading is recognisable rather than trusted. Its own conclusion — that nothing changed during the window, based on `find -newermt '3 hours ago'` returning 0 files — was the measurement that failed: it was taken over the wrong window and the wrong set.

run	tree	leaks	prose	short	name	already_public
1	pre-edit carriers	12745	12182	478	85	2252
2	post- edit carriers	12846	12280	481	85	2261
3	pre-edit carriers (same tree as run 1)	12867	12299	483	85	2261

THE GATE NOW CARRIES ITS OWN EVIDENCE OF STABILITY — the residual defect was its SILENCE, and that is fixed. It fingerprints the exact scanned set before and after the analysis passes. The enumeration is shared with the analysis by construction (`cb_ls_tracked / cb_ls_untracked` are called by both), because a second enumeration that merely agrees today would be false reassurance. A green run now prints positive evidence; a moving run emits undet rows **naming the changed paths** — paths only, never content — and says the counts are not reproducible. `--expect-corpus <file>` lets two figures be **proved** to come from the same tree, which is what makes any future re-baselining defensible. Precedence is unchanged and is asserted by a mutation: **a moving tree cannot mask a finding** (`LEAKS>0 → rc 1 outranks UNDET → rc 2`).

Its first real run caught the tree moving, which is the point:

```
LEAK – 15308 surviving match(es) (prose 14552, short 642, name
114);
    19 row(s) also could not be determined
corpus MOVED – 18 of 12535 enumerated file(s) changed between the
pre- and
    post-analysis fingerprints
```

8 APPEARED, 10 CHANGED, **17 of the 18 inside the private workshop submodule** — consistent with nine agents working in one checkout. `--prove-failure` is **0 at 59 passed / 28 mutations** (was 43 / 24).

HONEST BOUNDARY. The population has moved far past every figure recorded above, and **no row was judged**: no claim is made about what part of 12867 → 15308 is new content versus a moving tree, and a large document-export run authorised by the operator is adding files to this repository concurrently. **Treat 15308 as a reading of a tree that was measurably moving**, not as the count. The gate now says so itself, which is the whole improvement.

What survives the instability, and it is the part that matters: the class A direction split is stable to within 0.6 percentage points across all three runs, with **OUTWARD 6240–6241 and INWARD 1538–1544 in absolute terms**. The rows that move between runs land in UNDETERMINED. See “Class A direction” below.

Every earlier reading is **SUPERSEDED**: LEAK – 11878 (prose 11356, short 433, name 89) from earlier the same day, the **2026-09-02** LEAK – 11158 (prose 10833, short 268, name 57), and LEAK – 285 (prose 207, short 35, name 43). Note the name class fell **89 → 85** and then held at 85 across all three runs while every other class moved; that is the population changing, not a clearance. **Treat the percentages and the class shares as the durable figures and the absolute totals as ±150.**

THE POPULATION IS FULLY ATTRIBUTED — every row of each run is accounted for by path, not by sampling. Every run was taken with --json, so the tables here are aggregated from the gate’s own structured row set rather than re-parsed out of a log, and each aggregate matches its run’s own total exactly. **Figures below are from run 3 (12867);** read the class SHARES, which move by less than 0.3 percentage points across the three runs, rather than the counts.

By private source: **workshop 12590 (97.8%)**, ai_interviewing 277 (2.2%), monetization **0** — the 97.8 / 2.2 split is identical in all three runs. By public destination, specs/** holds **7890**, submodules/** ~2800, docs/** 502, scripts/** 489, CONTINUATION.md 376 and the four root carriers 408 (**102 each**, superseding the 81-per-carrier figure). Inside specs/**, **specs/002-** is 6104 and specs/001-** is 1786** — the two spec trees are **61.3% of the entire finding on their own** (61.3–61.5% across the three runs).

The load-bearing limitation is unchanged: this gate detects CO-OCCURRENCE, not DIRECTION. Governance text that originated in these PUBLIC carriers and propagated INTO a private submodule is indistinguishable, to it, from private content leaking OUT. **On 2026-09-03 the direction was measured independently, by git first-**

commit timestamp, for THREE buckets — the two below, and then class A, the largest, which had none until that day. Class A is the one to read: see “Class A direction” in the decision packet. Its answer is NOT uniformly outward.

- **Root carriers (324 matches over the four files).** CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md were first committed **2026-08-27T00:41:30**. All six sampled private counterparts were first committed **2026-09-01 or later** (earliest 2026-09-01T21:02:46). The public text pre-dates the private text by five days. **Direction: outward. Cascade, not disclosure.**
- **submodules/RAG (1741 matches). 1623 of them — 93% — come from one private directory, workshop/platform/upstream-contributions/**, whose name states its own role. submodules/RAG/pkg/grounding was first committed **2026-09-01T11:24:47**; that private staging directory first appeared **2026-09-01T21:02:46**, about ten hours LATER. **The public package is the older artifact.** submodules/LLMProvider has the same shape: 617 of its 692 matches come from workshop/platform/gates/.

The 94-matches-per-carrier figure this block used to carry is superseded: it is 81 per carrier, 324 across the four.

HONEST BOUNDARY (§11.4.6) — read this before quoting anything above. 1. **The two probes in this section are FILE-level and SAMPLED** — a first-commit comparison on 6 of 16 private sources for the carrier bucket, and on a directory for RAG. Neither is a per-string provenance trace. A string can move between files, and that was not traced for these two buckets. **The class A probe described in the packet below is a different and stronger instrument: TEXT-level, per-string, over the complete class with nothing sampled.** Do not quote the weaker method’s caveat as though it applied to class A, and do not quote class A’s rigour as though it applied to these two. 2. **The attribution table is complete; a row-by-row JUDGEMENT is not.** Every row is placed by path. **No row was judged this session, and none was redacted, allow-listed or re-baselined.** 3. **The judged population and today’s population are barely the same set.** The incident note’s two-wave assessment judged 232 matches (29 redacted as real disclosures, 203 judged not to be) on **2026-09-01**, and its own split was ai_interviewing **216** / workshop **16**. Today the split is workshop **12468** / ai_interviewing **277** — the dominant source has **inverted**. **So the unassessed remainder cannot be computed as <today's total> – 232.** That subtraction assumes the 232 judged rows are a subset of today’s ~12,800, and the source split shows they largely are not: 216 of the 232 were ai_interviewing rows, while 97.8% of today’s population

is workshop. **How much the two sets overlap was NOT measured.** Do not report any prose row as “cleared”. 4. **The name class is NOT cleared: 85 matches carrying 9 distinct withheld name identifiers, reported by digest only.** The figure this file carried — “89 matches, 8 distinct personal names” — is **superseded by measurement**, not by any clearance: the class went 57 → 89 → 85 while the detector was never touched. **A falling name count is not progress.** Direction was NOT established for a single name row, and cannot be with this instrument: the gate withholds the matched text by design, which is the same protection that keeps a real person’s name out of an archived artefact.

Re-baselining this figure is an **operator decision** and this session did not take it. See “Content-boundary re-baseline — the decision packet” below.

An EARLIER figure in this lineage — LEAK – 10293 surviving match(es) (prose 203, short 8535, name 1568) — was withdrawn before the 285 reading, and is recorded here so a stale one is recognisable rather than trusted. It was a MID-TUNING snapshot, captured while the detector was still being rewritten, and it measured the detector’s own noise floor rather than this tree. **The two numbers are not a before/after of any cleanup: nothing was redacted between them.** The short class fell 8535 → 35 and the name class 1568 → 43 because the detector learned to subtract keys already public in two or more public repositories, public git identity forms, a scale-free name-frequency floor and path-reference matches — each subtraction printed on every run, with its RECALL COST stated, so what the gate can no longer see is visible rather than assumed.

The prose class is the judged class — and the judging has NOT kept pace with it. At the 285 reading, 207 prose rows stood against an incident-note two-wave assessment covering 232 matches (29 redacted as real disclosures, 203 judged not to be) — so 207 and 203 were close but **not the same set**, and the difference was never accounted for. **At 10,833 prose rows the gap is no longer a rounding difference: the assessed population is 232 and the reported population is two orders of magnitude larger.** Do not read any prose figure as “rows individually cleared”. The judged rows are left VISIBLE rather than allow-listed, because an allow-list entry buys a green exit at the cost of hiding the row from the next reader. The short and name classes are substring noise the note characterises. **Do not “fix” this gate by adding exemptions, and do not add more.** .content-boundary-allow still carries the same **9 declared pairs**, pardoning **4081** matches as re-measured 2026-09-03 (the 4080 and 4193 figures are superseded; the

file is unmodified, so the count moved with the corpus, not with the allow-list). Every entry needs a real cross-boundary flow and a mandatory reason, and the operator decides.

Content-boundary re-baseline — the decision packet
(prepared 2026-09-03, NOT decided)

Prepared so an operator can decide; **nothing here was acted on**. No row was judged, no allow-list entry added, no baseline moved, and the gate still exits 1.

Three structural classes account for 97.3% of the rows, counted exactly rather than sampled — and the SHARES are stable across all three runs even though the totals are not. Each has a DIFFERENT correct disposition. Counts are run 3:

Class	Rows	Share	What it is	Direction evidence
A · spec trees in this public umbrella	7890	61.3%	specs/002-** 6104 + specs/001-** 1786, matching workshop/docs/session-evidence/, platform/backend/, pipeline/extract/	MEASURED 2026-09-03, text-level, complete, and REPRODUCED on all three runs: ~79% OUTWARD, ~19.5% INWARD. See “Class A direction” below. Not uniformly outward — roughly one row in five is private-first.
B · public reusables extracted from the private tree	2703	21.0%	submodules/RAG 1742, LLMProvider 696, passage 220, containers 19, verdict 9,	File-level and sampled. Direction measured for RAG: the public package pre-dates the

Class	Rows	Share	What it is	Direction evidence
			plus design-toolkit 17	private staging copy by ~10 h, and 93% come from workshop/platform/upstream-contributions/.
C · governance prose that propagated OUTWARD	1922	14.9%	umbrella docs/ ** 502, scripts/ ** 489, 4 root carriers 408, CONTINUATION.md 376, submodules/ constitution 125, Constitution.md 22 — matching workshop/docs/ **	File-level and sampled. The carriers are five days OLDER than every sampled private counterpart. This is the §11.4.157 cascade working.
Remainder — genuinely unassessed	352	2.7%	_tests 88, _content 61, design-system 58, _analysis 39, .specify 37, and a long tail	Not characterised. Not cleared. No direction probe run.

The 85 name-class rows do NOT sit in the remainder — they are spread across every class: A 50, B 19, remainder 14, C 2, identical in all three runs. No class is name-clean, so no class can be pardoned wholesale without pardoning name rows with it. The earlier split “A 55, remainder 18, B 14, C 2” over 89 rows is superseded.

Class A direction — measured 2026-09-03, and it is the strongest probe in this file

Method, stated so it can be attacked. Every one of the 7835 class A rows was dated on BOTH sides at **text level, not file level**: for each matched string, walk every commit that ever touched the file it was found in, normalise that blob with the gate’s own normalisation ([^A-Za-z0-9]+ → space, lowercased — validated against the gate on a 24-

of-24 sample), and take the earliest commit whose content actually contains the string. That dates the TEXT, not the file that carries it, which is the defect in the class B and C probes. It was then **widened to corpus level**: for every row, the earliest appearance of that string ANYWHERE in the public umbrella was compared against its earliest appearance ANYWHERE in the private repository, so a string that moved between files is still caught. 55 public and 489 private path-histories were walked for the pairwise dating; the widening ranged over the **9,414 file-revisions** the two histories hold (umbrella 8,213, workshop 1,201). **Nothing sampled, nothing extrapolated** — the walk is chronological and skips a blob only when no still-undated string could be affected by it, which is an exact pruning rather than a shortcut: a string already dated cannot be dated earlier by a later commit. Verified for the one case that could have biased it — no string was found in both private repositories, so the two-repo merge changed no date.

The probe was run independently against ALL THREE gate runs, and this is the answer to the instrument's own instability: the totals move, the direction does not.

run	class A rows	OUTWARD	INWARD	UNDETERMINED
1	7835	6241 (79.7%)	1544 (19.7%)	50 (0.6%)
2	7872	6240 (79.3%)	1538 (19.5%)	94 (1.2%)
3	7890	6240 (79.1%)	1538 (19.5%)	112 (1.4%)

OUTWARD varies by ONE row and INWARD by SIX across three runs whose totals differ by 122. Every extra row a noisier run emits lands in UNDETERMINED — these are strings not committed on one side, i.e. working-tree-only text. **The direction finding is therefore robust to the count instability**, and it is the figure to act on.

The widening moved rows in BOTH directions, which is why it is not a rationalisation: on run 1, 113 rows flipped INWARD → OUTWARD and **178 flipped OUTWARD → INWARD**. A one-sided correction would have been evidence of a biased probe.

Lead times. OUTWARD: median 6.1 h, max 611 h. INWARD: median 0.5 h, max 465 h, with **92% of inward rows inside 24 h** — spec and implementation are being written the same day, in both orders.

Where the INWARD rows sit, because that is what an operator must look at: of the 1544 on run 1, **1084 (70.2%) are private SOURCE CODE**, 153 private session-evidence briefs, 117 private training docs,

56 other private docs, 40 data/config, 26 other — and **68 are workshop/chapters/01/, the private teaching-session material**, all from one JSON artefact, landing in specs/001-.../tasks.md (45) and specs/002-.../tasks.md (23). By public file the inward rows concentrate in specs/002-.../contracts/knowledge-graph.md (431), specs/002-.../tasks.md (410) and specs/001-.../tasks.md (279). By gate class: 1468 prose, 76 short, 0 name.

HONEST BOUNDARY — three limits, and the first is the one that matters. 1. **“Private committed first” is NOT a finding of disclosure.** It establishes an ORDER, nothing more. A task line written in a commit message, a symbol name, or a design decision made in code and then written up in the spec all produce INWARD rows with no private content having left the boundary. **~1540 rows are not ~1540 leaks, and must never be reported as such.** What the number does establish is that the reassuring story — *“the specs are outward propagation”* — is true for four rows in five and **false for the fifth.** 2. **The 50 UNDETERMINED rows are the entire class A name population**, all in specs/001-.../research/transcription.md, all carrying ONE withheld identifier. They are undetermined because the gate REDACTS the matched text, so there is nothing to date — the protection and the blind spot are the same mechanism. A weaker FILE-level fallback splits them 30 outward / 20 inward, and **the 20 inward include 15 rows whose private side is an ai_interviewing document first committed 2026-08-12, twenty days before the public spec file, and 5 whose private side is the private recording-notes PDF first committed 2026-08-31.** That is file-level only, it is weaker than the text-level result, and it is **not** a determination — it is recorded because it points at the rows an operator should open first. 3. **No row was judged, redacted, allow-listed or re-baselined**, and the gate still exits 1.

What an operator is actually being asked to decide — four options, with the cost of each stated rather than implied:

1. **Do nothing.** The gate stays red and stays readable. Cost: the signal keeps degrading — 232 → 11,158 → 11,878 → ~12,800 in three days — until nobody reads it. **And the total is now known not to be reproducible run-to-run**, which makes “watch the number” a weaker plan than it looks.
2. **Allow-list classes B and C** (2700 + 1861 = 4561 rows, as declared pairs with reasons). Cost: **an allow-list entry hides the row from the next reader**, which is exactly the failure mode this file warns about elsewhere; it would pardon **21 name rows** along with them; and the direction evidence behind both classes is **file-level and sampled**, which is weaker than what class A now has.

3. **Teach the gate direction** — a git-timestamp or provenance pass, so outward propagation is subtracted with its RECALL COST printed, the way the detector already subtracts already-public keys. Cost: real work. **The class A probe is a working prototype of exactly this**, run offline over 9,414 file-revisions in under a minute, so the feasibility question is answered. It is still the only option that raises precision without hiding a row.
4. **Judge the ~1540 class A INWARD rows** as a fresh assessment wave — **not all ~7870**. The direction probe has cut the reading assignment by 80% and named where it concentrates: 1084 rows against private source code, 68 against the private teaching-session material, three public files holding 1120 of them. Cost: still the largest, and still the only option that could find a genuine new disclosure.

Recommended ordering — a recommendation, not a decision: 4 (now bounded to the ~1540 inward rows plus the 50 undetermined name rows), then 3, then 2 for whatever survives. The reason for putting 4 first has CHANGED: class A is no longer the class with no direction evidence — it now has the best evidence in this file, and that evidence is what makes the reading assignment small enough to be worth starting.

Do NOT read this packet as clearance. It places every row by path; it judges none. Options 2 and 4 both require reading private material and are operator work.

A gate that a human must read is worth more than one that is quietly green.

Project overview

vasic is the umbrella monorepo for two personal/portfolio sites and the shared tooling that builds, translates and validates them. vasic.digital/ is committed static HTML served as-is; milosvasic.ru/ is Jekyll source whose rendered _site/ is **GIT-IGNORED AND UNTRACKED, re-measured 2026-09-08**. `git -C milosvasic.ru ls-files _site` returns **0 files** and `check-ignore -v` resolves it to `.gitignore:70`, while 29 entries sit on disk unversioned. **The claim this sentence carried until today — “TRACKED, not git-ignored”, itself a withdrawal of an earlier “ignored” claim on 2026-09-06 — is now SUPERSEDED, and it was true when written.** What changed is not a measurement error but an EXECUTED DECISION: operator decision #13 of 2026-09-07 (“Untrack it, keep the ignore rule”) was carried out,

and this carrier was not updated with it. Record the lesson rather than only the fact: a decision executed without updating the document that describes the state leaves a carrier asserting the pre-decision world, and the next reader cannot tell a stale claim from a false one. `_tools/gen/` is the Go generator that renders the localized pages for both sites, `design-system/` holds the shared per-brand tokens and component CSS, and `_tests/` is the Playwright plus self-validating harness. English source content lives in `_content/` with per-language siblings in `_content_{lang}/`. See `README.md` for the full map.

Project-specific facts

These are facts about this repository, not additional rules. They exist so an agent does not have to guess. The authoritative source for each is `README.md`.

Owned submodules

`.gitmodules` declares **14** gitlinks (measured 2026-09-01). **11** are owned consumers of the governance cascade: `milosvasic.ru`, `vasic.digital`, `design-toolkit`, `ai_interviewing`, `monetization`, `workshop`, `submodules/containers`, `submodules/LLMProvider`, `submodules/RAG`, `submodules/verdict` and `submodules/passage`. `submodules/constitution` is the governance **source** — the head of the cascade, not a consumer. `submodules/superspec` is third-party (upstream WangX0111/superspec) and is outside the owned-submodule set for tagging and propagation purposes.

Four of those 11 joined on 2026-09-01, and the earlier “nine gitlinks, seven owned” figure is **WITHDRAWN**, not restated:

- `submodules/LLMProvider` and `submodules/RAG` were **adopted** under the §11.4.74 catalogue-check — reuse before reimplement. A consuming project that needs to talk to a model backend or rank documents against a query uses these and does not grow its own seam.
- `submodules/verdict` and `submodules/passage` are **newly created PUBLIC reusables**, extracted from `workshop/` rather than left embedded in it.

All four are **staged, not yet committed** at this writing, which is why `scripts/verify-manifest-pins.sh` compares against the INDEX (`git ls-files -s`) and not against HEAD: their refs are checkable today instead of sitting unverifiable until the adding commit lands. The verifier prints a NOTE naming each one as newly STAGED.

Nothing here is a hardcoded roster and nothing should become one. The fleet is DERIVED at run time from `.gitmodules` plus `helix-deps.yaml`, and the two are guarded against each other in both directions by cascade check C6, with C9 guarding every recorded `deps[].ref` against its live gitlink. Re-derive rather than reading a list:

```
git config -f .gitmodules --get-regexp 'submodule\.*\.path'
bash scripts/verify-governance-cascade.sh      # C1 classifies the
fleet from evidence
bash scripts/verify-manifest-pins.sh           # C9 standalone: refs
vs gitlinks
```

design-toolkit’s two mirrors are NOT in sync, and only one of them is public. As of 2026-09-03 BOTH HALVES ARE MEASURED FROM THIS TREE — the “cannot be probed” claim this section carried is WITHDRAWN as FALSE. Its GitHub origin `vasic-digital/design-toolkit` was flipped **public** on 2026-09-01 after a clean full-history privacy audit.

What changed is the checkout, not the mirror. A gitlab remote **IS** declared for `design-toolkit` today — added by that submodule’s own commit `7d9240e` `fix($11.4.6)`: record the GitLab mirror; absence needs its own evidence, which arrived with the gitlink. Two sentences that stood here are therefore withdrawn as false, not merely stale: *“this checkout wires up **no** GitLab remote”* and *“neither the visibility nor the lag can be re-measured here”*. **Nothing was added to this repository’s configuration to obtain the figures below** — the remote was already there and every probe is read-only.

Measured **2026-09-08**, and every figure this section carried before is now SUPERSEDED — **the mirror is SYNCED and the lag is ZERO:**

	GitHub origin	GitLab gitlab
visibility	public	private
HEAD	e721b3c7eba8	e721b3c7eba8
lag	—	0 / 0

```
git -C design-toolkit remote -v                                # github, gitlab,
origin (BOTH hosts), upstream
git -C design-toolkit rev-parse --short=12 HEAD # e721b3c7eba8
```



```
git -C design-toolkit ls-remote gitlab HEAD      # e721b3c7eba8
git -C design-toolkit ls-remote origin HEAD      # e721b3c7eba8
```

Two claims this section used to carry are WITHDRAWN AS FALSE, and both were true when written:

1. *“The lag is SEVEN commits”* (and the 5- and 6-commit readings before it). The mirror was synced on 2026-09-08 on an explicit operator decision, by a fast-forward push — 520c436c..e721b3c7, nothing forced. All three HEADs now name the same commit.
2. *“The claim that a gitlab remote IS declared is WITHDRAWN AS FALSE.”* That withdrawal is itself withdrawn. A gitlab remote **IS** declared today, and origin is configured to push to **both** hosts. What made the earlier statement true was that upstreams/gitlab.sh was inert — named .disabled, so no upstreams/*.sh glob picked it up. **The operator authorised enabling it**, the recipe was renamed, its origin-matching guard ran green (8 pass / 0 fail) BEFORE the first push, and the mirror was synced.

Direction was verified before acting, and it is the reason this was safe: the mirror is PRIVATE and the GitHub origin is PUBLIC, so every commit GitLab lacked was already published. A public-to-private flow carries no disclosure risk; the reverse would have.

The standing lesson survives all of it: do not quote a lag figure from this file. It has now read 5, 6, 7 and 0 across four measurements. Re-run the rev-list above — it needs no remote, because both commits are in this checkout’s object store.

TWO LINES OF THE BLOCK BELOW WERE WRONG AND ARE CORRECTED, re-measured 2026-09-06. The claim that “a gitlab remote IS declared for design-toolkit” is WITHDRAWN AS FALSE.

git -C design-toolkit remote -v returns exactly three names — github, origin and upstream — and **all three point at git@github.com:vasic-digital/design-toolkit.git**; git config --get remote.gitlab.url returns nothing. The earlier text credited that submodule’s commit 7d9240e with adding the remote. That commit is real and is titled *“record the GitLab mirror; absence needs its own evidence”*, but what it added is the inert RECIPE upstreams/gitlab.sh.disabled — note the suffix, so no upstreams/*.sh glob picks it up — not a git remote. **A recipe recording a mirror is not a remote, and this file conflated the two.**

The rest of the section SURVIVES, which is why only two lines change. The 6-commit lag is still re-derivable here WITHOUT any remote, because both commits are already in this checkout’s object

store: rev-list --left-right --count returns 0 / 6 today. glab is on PATH and its project-path probe needs no remote either. What is NOT re-derivable from this checkout is git ls-remote gitlab HEAD — there is no such remote to ask.

SUPERSEDED 2026-09-08. The block that stood here recorded a checkout with no gitlab remote and a 7-commit lag. Both are gone: the recipe was enabled on an operator decision and the mirror was synced. Re-derive rather than trusting any of it:

```
git -C design-toolkit remote -v
      # github, gitlab, origin (BOTH hosts), upstream
git -C design-toolkit ls-remote origin HEAD
      e721b3c7eba8...
git -C design-toolkit ls-remote gitlab HEAD
      e721b3c7eba8... - the remote EXISTS now
git -C design-toolkit rev-list --left-right --count
      e721b3c7...HEAD      # 0 0 - synced
ls design-toolkit/upstreams/
      gitlab.sh - no longer .disabled
glab api projects/vasic-digital%2Fdesign-toolkit
      # .visibility -> private; needs no remote
```

Honest boundary (§11.4.6), and it is narrower than it was: scripts/verify-submodule-remote-sync.sh probes the declared **origin** only, so its CURRENT verdict for this submodule is a true statement about GitHub and still says **nothing** about the mirror — the gate did not gain mirror awareness, the operator gained a way to ask by hand. upstreams/gitlab.sh in that submodule is **.disabled on purpose**, so no tooling pushes to the mirror and the gap does not close by itself; renaming it is an operator decision that starts publishing to a private mirror. **Do not write as though the two mirrors match, and do not quote “5 commits” as a live figure — re-run the block above.**

submodules/containers (vasic-digital/containers) was added under §11.4.76, which mandates that module for ANY containerised workload and forbids reimplementing it.

The sentence that used to finish that paragraph — “a hand-rolled containerfile is a violation, not merely an inferior choice” — is WITHDRAWN as of 2026-09-04. It overstated the anchor, and an overstated rule is as unhelpful as a missed one. §11.4.76 was re-read verbatim. Its list of what the module is authoritative for is “runtime auto-detection, endpoint discovery, lifecycle/health management, compose orchestration, cross-build, emulator integration, and on-demand service boot”; clause 4 forbids adding a missing “runtime ... or lifecycle primitive” “as a parallel implementation inside the consuming project”; and the closing line names the offence as “reinventing

compose orchestration in-project". A **container image recipe appears in none of those**. Nor could it: the module ships *.Containerfile itself under pkg/crossbuild, and its own planned gate CM-CONTAINERS-USED scans a Dockerfile*-touching change **for an import of digital.vasic.containers/...** — it treats such a file as a trigger to look for the import, never as the violation. Read it precisely: **what is forbidden is the hand-rolled runtime, lifecycle and orchestration logic AROUND a container file, not the file**.

Two claims this block carried are WITHDRAWN as of 2026-09-03, and both were false in the direction that flattered the tree.

- *"The submodule is present for the workload specs/001-... plans."* **It has a real consumer now.** `_tools/containers/` was committed 2026-09-03 (460266c) with 7 tracked files — `go.mod` requiring `digital.vasic.containers`, `go.sum`, `cmd/site-build/main.go`, `compose/compose.sites.yml`, `compose/jekyll-build.sh` — and its own README records the measurement that made it necessary: before it existed, `grep -rn 'digital.vasic.containers'` outside `submodules/` and `workshop/` returned **zero** hits. **The gitlink was declared, manifest-pinned, cascade-verified and unused** — every gate green over a submodule nothing consumed.
- *"Nothing in this repository has yet been built or run in a container."* **False.** A fleet container is running on this host right now: `podman ps` shows `workshop-curriculum_platform_1` (image `docker.io/library/alpine:3.20`) **Up 2 hours, healthy**, serving `/opt/workshop/bin/workshop-server` — the private workshop submodule's platform workload, not the umbrella root's.

And a third fact this block never carried at all: the umbrella root has been shipping container files and their hand-rolled runtime logic since 2026-06-26. `_tools/helixtranslate-container/Containerfile`, `Containerfile.translator` and `run.sh` are **tracked**, first committed in 32dfdbd — three months before `submodules/containers` was adopted. `grep` finds **no** reference to `submodules/containers` or `digital.vasic.containers` anywhere under that directory. It is recorded here rather than left for the next reader to discover.

A second clause of that paragraph is WITHDRAWN, by name. It read *"referenced by 10 tracked files including `_tools/distribute-helixtranslate.sh` and four `_tools/gen/` translation scripts"*, and it silently merged two different paths. Re-measured 2026-09-04:

- **10 tracked files reference the DIRECTORY `_tools/helixtranslate-container/`** — the four carriers, `CONTINUATION.md`,

`_analysis/CONTAINER-DISTRIBUTION.md`, `docs/environment-adaptability/AUDIT.md`, `docs/workshop-curriculum/RECON.md`, `scripts/audit-environment-assumptions.sh`, and `_tools/distribute-helixtranslate.sh` (whose reference is composed at line 65 as `ASSETS="$HERE/helixtranslate-container"`, which a literal-path grep misses). The 10 is right; nine of the ten are documentation or an audit rule.

- **No `_tools/gen/` file is among them.** Those scripts reference `_tools/helixtranslate-container.sh` — a **different file**, the SSH shim — and there are **six** of them, not four: `translate-ui.py`, `translate_ui_all.py`, `translate_ui_batch.py`, `translate_ui_chunked.py`, `translate_ui_slow.py`, `repair_ui_terms.py`.

The surface was then judged file by file, and it splits three ways rather than two. Full table, with the measurement behind each verdict: `_tools/containers/README.md`. In brief — the two Containerfiles are **not** violations (see the withdrawal above); `_tools/helixtranslate-local.sh` **was** one and is **CONVERTED**, its frozen podman literal replaced by the module's own `runtime.AutoDetect` reached through `_tools/containers/cmd/runtime-probe`, verified on this host at all three exit codes; `_tools/distribute-helixtranslate.sh` is one whose replacement exists but is **UNVERIFIED**; `_tools/helixtranslate-container.sh` and `_tools/helixtranslate-container/run.sh` are violations that carry declared exceptions — **and the words “cannot be closed from this tree” that this sentence applied to BOTH of them are WITHDRAWN for `run.sh`, by the measurement below; they were true when written and still hold for `_tools/helixtranslate-container.sh` — and `_tools/translate-fleet.sh` is **not** a violation.**

That last verdict was reversed by measurement mid-task and the reversal is the point. Its `host == amber.local ? docker : podman` looked like hand-rolled runtime detection. It is not: the module's own `pkg/remote.RemoteHost` carries a declared `Runtime` string field — “*the container runtime on this host*” — and ships **no** remote runtime detector to defer to

(`grep -rn 'DetectRuntime\|AutoDetect' pkg/remote pkg/discovery` outside tests matches nothing). Naming a REMOTE host's runtime is configuration the module expects a consumer to supply. Naming the LOCAL host's runtime is not, which is exactly why the same-looking literal in `helixtranslate-local.sh` was a real violation and was converted.

THE CLAIM THIS BLOCK CARRIED IS WITHDRAWN BY NAME, AND IT IS FALSE TODAY RATHER THAN MERELY STALE. It was TRUE when written, at the pin it named. The withdrawn sentence read, verbatim:

*“Measured at gitlink d940b51fc247c285c805799452992da8d09c75b9: pkg/runtime’s ContainerRuntime interface declares Name, Version, IsAvailable, Start, Stop, Remove, Status, List, Stats, Exec and Logs — **no Run, no Create**, so there is no ephemeral-run primitive at all; Exec(ctx, id, cmd []string) accepts no stdin; and the module’s only WithStdin (pkg/remote/connection/interface.go:146) sits in an interfaces-and-options-only package that nothing implements and no constructor returns.”*

The gitlink moved and the primitive EXISTS. Re-measured **2026-09-09** at gitlink 7f5922563d8bec866b1a25eac483590c9a212817 — git ls-files -s submodules/containers and git -C submodules/containers rev-parse HEAD agree. The ephemeral-run primitive landed upstream in 6d13ad03528c (2026-09-04), **two commits before the current pin**, so this repository consumed the fix without noticing it and kept asserting the gap for five days. pkg/runtime/run.go and pkg/runtime/run_test.go exist, and pkg/runtime/runtime.go now declares on the interface itself:

```
Run(  
    ctx context.Context, image string, cmd []string,  
    opts ...RunOption,  
) (*ExecResult, error)
```

WithRunStdin(io.Reader) supplies the document **and** is what puts -i on the argv (pkg/runtime/run.go:119, and args = append(args, "-i") at line 204); StdinExecutor / ExecuteWithStdin is really implemented, by defaultExecutor on the os/exec path (run.go:53). Missing stdin support is refused with ErrStdinUnsupported **before the command runs**, never degraded to an empty-stdin success.

WHAT REMAINS TRUE, and must not be over-corrected away. Two halves of the old claim survive verbatim at the current pin, and they are the reason this is a SPLIT rather than a clean withdrawal:

- Exec(ctx, id, cmd []string) **still accepts no stdin.** Read the interface; the signature is unchanged.
- **The REMOTE-stdin gap is OPEN.** remote.RemoteRuntime.Run (pkg/remote/runtime.go, around line 240) *explicitly refuses* a caller that passes WithRunStdin, and says why in its own words:
“RemoteExecutor exposes Execute ... and ExecuteStream ...; neither

accepts an io.Reader, so there is no seam through which a local document could be streamed to the remote container's standard input." It returns an error wrapping `runtime.ErrStdinUnsupported` rather than a zero-exit result from an empty stdin — a refusal worth recording as good practice, and the same discipline the LOCAL path shows. `pkg/remote/connection` is **still** interfaces and option builders only: `WithStdin` at `interface.go:146` stands, the package holds four files and none of them implements the `Connection` interface.

So the consequence for the two exception-carrying scripts is a SPLIT, not one verdict:

- **`_tools/helixtranslate-container/run.sh` — its declared §11.4.76 exception REASON no longer holds.** That file is installed on, and runs ON, the remote host, so from that host's own view its `run --rm -i` is a LOCAL run. It is **convertible in principle** today via `runtime.Run` with `WithRunStdin / WithRunVolumes / WithRunEntrypoint / WithRunExtraArgs`. Its header has been corrected to withdraw the reason while stating that the conversion remains unperformed.
- **`_tools/helixtranslate-container.sh` — its exception STANDS,** with its reason narrowed by measurement to the remote-stdin gap named above. It is the local half, a raw `ssh ... < "$IN"`, and is **still not convertible**.

NEITHER SCRIPT WAS CONVERTED, and that is a decision with evidence rather than an omission. Re-measured 2026-09-09 on this host: `podman images | grep -i helixtranslate` matches **zero** rows, `getent hosts` fails to resolve both `thinker.local` and `amber.local`, and `ssh -o BatchMode=yes` returns **rc 255** for both. Absent image, unreachable hosts — that is an **rc 2**, and **a 2 is never a pass**. A working script must not be replaced by a rewrite that cannot be exercised end to end; a convertible-in-principle finding is not a verified-in-practice conversion, and this file will not blur the two.

One honest upstream observation, recorded as an observation and not as a defect (§11.4.6). `defaultExecutor.ExecuteWithStdin` builds its child with `exec.CommandContext`, then folds an `*exec.ExitError` into `exitCode` and sets `err = nil` (`pkg/runtime/run.go:53-73`). A CANCELLED run is killed by the context and therefore comes back as a signal death — `ExitCode = -1` with a **nil error** — so a caller must consult `ctx.Err()` to tell cancellation from a container that exited `-1`. **Cancellation**

genuinely works; whether reporting it this way is deliberate is **UNCONFIRMED** — it was not asked upstream and no comment in the module states it.

Honest boundary (§11.4.6): whether those particular images were ever built or run CANNOT BE DETERMINED from this checkout, and the conversion did not change that — and it is UNCHANGED by the primitive landing upstream. Re-measured 2026-09-04 and **re-confirmed identically 2026-09-09**: podman images on this host matches helixtranslate **zero** times, and both remote hosts are **unreachable** — getent hosts fails to resolve thinker.local and amber.local, and ssh -o BatchMode=yes returns **rc 255** for both, which is a measured absence and not a broken resolver. So the converted helixtranslate-local.sh is verified for runtime resolution and exit codes only, **never for an actual translation**, and it says so in its own header. Absent image, absent evidence: this is a 2, and a 2 is never a pass. Re-derive rather than trusting any of the above:

```
git ls-files | grep -iE 'containerfile|dockerfile|compose'
# the tracked surface
git ls-files _tools/containers/
# the consumer
git grep -l '_tools/helixtranslate-container/' -- . # 9
literal + 1 composed = 10
podman ps --format '{{.Names}} {{.Image}} {{.Status}}' #
what is actually running
podman images | grep -i helixtranslate #
expect no rows on this host
( cd _tools/containers && go build -o bin/runtime-probe ./cmd/
runtime-probe \
&& ./bin/runtime-probe )
# the module's own answer
for h in thinker.local amber.local; do getent hosts "$h" || echo
"$h unresolvable"; done
```

Submodule-vs-remote drift — a gate exists now, and it is RED

scripts/verify-submodule-remote-sync.sh compares every declared gitlink to its remote (git ls-remote on the declared branch, else HEAD). **No other instrument in this tree makes that comparison, and none did until 2026-09-01.** C9 and scripts/verify-manifest-pins.sh compare helix-deps.yaml to the **local gitlink**; a submodule arbitrarily far behind upstream is invisible to both, and both were green throughout the period this gate now reports as drift. A green manifest check has never been evidence about a remote. The new gate is three-

valued — 0 all current, 1 real drift, 2 could not determine — and is registered as submodule-remote-sync in scripts/check-registry.tsv with a --prove-failure paired proof, which is what R5 required of it.

It has now exited 0 twice and 1 three times, and the SAME submodule caused every red. The “it exits 0 as of 2026-09-02, the first time it has ever done so” that stood here is **WITHDRAWN as current**; so is the “back to 1 on 2026-09-03” that replaced it. Five readings, all real, kept because the movement is the record — and because readings 3→4→5 are the clearest evidence in this file that a green here does not survive the day:

```
bash scripts/verify-submodule-remote-sync.sh
# 1 - 6 CURRENT, 6 DRIFT, 0 UNDETERMINED (2026-09-01)
# 1 - 11 CURRENT, 1 DRIFT, 0 UNDETERMINED (2026-09-02, after
      five bumps landed)
# 0 - 12 CURRENT, 0 DRIFT, 0 UNDETERMINED (2026-09-02, after the
      authorized fast-forward)
# 1 - 11 CURRENT, 1 DRIFT, 0 UNDETERMINED (2026-09-03, the SAME
      pin went stale again)
# 0 - 12 CURRENT, 0 DRIFT, 0 UNDETERMINED (2026-09-03, after the
      FOURTH authorized fast-forward)
```

The 2026-09-03 DRIFT row, verbatim, before it was closed: submodules/constitution 3be10826f3d2 2887b42e9349 DIFFERS. The gate’s own words at the time — “*The difference is DETERMINED; the direction is NOT, because the remote commit is not in this checkout’s object store.*” **The operator then authorized the --fetch, which resolved it as 1 behind / 0 divergent, and git merge --ff-only closed it.** That is the designed sequence: the gate states what it cannot know, an operator decides, and only then does a pin move.

The 2026-09-01 run named design-toolkit, ai_interviewing, submodules/containers, submodules/LLMProvider and submodules/RAG as **BEHIND** and fast-forwardable, and submodules/constitution as **DIFFERS** against remote b9096acd98d2. All five BEHIND rows were bumped and matched their remotes. The last row was submodules/constitution, reading f5876a3b700e vs 3be10826f3d2 **DIFFERS** with the direction UNDETERMINED because the remote commit was not in this checkout’s object store; the operator authorized the --fetch, which classified it as **2 behind, 0 divergent**, and git merge --ff-only closed it. submodules/superspec is probed and reported as a third-party NOTE, never a verdict input, which is why 12 owned gitlinks are probed of 14 declared.

Do not bank the green — and note that this file already said so, and was right. This pin has now gone stale within a day on **three** consecutive occasions, the third measured 2026-09-03 against a green

recorded 2026-09-02. The correct standing conclusion is that the constitution pin is an operator decision that RECURS, not a task that completes.

Note what this gate is blind to even when a row reads CURRENT. It probes the declared **origin** only. design-toolkit has a GitLab mirror, and the parenthetical that stood here — *“this checkout declares no remote for (git -C design-toolkit remote -v → origin, GitHub, only)”* — is **WITHDRAWN as false**: a gitlab remote IS declared there today. What survives is the real limitation: the gate does not probe it. Re-measured 2026-09-03, this row reads design-toolkit e721b3c7eba8 e721b3c7eba8 CURRENT. **The “6 commits behind at 520c436c2c2a” reading is SUPERSEDED** — the mirror was synced 2026-09-08 and now sits at the same commit. But the limitation the sentence illustrated is UNCHANGED and is the point: the gate reads CURRENT because it asked GitHub, and it would read CURRENT just the same if the mirror were a thousand commits behind. **A gate reporting CURRENT is not evidence about a remote it never asked** — and that is true today, when the mirror happens to agree, exactly as it was when the mirror did not. See “Owned submodules” above.

Bumping a gitlink is an operator decision and this gate does not make one. **Do not silence the 1** by deleting the gate, by allow-listing a submodule, or by bumping pins to make it green.

Build and test entry points

The gates mirror the definitions preserved in .github/workflows/ci.yml.disabled and are run from the repository root. Toolchains: Go 1.26, Node 20, Ruby 3.3 + Bundler, poppler-utils, tesseract-ocr.

There is no server-side CI at this umbrella root. As of 2026-08-27 the workflow is disabled per §11.4.156(B) and enforcement is a **local pre-push hook**. .git/hooks/ is not tracked by git, so a fresh clone runs no gates until bash scripts/pre-push-gates.sh --install is run, and git push --no-verify bypasses the hook with no record — run these by hand until you have confirmed the hook is in place. (This says nothing about the site submodules: milosvasic.ru deliberately keeps an active deploy workflow — see **Deploys**.)

```
cd _tools/gen && go test ./... && cd -           # Go unit tests
              (generator)
bash _tools/audit-hardcoding.sh                   # hardcoding audit
bash _tools/translate/reproducibility-selftest.sh
```

```
bash _tools/portfolio/self-validate.sh
bash _tests/run-harness-selfvalidation.sh      # harness self-
validation
```

Playwright (chromium) additionally requires `npm ci` and `npx playwright install chromium` inside `_tests/`, and a built `milosvasic.ru/_site` before the suite is served.

TWO OF THOSE TOOLCHAINS ARE ABSENT ON THIS HOST, and both make a gate report rc 2 — measured 2026-09-06, and rc 2 is never a pass.

- **tesseract is NOT INSTALLED**, so **gate 5 exits 2**. poppler-utils IS present — `pdftotext`, `pdfinfo` and `pdftoppm` all resolve — so the export validator runs 5 of its 6 checks and skips exactly one:
[SKIP] FULL-VISUAL/visual.ocr — tesseract missing — cannot OCR rendered pages, giving 5 PASS / 0 FAIL / 1 SKIP of 6 and `verdict=UNDETERMINED`. **This is the gate working, not failing**: its golden-bad arm still DETECTS correctly (rc 1, naming two CONTENT faults), so the detector is live; it simply refuses to call the golden-good arm a pass while a check family did not run. Remedy is a package install and therefore an operator action: `sudo apt-get install -y tesseract-ocr`.
- **bundle, bundler and jekyll are all absent**, so gate 6's precondition reports CANNOT DETERMINE — see the `_site` section below.

Re-derive rather than trusting either line; a host changes under you:

```
for b in pdftotext pdfinfo pdftoppm tesseract bundle jekyll; do
  printf '%-10s %s\n' "$b" "$(command -v "$b" || echo MISSING)";
done
bash _tests/run-harness-selfvalidation.sh
# 0 pass · 1 finding · 2 could-not-determine
```

Governance and adaptability instruments

Separate from the pre-push gates. Every one is three-valued — **0 clean, 1 a real finding, 2 COULD NOT DETERMINE, and 2 is never a pass. The whole suite was re-run on 2026-09-03 and the results below are that run.** Four verdicts moved since 2026-09-02 and every superseded figure is named where it stood, not silently replaced. They remain dated observations, not standing facts.

```
bash scripts/continuation-check.sh      # 0 — 8 PASS/0
DRIFT/0 UNDET/7 NOTE
bash scripts/verify-governance-cascade.sh  # 0 — 12 PASS/0
FAIL/0 ENV/8 NOTE (C0..C9)
```

```

bash scripts/verify-manifest-pins.sh          # 0 – C9
standalone: 12 MATCH/0 DRIFT/0 UNDET of 12
bash scripts/verify-check-registry.sh         # 0 – 57 PASS/0
FAIL/1 DEBT (2026-09-06; 43/45 superseded)
bash scripts/audit-hardcoded-paths.sh         # 0 – BACK TO
GREEN by a REAL FIX, not a re-baseline; see below
bash scripts/audit-environment-assumptions.sh # 0 – 567 allow-
listed, 666 baselined, 2247 files
bash scripts/verify-content-boundary.sh        # 1 – RED BY
DESIGN, see "Content boundary"
bash scripts/verify-submodule-remote-sync.sh  # 0 – RED, THEN
GREEN AGAIN: 12 CURRENT/0 DRIFT
bash scripts/verify-provider-ci.sh            # 1 – 1 CONFIRMED/
6 UNVERIFIED/2 HISTORICAL
bash scripts/verify-pretooluse-guard.sh        # 0 – 8 PASS/0
FAIL/0 UNDET (2026-09-09; G12)
bash scripts/verify-all-constitution-rules.sh # 1 – 173 PASS/96
FAIL/2 ERROR of 271 gates
bash scripts/lumen-index-doctor.sh            # semantic index
health
bash scripts/ollama-tune.sh                   # local inference
host tuning

```

Four verdicts that moved on 2026-09-03, stated as withdrawals:

Instrument	Was (2026-09-01/02)	Is (2026-09-03)
verify-check-registry.sh	0 — 41 PASS	0 — 43, then 45 PASS within the same session
audit-hardcoded-paths.sh	0 — 6 file(s) allowed	1 — 1 occurrence, 12 file(s) allowed
audit-environment-assumptions.sh	0 — 531 allow-listed, 2166 files, 683 baselined	0 — 567 allow-listed, 2247 files, 666 baselined
verify-submodule-remote-sync.sh	0 — 12 CURRENT	1 — 11 CURRENT / 1 DRIFT, then 0 — 12 CURRENT after the fourth authorized fast-forward, both on 2026-09-03
verify-content-boundary.sh	1 — 11878 (prose 11356, short 433, name 89)	1 — 15308 (prose 14552, short 642, name 114) with corpus MOVED – 18 of 12535 files; RED BY DESIGN, and the gate now REPORTS its

Instrument	Was (2026-09-01/02)	Is (2026-09-03)
		own instability instead of hiding it

audit-hardcoded-paths.sh went RED and the finding is REAL, not a re-baseline. Exit 1, 1 occurrence(s) across 1 file(s), scanning 6054 files across 14 repositories. The single occurrence is workshop/chapters/01/transcript/accuracy-plan.json:3, a JSON value carrying this developer host's absolute checkout path. **It is INSIDE the private workshop submodule and was NOT edited from here** — a finding inside a submodule is fixed inside that submodule and returns as a gitlink bump, which is an operator decision. Two other figures in the same run also moved and are recorded rather than glossed: files **explicitly allowed** 6 → 12, and **377 baselined occurrences** are printed on every run as declared, known, unfixed debt.

Five of those need reading carefully rather than glancing at:

- **verify-check-registry.sh used to print 5 DEBT rows on every run. As of 2026-09-02 it prints ZERO — all five owed paired proofs were written, and each was verified in BOTH directions: it passes on this tree AND it fails against a deliberately weakened throwaway copy of the gate it guards.** Every mutation is DATA rather than a code edit, so each control is green by construction — the “inoperative proof” defect the registry itself documents. setup-agents-wizard-suite also owed **three-valued** and now demonstrates rc 2 four ways. Re-measured **2026-09-03: 43 PASS / 0 FAIL / 0 DEBT / 0 UNDET / 0 NOTE**, exit 0. The **41 PASS** figure measured 2026-09-02 is SUPERSEDED — two further checks were registered in the intervening day, which is R5 working rather than drift; the PASS count has now moved 23 → 25 → 31 → 41 → 43 for that reason alone. The five debt rows became check rows, and no new unregistered *.sh appeared, so R5 stays green.

--run-proofs now exits 0 — re-measured 2026-09-03 at **65 PASS / 0 FAIL / 0 DEBT / 0 UNDET / 0 NOTE**. **The reading this bullet carried — “still exits 1, at 54 PASS / 5 FAIL” — is WITHDRAWN**, and all five named rows were re-checked individually: provider-ci, submodule-remote-sync, mutation-anchor-rot, private-object-exposure and remedy-executability each now show **3 PASS / 0 FAIL**. The one previously called REAL is genuinely resolved: provider-ci no longer fails “M6b no repository or owner name in body”; its --run-proofs row reports PASS M1 verdict is not constant ... got 5 distinct. The other four were described as false positives of the hollow-proof heuristic (proofs that ran and returned 0 but

summarised in prose rather than the $M_{<n}$ / “N mutations” form) and now report in the recognised form. **The heuristic was NOT loosened** — it is the only thing between the registry and a proof that returns 0 while exercising nothing. **This run takes roughly an hour**; a plain run verifies proof *structure* only and is not a substitute.

Honest boundary (§11.4.6): a 0 here means every registered proof executed and reported, not that this tree is defect-free. The sweep’s **DROP** direction remains open and is printed on every run: its ADD direction is proved, while assertion **L1** deletes a gate, measures the count falling $3 \rightarrow 2$, and records that the sweep still exits 0 because it keeps **no expected-gate ledger**.

The zero-DEBT history is worth keeping. A zero means “every registered check is accounted for”, **not** “every check has a working paired proof” — a plain run verifies proof *structure* only and says so; --run-proofs actually executes them, --strict makes debt block. Its R5 anti-drift rule means **a new *.sh under scripts/ fails the registry until it is registered** in scripts/check-registry.tsv (documented in docs/check-registry.md) as a check, a debt, or an exemption. Measured across 2026-09-01 in three states, all real: exit **0** at 21 PASS / 0 FAIL / 5 DEBT in the morning; then exit **1** at 20 PASS / 1 FAIL / 5 DEBT once an unregistered scripts/verify-content-boundary.sh appeared — R5 doing its job; then exit **0** once that gate was registered. Re-measured later the same day it is exit **0** at **31 PASS / 0 FAIL / 5 DEBT / 0 UNDET / 0 NOTE** (re-measured 2026-09-02; the **25 PASS** figure carried here earlier is superseded) — the PASS count moved $23 \rightarrow 25 \rightarrow 31$ as further checks landed and were registered, which is R5 working rather than drift, and the “23 PASS” figure this bullet used to carry is superseded, not wrong-at-the-time. **Both the FAIL and the 5 DEBT rows are now cleared** — the five that owed a proof were constitution-rules-sweep, lumen-index-doctor, ollama-tune, prepush-gates and setup-agents-wizard-suite.

- **audit-environment-assumptions.sh went RED at 17 and was taken back to GREEN in the same session — 12 by REAL FIXES, 5 by reasoned exemptions. Read that split before quoting the exit code.** The instrument has now moved six times without the audit itself being edited once. Every state is recorded, because the movement is the point.

Current state, re-measured 2026-09-03: exit 0, no NEW frozen environment assumptions (567 justified occurrence(s) allow-listed), scanning **2247 files across 14 repositories in 12 classes**, printing **666 baselined occurrences** and **1 out-of-scope third-party assumption**.

The 2026-09-02 reading — 531 allow-listed, 2166 files, 683 baselined — is SUPERSEDED, and all three numbers moved without the audit being touched. `git status --short scripts/audit-environment-assumptions.sh` is EMPTY, so the instrument is byte-identical to the one that produced the old figures; the scanned fleet grew (2166 → 2247 files) because the submodules under it did. Baselined debt fell 683 → **666** (−17) and the allow-list grew 531 → **567**; **this file does NOT claim which of those is a fix and which is a new exemption — that was not measured this session.** Do not read the −17 as seventeen defects repaired.

How the 17 were cleared — the distinction is the whole point, because an allow row is not a fix:

Resolution	Count	What
Real fix	11	pipeline/benchmark/{build_bench_nomic_full,run_retrieval_benchmark}.py — every <code>argparse</code> default now <code>os.environ.get(...)</code> (<code>\$OLLAMA_URL</code> , <code>\$BENCH_EMBED_MODEL</code> , <code>\$WORKSHOP_BASE_URL</code> , <code>\$BENCH_ALT_MODEL</code> , <code>\$BENCH_CURRENT_MODEL</code>), each keeping its former literal as last resort so existing callers are unchanged; prints and docstrings now report the RESOLVED model instead of a frozen name
Real fix	1	platform/backend/gates/prove-evidence-precision-mutation.sh — bare GNU-only <code>sed -i</code> → portable <code>sed -i.bak ... && rm -f "\$SED_BAK"</code>
Exemption	1	platform/backend/internal/api/suggest_gates_test.go — <code>net.Listen("tcp", "127.0.0.1:0")</code> asks the KERNEL for a port and reads it back with <code>ln.Addr()</code> , injecting it via <code>t.Setenv</code> ; that is an override being exercised, not an assumption, and <code>go build</code> excludes <code>*_test.go</code> from every production binary
Exemption	4	the three <code>pipeline/benchmark/*.json</code> result files — a benchmark result records WHICH MODEL WAS MEASURED; rewriting that name would falsify the record rather than adapt anything. Verified before exempting: only <code>positive_queries</code> / <code>negative_queries</code>

Resolution	Count	What
		are ever dereferenced, the <code>_meta</code> object is never read back

The 5 exemptions are recorded debt-free but they are still exemptions, each carrying a # REASON: naming the evidence. **No # BASELINE: row was added and the 728-row baseline was not touched.** A green exit bought with five reasoned rules is not the same artifact as a green exit bought with twelve fixes; both happened here, and the table above is why the distinction survives.

The red that preceded it, and its correctly-attributed cause. The run that opened this session exited 1 with 17 frozen environment assumption(s). All 17 were inside workshop/, in seven files: pipeline/benchmark/build_bench_nomic_full.py (6), pipeline/benchmark/run_retrieval_benchmark.py (5), pipeline/benchmark/retrieval_benchmark.json (2), pipeline/benchmark/results.json (1), pipeline/benchmark/results_nomic_full.json (1), platform/backend/internal/api/suggest_gates_test.go (1) and platform/backend/gates/prove-evidence-precision-mutation.sh (1).

Attribute this carefully — a stronger claim was drafted and WITHDRAWN before it stood. It read “the workshop gitlink bumps pulled the benchmark into scope”, resting on all seven files being ABSENT at 55076bf943a5. They are — but 55076bf943a5 is the gitlink CONTINUATION.md still *recorded*, not the one in effect at the previous Synced-Commit, which was 6130d6b55d693 (git rev-parse 1b90daa6:workshop). Against that correct baseline, only 2 of the 7 files are new (build_bench_nomic_full.py, results_nomic_full.json = **7 findings**); the other 5 files / **10 findings** were already present and simply went unseen. **Comparing against a stale reference point does not just misstate a number — it manufactures causality.** The fleet-derivation mechanism is real for the 7; the 10 are the more uncomfortable half.

Most were frozen model names (nomic-embed-text, jina-embeddings-code-cpu) and endpoints (127.0.0.1:11434, 127.0.0.1:8087) in benchmark code and its committed JSON results. **They were judged separately rather than blanket-allowed by directory** — the argparse defaults were the real freezes and were fixed; a results file recording which model was measured is closer to a RESULT than an assumption and was exempted with its evidence. Same directory, opposite verdicts, which is the outcome a blanket rule would have destroyed.

The earlier cycle, kept because each dead figure is a lesson: 7 frozen GNU-vs-BSD assumptions (six in scripts/verify-check-registry.sh — GNU-only in-place sed -i, stat -c '%a' — and one in scripts/verify-manifest-pins.sh) were fixed with portable helpers, taking the audit **green (exit 0)** at 1763 files across 14 repositories; then, re-measured 2026-09-01, it went **red (exit 1)** at **10 frozen environment assumptions**, scanning **1794 files across 14 repositories in 12 classes** — all 10 inside submodules/LLMProvider (pkg/providers/ai21/ai21_test.go, pkg/providers/openrouter/openrouter_test.go, scripts/prove-offline-discovery.sh), which had just joined the owned fleet and brought its own occurrences into scope with no edit to the audit itself. **That figure is now WITHDRAWN, not restated.** An intervening run on 2026-09-02 exited **0** at **1987 files across 14 repositories in 12 classes**, and none of the three previously-cited submodules/LLMProvider occurrences was flagged. **Careful with that sentence: the 10 LLMProvider FINDINGS are gone, but that submodule still contributes 14 files to the 728-row BASELINE** (13 under pkg/providers/, plus challenges/scripts/host_no_auto_suspend_challenge.sh), re-verified 2026-09-02. A finding fixed upstream and a baselined debt row are different states printed in different sections — an empty findings list is not a clean submodule. That submodule's own commit log (git -C submodules/LLMProvider log --oneline) shows the fix landed upstream: test(ai21): stop asserting AI21's uptime in a unit test, test(settings,codestral,ollama): stop freezing real model ids and endpoints in fixtures, and fix(providers): default model and endpoint were frozen with no override layer (F23, F24) — a real fix in the consumed submodule, not a re-baselining in this tree. Read the rest of today's output before calling anything clean. **Both long-standing residues were worked this session and BOTH MOVED, by real fixes rather than by re-baselining:**

- **STALE allow rules 2 → 0** (rule count 444 → 436). Both were DELETED from the embedded ALLOW_RULES heredoc, each leaving a tombstone comment carrying the evidence. scripts/audit-hardcoded-paths.sh * * was a blanket whole-file exemption that suppressed NOTHING — its one class hit sits inside a # comment, which the scanner blanks for .sh before any class is tested. _tools/gen/review_ui_all.py MODEL * named a model id living only in a module docstring with no quote character on the line, and git log -p --all shows that was the

- only occurrence the file has EVER had. Neither deletion surfaced a new finding, and --strict-allow-list moved **1 → 0**.
- **Baselined occurrences 728 → 683** (-45), across 216 files. **F13 and F14 are CLOSED**. The slice was chosen by measurement, not taste: of the 728, **638 sit inside submodules** (472 containers, 102 constitution, 20 workshop, 17 LLMPProvider, 14 ai_interviewing) and can only be fixed by an upstream commit returning as a gitlink bump, leaving **90** umbrella-root-owned — of which `_tests/` ENDPOINT was exactly half. The fix is a new `_tests/env.js`, one source of truth for PORTS **and** BASES, each process.env-derived with the former literal as documented default and an unparseable port throwing at `require()` instead of silently defaulting. Covering both closes the half no allow rule could express: the port a config **bound** and the base a spec **requested** were independent literals free to disagree, and a disagreement surfaced as an assertion about the *site*. Measured: `node --check` on 29 files; **237 passed** on chromium at default ports; **12 passed / 41 passed at non-default VD_PORT=9401 MV_PORT=9082**, impossible before and the entire content of F13. Proof the SOURCE moved rather than the ledger: `git show HEAD:scripts/audit-environment-assumptions.sh` run against the fixed tree reports **8** stale rules. `_tests/env.js` **is now TRACKED** — added in commit 402a8c7, confirmed 2026-09-03 with `git ls-files --error-unmatch _tests/env.js`. The note that stood here ("`_tests/env.js` is **UNTRACKED**, so `git ls-files` does not yet show it to the gate") is **WITHDRAWN**: every gate that enumerates via `git ls-files` now sees it.

Still printed on every run by design, re-measured 2026-09-03: **666 baselined occurrences** — declared, known, unfixed defects; the **683** figure is superseded — and **1 frozen assumption inside a third-party gitlink** (submodules/superspec/.github/workflows/ci.yml, a pinned python-version: "3.12"), reported out-of-scope under §11.4.156(C) / §11.4.29 so it is never silently omitted. **A baseline is recorded debt, not a justification, and an allow row is not a fix — and ANY exit code here, green or red, is a measurement of today, not a guarantee about tomorrow's fleet. This instrument's verdict has changed five times without the audit itself being edited once; the fleet moves under it. Re-run it, never quote it.**

- **verify-content-boundary.sh exits 1 and is MEANT to**. Re-measured 2026-09-04: **15308** (prose 14552, short 642, name 114) with **19 undetermined rows**, of which the load-bearing one is corpus MOVED — 18 of 12535 enumerated file(s) changed. The

earlier claim that **“the total is now known not to be reproducible run-to-run”** is **SUPERSEDED**: the algorithm is deterministic — four runs on a frozen snapshot were byte-identical — and the movement is concurrent editing, which the gate now measures and names rather than absorbing silently. See the “Content boundary” section above. A 1 from this gate is a reading assignment, not a regression — but the *reading* has not kept pace with the *population*, and that gap is still the finding. **What DID move on 2026-09-03 is the reading assignment’s SHAPE**: class A, 61.5% of the population, now has text-level direction evidence over its complete row set — 79.7% outward, ~19.5% inward — so the rows an operator must actually open number about 1540, not about 7870. Nothing was judged, allow-listed or re-baselined to achieve that.

- **verify-submodule-remote-sync.sh exits 0 as of 2026-09-03, and that is a measurement of today rather than a property of the tree.** Re-measured after the fourth authorized fast-forward: **12 CURRENT / 0 DRIFT / 0 UNDETERMINED**. Both **“six owned gitlinks are out of sync”** (2026-09-01) and **“11 CURRENT / 1 DRIFT”** (2026-09-03 morning) are WITHDRAWN as current. **This gate has now gone red on the constitution pin four times; treat a green as perishable.** Bumping is an operator decision. See “Submodule-vs-remote drift” above. It is the ONLY instrument on this list that looks at a remote at all, and even it probes origin only.
- **audit-hardcoded-paths.sh exits 1 as of 2026-09-03 and the 1 is REAL.** One occurrence, in workshop/chapters/01/transcript/accuracy-plan.json — a developer-host absolute path frozen into a tracked JSON value inside a private submodule. Not fixable from this tree; see the block above the bullet list.
- **verify-provider-ci.sh exits 1 and the 1 is CONFIRMED, not undetermined.** One standing provider-side trigger, on vasic-digital/vasic-digital.github.io. Six rows are separately UNVERIFIED and are printed regardless of exit code, by the script’s own documented rule: *“Precedence: CONFIRMED (1) outranks UNDETERMINED (2) outranks clean (0)”* — so **the 1 hides no 2 here, and the 6 UNVERIFIED rows are still open.** See G4 below.

Both audits, and gate E, once had the same blind-instrument defect and all three are fixed: gate E enforced §11.4.156 only at the umbrella root, because `git ls-files` sees a gitlink as ONE entry — an active `pages.yml` in `milosvasic.ru` and an active `ci.yml` in `superspec` both passed it — and the two audits scanned **zero files inside the 9 submodules declared at the time**. All three now derive the fleet from `.gitmodules +`

helix-deps.yaml, which is why the environment audit's coverage moved with the fleet on its own — 14 repositories at the 2026-09-01 run, with no edit to the audit. scripts/pre-push-gates.sh's run_gate likewise gained a distinct **UNDET** verdict for a child rc=2 instead of mapping every non-zero return into FAILED, and gate 6 classifies network-class failures as rc=2. **A blind instrument reports PASS.**

Deploys

Deploys are driven by bash _tools/deploy-langs.sh. It regenerates EN plus every complete language into both site submodules, commits and pushes each site only when something changed, then validates the LIVE sites. --dry-run previews without committing.

milosvasic.ru self-publishes on push, and that must not change (verified 2026-08-27, re-verified 2026-09-01: git ls-files '.github/workflows/*' in that submodule returns .github/workflows/pages.yml — one file, still under its active name). Its .github/workflows/pages.yml is the custom GitHub Pages build+deploy action and is **ACTIVE**. It was briefly renamed to a .disabled name under the 2026-08-27 11.4.156 decision (*“Comply — disable both, enforce locally.”*); that half of the decision was **reversed the same day** on the operator's overriding directive:

“Make sure all pages websites work flawlessly! No website can be broken! All websites we have here are running deployed in production!”

The material fact behind the reversal: gh api repos/milos85vasic/milosvasic.ru/pages returns build_type: "workflow" (**re-measured 2026-09-03 via scripts/verify-provider-ci.sh: pages=enabled build_type=workflow source=main:/ status=built — unchanged**) — that workflow is the **sole** publish path for the live site. There is no gh-pages branch and no docs/ folder, and the repository root is Jekyll SOURCE (Liquid + front matter), so it cannot be served raw from a branch. _tools/deploy-langs.sh is **not** a substitute: it generates, commits and pushes source, then sleeps waiting for the server to rebuild — it covers generation and push, none of the publish step. **Do not disable, rename, or otherwise “fix” pages.yml.** vasic.digital needs no build step (committed static HTML), but its Pages source is still build_type: "legacy" (**re-measured 2026-09-03: pages=enabled build_type=legacy source=main:/ status=built, and the probe classifies it CONFIRMED — 41 provider-generated pages build and deployment runs in the last 30 days, newest 2026-09-01T19:54:02Z**), so

every push still triggers a provider-side run even though `git ls-files '.github/workflows/*'` in that submodule returns **zero** files. **This is the fleet's only CONFIRMED standing provider-side trigger.**

design-toolkit adds **no** §11.4.156 CI surface of its own: the same `git ls-files '.github/workflows/*'` returns **0** there too (re-measured 2026-09-03), so its flip to public created no new provider-side trigger — and the 2026-09-03 probe confirms it from the provider side as well, NONE on both its GitHub origin and its GitLab mirror (`jobs_enabled=true`, but `active-pipeline-schedules=0` and `pipelines(listed=0, in-window=0)`).

Every `build_type` and `run-count` in this file is a DATED OBSERVATION, not a standing fact. Provider settings change outside this tree and nothing here can see that happen. Re-measure before relying on any of them:

```
bash scripts/verify-provider-ci.sh  # 0 = none found, 1 =  
                                   confirmed, 2 = could not determine
```

It enumerates the owned repositories from this checkout's own remotes and queries the provider. Exit 2 means UNVERIFIED — it is not a pass and must never be recorded as one. `scripts/setup-agents-wizard.sh` runs it as Step 9 and turns a confirmed finding into a manual step, because changing a provider setting is an operator action in a provider UI.

Honest boundary — this repository is NOT yet fully compliant

Anchor 11.4.6 forbids reporting a state you have not verified, so this section states plainly what is still missing rather than letting the pointer above imply full adoption. The full audit is [docs/constitution-adoption/INVENTORY.md](#).

Known open gaps at the time this carrier was created (identifiers are the inventory's own):

- G3 — CLOSED (verified 2026-08-27). Both halves of the §11.4.32 sweep contract now exist and run. `scripts/verify-all-constitution-rules.sh` discovers its gates dynamically, and `scripts/verify-governance-cascade.sh` supplies step 1: re-measured 2026-09-01 as **12 PASS / 0 FAIL / 0 ENV / 8 NOTE**, exit 0, with a paired-mutation proof (`--prove-failure`) that catches seeded violations as `rc=1` and reports an environment fault as `rc=2` rather

than accusing the tree. It grew from 10 checks to 12 with **C8** (the four carriers INSIDE each owned submodule agree with each other, normalising the per-agent header rather than splitting at a fixed line, because submodule carriers do not share this root's header geometry) and **C9** (every `helix-deps.yaml` `deps[].ref` equals its live gitlink; re-derivable standalone as `bash scripts/verify-manifest-pins.sh`). Step 1 is a measured PASS, not a SKIP-with-reason; OC-3 is resolved. **C5 ROOT-LOCKSTEP was tightened on 2026-09-01 and no longer splits at a hardcoded line 24.** It now measures where the four carriers actually converge and enforces from there, and prints a NOTE saying so — *“declared split 19 equals the measured convergence line 19 — the gate enforces exactly the range it claims, with no ungated identical remainder (ceiling 24)”*. That closes a real blind spot: a line-24 window is a strict SUBSET of the shared region, so the old form could pass on carriers that differed at lines 19–23. Lines 1–18 are the per-agent opening block. **Edit the shared region as ONE artifact** — split the head off, edit the tail once, recompose — rather than editing four files by hand and trusting a gate to catch the slip:

```
for N in 18 19; do
  printf 'from %s: ' "$N"
  for f in CLAUDE.md AGENTS.md QWEN.md GEMINI.md; do tail -n +
    $N "$f" | sha256sum; done \
    | sort -u | wc -l
done # expect 4, then 1
```

This closes the sweep CONTRACT, not the sweep's verdict. For the sweep's own PASS/FAIL split see “The sweep's own result” below. The step-1 caller was ALSO fixed: it previously mapped any non-zero rc to STEP1 FAIL, collapsing the verifier's three-valued contract and reporting a broken check as a governance violation. It now branches on rc and has a distinct ERROR state that exits non-zero without accusing the tree. The earlier “37 PASS / 21 FAIL / 0 ERROR out of 58” figure is WITHDRAWN, not restated: the true pre-fast-forward split was 36/22, and the gate population moved 57 → 286 when the constitution was fast-forwarded, so no old split is comparable to a present-day run. No completed post-fast-forward split is claimed here.

- G4 — PARTIAL (decided 2026-08-27; **partially reversed the same day**). `github/workflows/ci.yaml` is active at the repository root, which conflicts with anchor 11.4.156(A). Resolving it is an operator decision (disable per 11.4.156(B), or record an explicit override in a project Constitution). **The override half of that sentence was wrong: no such option exists.** 11.4.156 refuses the exemption

vocabulary by name (“No escape hatch — no `--allow-ci ... --ci-exempt` flag”), and a consumer carrier may only extend inherited rules, never weaken or override them — so a project-local override is structurally impossible, not merely disfavoured. First operator decision: *“Comply — disable both, enforce locally.”* Both `.github/workflows/ci.yml` and `milosvasic.ru/.github/workflows/pages.yml` are renamed to a non-active `.disabled` name. **Reversed in part on 2026-08-27** after a material fact emerged: `gh api repos/milos85vasic/milosvasic.ru/pages` returns `build_type: "workflow"`, so `pages.yml` is the **SOLE** publish path for the live production site (no `gh-pages` branch, no `docs/` folder, root is Jekyll source; `_tools/deploy-langs.sh` pushes source and waits — it does not publish). The operator’s overriding directive: *“Make sure all pages websites work flawlessly! No website can be broken! All websites we have here are running deployed in production!”* **Current state.** `.github/workflows/ci.yml` → `ci.yml.disabled` and the gates run from a local pre-push hook — **the umbrella root complies** with 11.4.156. `milosvasic.ru/.github/workflows/pages.yml` is **ACTIVE and will stay active**; that submodule is a **known, documented deviation** from 11.4.156, taken deliberately for production uptime. It is **not** an Override §11.4.156 — the rule forbids one — and must never be written up as one. `vasic.digital` is non-compliant at the **provider** level (`build_type: "legacy"`; every push triggers a `pages` build and deployment run) with **no file-level remedy**, because it has zero workflow files. **Not CLOSED** — honest boundary (11.4.6): file-level disabling stops FILE-triggered runs but does not reach provider-side settings (org-default required workflows, branch-protection required checks, the GitHub Pages source setting, provider-side scheduled exports). That boundary is unchanged: those settings are still operator-only to CHANGE, in a provider UI. What is no longer hand-asserted is their STATUS. It is measured on demand by `bash scripts/verify-provider-ci.sh` (0 = none found, 1 = provider-side triggering confirmed, 2 = could not determine — **not** a pass), which the setup wizard runs as Step 9 and which surfaces a confirmed finding as a manual step. Do not quote a status from this document as current; run the check.

RE-MEASURED 2026-09-03 — exit 1, and the verdict per repository is NOT what this entry’s prose implies. The probe covered **22 repositories / 40 upstream rows** over a 30-day window, with `gh` and `glab` both authenticated: **1 CONFIRMED · 6 UNVERIFIED · 2 HISTORICAL · 30 no-trigger · 1 out-of-scope.**

Repository	Reading 2026-09-03	Meaning
vasic-digital/ vasic-digital.github.io	CONFIRMED — pages enabled, build_type=legacy, 41 provider-generated runs in window, newest 2026-09-01T19:54:02Z	The only standing provider-side trigger in the fleet. Confirms this entry's legacy claim, with no file- level remedy: <code>git ls-files '.github/ workflows/*'</code> returns 0.
milos85vasic/ milosvasic.ru	HISTORICAL , <i>not</i> confirmed — pages enabled, build_type=workflow (re-verified), 34 provider-generated runs since 2026-08-04 but newest 2026-08-06T20:29:49Z	<code>pages.yml</code> remains the SOLE publish path and MUST NOT be disabled, renamed or “fixed” . The gate's own words: those runs are <i>“a fact about the past, NOT a claim that a push today triggers one”</i> .
milos85vasic/vasic (this root)	clean — active=0 inert=1, 0 provider- generated runs	The umbrella half of G4 holds.
6 rows on gitflic.ru / gitverse.ru	UNVERIFIED	No read-only API adapter is registered for those hosts. rc-2 material — not a pass.

Two findings worth acting on, neither acted on here. (a) The probe's own remediation output flags **21 repositories where Actions are ENABLED while the tree declares zero active workflow files** — disabling Actions is the only file-independent way to stop provider-generated runs, but it would also stop any Pages *workflow* build, so `milosvasic.ru` must be excluded from any such sweep. (b) The `milosvasic.ru` HISTORICAL rows mean **34 runs occurred that no setting readable today explains**; confirming that the setting which produced them is genuinely gone is operator-only.

Cost of the umbrella half: no server-side enforcement on push or PR; .git/hooks/ is untracked, so a fresh clone is unprotected until bash scripts/pre-push-gates.sh --install is run, and git push --no-verify bypasses the hook. Record: docs/constitution-adoption/DECISION-11-4-156-COMPLY.md; CI surface inventory: docs/constitution-adoption/CI-INVENTORY-11-4-156.md.

- G5 — PARTIALLY CLOSED (verified 2026-08-26). A commit wrapper IS available: commit resolves on PATH to \$SUBMODULES_HOME/Upstreamable/commit, which chains to Software-Toolkit/Utils/Git/commit.sh (git add . + git commit) and then push_all.sh. push_all.sh reads the upstreams/*.sh recipes, and this root DOES have a tracked upstreams/GitHub.sh, so commit "<message>" works here and pushes to the github remote plus tags. What is still missing at this root is git HOOKS, not the wrapper. Note the wrapper runs git add ., which stages everything untracked — keep .gitignore accurate before using it.
- G6 — CLOSED (verified 2026-08-27). helix-deps.yaml exists at the repository root and parses under yaml.safe_load.
- G7 — CLOSED (re-measured 2026-09-01; **44 of 44**). The closure condition this entry itself named — “workshop/ is onboarded and that verifier exits 0” — is met. bash scripts/verify-governance-cascade.sh exits **0** with 12 PASS / 0 FAIL / 0 ENV / 8 NOTE: C1 classifies all **13** declared gitlinks FROM EVIDENCE (**11 owned**, 1 governance source, 1 third-party, no hardcoded roster); C2 finds all **11 × 4 = 44** owned-submodule carriers present and non-empty; C3 finds all 44 accepted by the canonical is_pointer_carrier predicate (§11.4.35 inv. 6); C6 finds helix-deps.yaml and .gitmodules in agreement both ways (13 submodules = 12 deps[] + third-party comments); C7 classifies the 1 nested gitlink under the 11 owned submodules; C8 finds every owned submodule internally in four-carrier lockstep. **Every earlier count is WITHDRAWN, not restated**, and the sequence is kept because the reason each one died is the lesson. (a) “only vasic.digital/QWEN.md exists, and it references no anchor” — false on both halves. (b) “CLOSED, 5 × 4 = 20/20” — wrong, because it enumerated a HARDCODED list of five submodules instead of deriving the fleet, the same defect class the roster itself is criticised for, committed while correcting (a).

(c) “PARTIAL, 20/24” — correct when written, superseded once workshop/ was onboarded and submodules/containers was added. A “24/24” figure circulated verbally is likewise superseded: 24 was 6 × 4, measured before submodules/containers joined the owned fleet.

(d) “**28 of 28**” (7×4) — correct on the morning of 2026-09-01, superseded the same day when submodules/LLMProvider, submodules/RAG, submodules/verdict and submodules/passage joined the owned fleet. A “36/36” figure circulated verbally is likewise superseded and was never measured: 36 is 9×4 , and the owned fleet is 11, not 9. The gate caught (b) before any human did, and the fleet moved under (d) within hours — which is the point of it. **Do not re-derive this number by hand — run the verifier and quote what it prints.**

- G8 — the markdown export mandate (11.4.65) is unmet across the repository. Unchanged and with no work in flight.
- G12 — **CLOSED 2026-09-09**, and the closure rests on the guard being observed to REFUSE rather than on a config key being present. The §11.4.109 PreToolUse guard is wired in .claude/settings.json — matcher *, command bash "\$CLAUDE_PROJECT_DIR/submodules/constitution/scripts/hooks/guard-forbidden-commands.sh" — **by reference, never copied**, because §11.4.109(A) says a copy diverges silently.

The wiring location was chosen, not defaulted to, and the reason is a lesson this file already records elsewhere. .claude/settings.json is **TRACKED** (git ls-files --error-unmatch .claude/settings.json succeeds), so a fresh clone inherits the guard. The host-wide ~/.claude*/settings.json was rejected: it would guard one developer’s machine and leave every clone unprotected with nobody told — precisely the .git/hooks/ hole recorded under “Build and test entry points” above.

docs/AGENT_GUARDRAILS.md now carries §11.4.109 **(B)** the SUBAGENT CONSTITUTIONAL PREAMBLE and **(C)** the ORCHESTRATOR PRE-ACTION CHECKLIST, with the 11.4.109 anchor literal.

Measured 2026-09-09, not asserted:

```
bash submodules/constitution/scripts/hooks/
      test_guard_forbidden_commands.sh
#   rc 0 — PASS=70 FAIL=0      (the upstream author's own
      hermetic suite)
bash scripts/verify-pretooluse-guard.sh
#   rc 0 — 8 PASS / 0 FAIL / 0 UNDET
bash scripts/verify-pretooluse-guard.sh --prove-failure
#   rc 0 — 15 cases: 12 mutations caught, 2 rc-2 states,
      control green
```

scripts/verify-pretooluse-guard.sh is the §11.4.234(A)(3) dedicated hook-validation stage, registered as pretooluse-guard in scripts/check-registry.tsv. It **executes** the wired guard against five forbidden probes (force-push, --force-with-lease, --no-verify, privilege escalation, host-power) and requires rc 2 from each; requires rc 0 from a benign call, because a guard that blocks everything is switched off within the hour; and requires that the #guardrails:allow marker does **not** downgrade host-power. Its battery proves it catches a **neutered** guard (edited to exit 0, settings untouched) and an **untracked** settings file. A validator that passes while the hook is inert is the false-green control §11.4.201 forbids.

Honest boundary (§11.4.6), three limits the closure does not hide. (a) The guard's refusal text cites the consuming-project clauses §6.T.3 / §6.U rather than the universal §11.4.113 / §11.4 anchors — upstream text inside a consumed submodule, REPORTED and not patched, per the design-toolkit precedent above. The class is blocked either way. (b) The validator executes the canonical script at its canonical path; it does not re-implement the harness's \$CLAUDE_PROJECT_DIR expansion, so it proves the script refuses, not that a future harness expands the path correctly. (c) The 70-case upstream harness is asserted PRESENT and EXECUTED as a separately named stage, printed on every green run — a recorded deferral per §11.4.234(C), never a dropped gate.

What is NOT closed by this. §11.4.109's own mechanical-enforcement list also requires the 11.4.109 literal in every per-scope governance file of every owned submodule (item 5). That was NOT measured this session and no claim is made about it.

Content boundary incident, 2026-09-01 — OPEN, and not closeable by a commit

This is not an inventory G item; it is an incident, and it is recorded here because it is the largest un-remedied fact about this repository.

Private material was written into this PUBLIC repository, committed, and pushed. Three classes crossed the boundary: verbatim prose quoted out of a private meeting-notes PDF inside the private workshop submodule; SQL and script headers copied out of a private repository; and **a third party's full real name** — a person who

is not the repository owner and who did not publish it. That name is deliberately **never written into any file in this tree**; it is referred to only as recorded in the incident note, by location.

The working tree is redacted. The commits are not. Redaction is containment, not a remedy: the content remains in the history of a **public remote**, and history is not editable after a push. Do not read the redaction, or any green gate, as closure. Removing it from history is a rewrite of public history — an operator decision, prepared and unexecuted.

Full forensics, the exact pre-redaction blob/line locations, the two-wave assessment (232 matches assessed; 29 redacted as real disclosures; 203 judged not to be) and the unexecuted remediation plan: [docs/content-boundary-incident-2026-09-01.md](#). The standing rule and its fleet map: [docs/content-boundary.md](#).

The gate that now watches for a recurrence is `scripts/verify-content-boundary.sh`, registered as `content-boundary` in `scripts/check-registry.tsv` with a `--prove-failure` paired proof and a three-valued exit. **It exits 1 today, by design** — see the “Content boundary” section near the top of this file. It was written after the leak, so it is evidence that the class is now watched, and evidence of nothing at all about what is already public.

The sweep’s own result

`scripts/verify-all-constitution-rules.sh` discovers its gates dynamically, and that population moved 57 → 286 when the constitution was fast-forwarded, so **every split published before the fast-forward is withdrawn and none of them is comparable to a present-day run**. Do not restate an old number.

A full run COMPLETED on 2026-09-03 and this is its split, observed rather than reported: exit 1 — SWEEP: FAIL – 96 FAIL + 2 ERROR out of 271 gate(s), i.e. 173 PASS / 96 FAIL / 2 ERROR of 271.

The “186 PASS / 95 FAIL / 6 ERROR of 287 gates” this block carried is WITHDRAWN, and it was never a state this file observed — it was recorded as a *reported prior measurement* from a run that had not finished. It is not comparable to the run above in any case: **the gate population itself moved 287 → 271**, so no per-count comparison

between the two is meaningful. The population has now moved twice (57 → 286 → 287 → 271) without this repository editing the sweep, because the sweep discovers its gates from the constitution submodule.

Three FAILs from the 2026-09-03 run are named here because they are actionable in THIS tree, unlike the two classes below:

cm_readme_badge_row_at_top (§11.4.259, no badge row atop README.md),
cm_zero_findings_audit_sweep (§11.4.261, no scripts/audit/
zero_findings_sweep.sh) and cm_zero_findings_monotone_ratchet
(§11.4.261(C), no docs/findings/zero_findings_ratchet.tsv). **None was acted on this session.**

Two classes of failure inside the run are known and **cannot be cleared by any commit this repository can make**:

1. **Third-party and staged carriers** — submodules/superspec/examples/..., milosvasic.ru/Upstreamable/..., and the vendored spec-kit extension copy under .specify/extensions/. scripts/verify-governance-cascade.sh reports these as *known-unclearable* and excludes them from its own verdict; the propagation-gate family inside the sweep still counts them as FAILs. Excluded there means **not double-counted, never suppressed** (§11.4.6).
2. **Defects internal to the constitution submodule**, which is upstream code this repository consumes rather than owns.

Neither class licenses treating the sweep as green.

Host capability — verify it, never assume it

Re-measured on the development host 2026-09-01, **after** an ASR stack and a generative model were installed the same day. **Two claims this section used to carry** — “No ASR engine” and “No generative model” — are **WITHDRAWN, not restated**. What follows constrains what can actually be built here:

- **A dual ASR stack IS installed — inside a project venv, not on the system interpreter.** Both engines target the same model family and neither needs a GPU:
 - **faster-whisper 1.2.1** on **CTranslate2 4.8.2**, in workshop/pipeline/venv (Python 3.14.6), with the CT2 weights at workshop/pipeline/models/ct2/faster-whisper-large-v3-turbo (1.6 GB). It installed from **prebuilt wheels**; torch is **not** present even in that venv, and is not needed.

- **whisper.cpp v1.9.1**, built **CPU-only** at workshop/pipeline/engines/whisper.cpp, binary build/bin/whisper-cli, with models/ggml/ggml-large-v3-turbo-q8_0.bin (834 MB).
- **The system python3 still has no ASR engine, and that distinction is the whole point.** faster_whisper, whisper, vosk, transformers and torch are all still absent from python3's import path; the engine resolves only through workshop/pipeline/venv/bin/python. A probe that runs bare python3 will conclude there is no transcriber, and be wrong.
- **/usr/bin/whisper exists and is a trap. This is unchanged and still load-bearing.** It is a GTK desktop notification client (© 2023 Lorenzo Paderi, GPL), not a transcriber. command -v whisper succeeds and proves nothing — and it now succeeds on a host that genuinely HAS a transcriber, which makes the trap worse, not better: the name on PATH still resolves to the wrong program. A capability probe that tests only for a name on PATH will report an engine that is not there while missing the two that are.
- **A generative model IS available locally, and it is not merely present — it is wired in.** ollama list holds **three** models, not two: qwen2.5:3b-instruct-q4_K_M (1.9 GB / 1841 MiB, **generative**, pulled 2026-09-01) alongside the two **embedding** models ordis/jina-embeddings-v2-base-code and jina-embeddings-code-cpu (323 MB each). “Nothing local can generate prose” is false as of that pull. What a 3B q4 instruct model is fit FOR is a separate question this line does not answer. **Re-measured 2026-09-02:** the running container workshop-curriculum_platform_1 has this model in its own argv — podman inspect workshop-curriculum_platform_1 --format '{{.Config.Cmd}}' shows /opt/workshop/bin/workshop-server ... -ollama http://127.0.0.1:11434 -answer-provider ollama -answer-model qwen2.5:3b-instruct-q4_K_M ... — so the workshop platform's answer path is actually wired to this model today, not merely capable of being wired to it.
- podman is present; docker is absent (command -v podman docker). No image for this repository exists.

Re-derive all of it rather than trusting the list — a host changes under you, and this section has already been wrong once in each direction:

```
ollama list
command -v podman docker
python3 -c 'import faster_whisper' # expect
ModuleNotFoundError
workshop/pipeline/venv/bin/python -c 'import faster_whisper,
    ctranslate2; \
    print(faster_whisper.__version__, ctranslate2.__version__)'
git -C workshop/pipeline/engines/whisper.cpp describe --tags
```

```
workshop/pipeline/engines/whisper.cpp/build/bin/whisper-cli --
  help | head -3
podman inspect workshop-curriculum_platform_1 --format
  '{{.Config.Cmd}}' # confirms -answer-model wiring
```

Live-production test coverage moved out of gate 6, into deploy

_tests/playwright.config.js carries testIgnore: /(restyle-seo-regression|v170-fixes|v171-hardcoding)\.spec\.js/, so **gate 6 does not run those three specs — 32 test() blocks and 97 expect() call sites**, counted 2026-09-01. They are claimed by _tests/playwright.live.config.js, whose testMatch names all four live specs.

They run at deploy time, and that IS committed. _tools/deploy-langs.sh declares LIVE_SPECS="all-languages-link-integrity restyle-seo-regression v170-fixes v171-hardcoding", pre-flights each spec file, and invokes the live config with no spec name on the command line, so the whole testMatch set executes.

The warning that stood here is WITHDRAWN as FALSE, and it actively misinformed. It read: “*via an UNCOMMITTED working-tree change made on 2026-09-01 ... At HEAD that is not yet true — git show HEAD: _tools/deploy-langs.sh names all-languages-link-integrity.spec.js alone. Until the change is committed, a fresh clone runs the other three **nowhere**.*” Re-measured **2026-09-03**: git show HEAD: _tools/deploy-langs.sh is **byte-identical** to the working tree (both sha256 5b64034fd9c368f4...), and HEAD carries all four specs on line 513 with the presence pre-flight beside it. It landed in commit **695c22d** on 2026-09-01 — so the claim was already false, or became false, the day it was written. **A fresh clone DOES run all four.** Re-derive rather than trusting either version of this paragraph:

```
git show HEAD:_tools/deploy-langs.sh | grep -n 'LIVE_SPECS='
cmp <(git show HEAD:_tools/deploy-langs.sh) _tools/deploy-
  langs.sh && echo identical
```

A figure of “86” circulates for this suite. It is a Playwright **test-case pass count** (“86 passed / 2 failed”) observed across the four-spec live run — not a count of assertions in the three deferred specs, which is 97. Both are real; they do not measure the same thing.

OPEN DEFECT, recorded and NOT fixed — gate 6 validates a FRAGMENT for milosvasic.ru, and a failing build step is silently tolerated. The word “STALE” that stood here is WITHDRAWN: the defect is COMPLETENESS, not age. Re-measured 2026-09-06 — source and _site are timestamped 89 ms apart, so nothing is stale; _site

simply holds **8 tracked files and exactly one rendered page** against ~180 front-matter source files. Served alone it answers **1 of 525** sitemap URLs; the other 524 are 404, as are all 180 article fragments and all 45 download PDFs, and it ships no assets/od/ and no assets/js/, so even its one page renders without stylesheets. Measured 2026-09-03 on this host:

```
command -v jekyll                                # (nothing –
not on PATH)
cd milosvasic.ru && bundle exec jekyll --version  # rc 127:
"command not found: jekyll"
cd milosvasic.ru && bundle check                  # missing:
jekyll 4.4.1, jekyll-seo-tag 2.8.0, jekyll-feed 0.17.0
stat -c '%y' milosvasic.ru/_site/index.html     # 2026-08-28
08:52:45
stat -c '%y' milosvasic.ru/index.html           # 2026-09-03
19:03:35
```

Jekyll cannot build on this host at all, and the gap is ONE STEP WIDER than this paragraph used to record. The claim “*ruby and bundle are present; the jekyll gem executable is not*” is WITHDRAWN: re-measured 2026-09-06, ruby and gem resolve to /usr/bin, but **bundle, bundler AND jekyll are all absent from PATH**. So the remedy named below — `bundle install` inside `milosvasic.ru` — cannot even be STARTED today; there is no bundle to run. The consequence is the part that matters: `_tests/playwright.config.js:17` sets `MV_ROOT = milosvasic.ru/_site` and line 62 serves exactly that directory to the suite — so **the milosvasic.ru half of gate 6 is asserting against a fragment, not against the site**. A green gate 6 on that half was evidence about eight files.

PRODUCTION IS NOT AFFECTED, and this was CHECKED rather than assumed. `.github/workflows/pages.yml` runs `bundle exec jekyll build --destination _site` and uploads the REGENERATED tree, so the committed `_site` is overwritten by that build and never read. The entire cost is local.

The REPORTING is fixed even though the build is not. Sixty-nine red tests for one missing build is an infrastructure fault wearing content-failure red. `_tests/preflight.js` (added 2026-09-06, §1.1 proof `_tests/prove-preflight.sh` at 8 passed / 0 failed) answers only “can this suite honestly run?” — **0 READY or 2 CANNOT DETERMINE, never 1**, because a preflight has read no site content and may make no claim about it. `scripts/pre-push-gates.sh` now calls it, so gate 6 SKIPS with a stated reason instead of emitting 69 content failures, and `PREPUSH_STRICT=1` still promotes that SKIP to a blocking failure.

Two holes in that gate's own precondition are closed with it. It tested for `milosvasic.ru/_site/index.html` — *the one file a fragment is guaranteed to have* — so the fragment passed it. And nothing asked whether a browser could START: `@playwright/test` is installed here and **webkit cannot launch at all**

(Host system is missing dependencies to run browsers), the failure being in system libraries rather than npm, so a package-presence check was checking the wrong thing.

The figure “302 passed / 0 failed / 18 skipped” for this suite is WITHDRAWN. Gate 6's own command re-run verbatim on 2026-09-06 gives **69 failed / 163 passed / 1 flaky / 6 skipped**, and **69 of the 70 listed failure lines name milosvasic.ru** — not one is a `vasic.digital` defect.

The failure is tolerated by design, and that is arguably the worse half. `_tools/deploy-langs.sh:375` runs `( cd milosvasic.ru && jekyll build ... ) || build_warn "jekyll _site rebuild",` and `build_warn` only increments a counter; `report_build_warns` then prints `N build step(s) failed` and were tolerated and the script **still exits 0**. A `--dry-run` therefore reports success while having built nothing.

PRODUCTION IS NOT AFFECTED. `milosvasic.ru` publishes through its own server-side `pages.yml` workflow (`build_type=workflow`, re-verified 2026-09-03), which nothing here touched. This defect is confined to LOCAL build and LOCAL test evidence.

Remedy — `bundle install` inside `milosvasic.ru` — is an ENVIRONMENT change and was NOT made. Neither was the tolerate-and-continue behaviour altered: silently downgrading a failed build to a warning is a separate decision, and changing it could break the deploy path for a live production site. Both are operator decisions.

The gate-6 deferral is sound and should not be reverted blindly: the measured cause was runner reachability (12 `net::ERR_TIMED_OUT`, 8 `net::ERR_NAME_NOT_RESOLVED`, 5 `EAI_AGAIN`, 77 sixty-second `page.goto` timeouts, and **zero** genuine assertion failures) while `curl` reached both live sites with `http=200`. Gating a push of undeployed source on public DNS is the defect. The question was *where* those specs run, not *whether* — and the answer is: after deployment, against what was actually shipped.

SpecKit feature 001 is planned, not built

specs/001-workshop-curriculum-platform/ carries a full planning set — spec.md (42 FR-###, 18 numbered SC-### plus SC-016a), plan.md, research.md and three research files, data-model.md, three contracts/, quickstart.md, tasks.md (91 T###) and checklists/.

This section's heading is now WRONG and is kept only until the work settles. The claim it used to carry — “analyze, implement and review have not run, and not one of the 91 tasks has been executed” — is **WITHDRAWN**. Measured 2026-09-01: specs/001-workshop-curriculum-platform/analysis.md exists (15,555 bytes), and `git -C workshop status --short` shows an untracked pipeline/ tree carrying real implementation — `run_faster_whisper.py`, `run_whispercpp.sh`, `build_whispercpp.sh`, `compare_engines.py`, `detect_media.sh`, `calibrate.sh`, `requirements.txt`, `CALIBRATION.md` — plus untracked curriculum/, platform/, evidence/ and seven docs/ pages. Building is underway.

No completion claim is made here, and none should be inferred.

Everything named above is UNTRACKED in both repositories (`git ls-files` returns nothing for `analysis.md`; workshop's own status marks the rest ??), the work was in flight while this was written, and no task-by-task status was measured. Re-derive before relying on any of it:

```
ls specs/001-workshop-curriculum-platform/  
git -C workshop status --short
```

Treat every capability `tasks.md` assumes as unverified until the “Host capability” checks above say otherwise — and note those checks moved twice on 2026-09-01, so re-run them rather than reading them.

CONTINUATION.md DOES exist at this root and is now §12.10-conformant (verified 2026-08-31, re-verified 2026-09-01). An earlier revision of this file claimed it did not; that claim was already false when written — a thin 725-byte session note was tracked at 18c84ff, but it satisfied none of §12.10's six mandatory protections, so the gap was real even though the stated reason for it was not. The document now carries the \$0 resumption prompt, the \$3 active-work detail, and the top-of-file timestamp the rule requires. `scripts/continuation-check.sh` guards it against going stale: it cross-checks the gap statuses against these carriers, verifies every entry point it names still exists and is executable, and walks `git log` to catch a commit that changed a watched governance file WITHOUT updating CONTINUATION.md (§12.10 protection 2). Three-valued: 0 in sync, 1 drift, 2 could not determine. Honest boundary (§11.4.6): §12.10 protection 1 is internally

inconsistent — it names `docs/CONTINUATION.md` while requiring the file “at the project root”. This repository resolves that in favour of the ROOT and records the ambiguity rather than silently picking a side.

Do not claim this repository passes a constitutional gate you have not actually run. Do not treat the absence of a gate as a pass.

Project overrides of universal rules

None. No clause of the universal constitution is overridden by this project. Any future override **MUST** be recorded explicitly here with its justification, per the `Override §X.Y` form in `submodules/constitution/templates/Constitution.project.md.template`.

Do not propose an Override §11.4.156. One was sought on 2026-08-27 and does not exist: 11.4.156 names and refuses the exemption vocabulary, and the inheritance contract at the head of this file (“extend them — they do NOT weaken or override any universal clause”) makes a project-local override of an inherited clause structurally impossible. The umbrella root was brought into compliance instead. `milosvasic.ru` keeps an active deploy workflow and is a **known, documented deviation** taken for production uptime; `vasic.digital` is non-compliant at the provider level with no file-level remedy. **A documented deviation is not an override** — neither may be recorded, reported, or rationalised as one. See G4 above.