

tasks

Tasks: A Repository That Clones, and a Fleet With No Unexplained Red

Input: [spec.md](#) · [checklists/requirements.md](#)

How to read this file

This file was written AFTER much of the work landed, and it says so rather than pretending otherwise. Spec 005 had no `tasks.md` at all: its 26 FR and 15 SC were unmapped, its work was done ad-hoc by parallel agents, and completion was therefore not measurable. This file reconciles the requirements against what is actually in the tree — it is not a greenfield plan and must never be read as one.

[x] means a command was run and its exit code or output observed, with the evidence on the task line. **[] carries an explicit state marker** — one of **PARTIAL**, **OPEN** or **BLOCKED** — so that an unticked box says *which kind* of not-done it is. That distinction is FR-006 applied to this document itself.

Every ticked line carries its evidence, and every measurement claim carries its population — source (files in the tree), `in_process` (a program run outside the deployed service) or `served` (measured through the running service). Where a population could not be determined from the evidence as recorded, the line says (`population: unstated`) rather than guessing. A ticked box with no evidence is the bluff this project forbids, and an evidence line with no population is the defect the constitution principle *Source Is Not Served* exists to catch.

Dates. Every re-measurement below was taken **2026-09-08** from the umbrella root at <repo-root>. Three of them contradict an earlier recorded state, and the contradiction is written on the line rather than absorbed.

Task Format

[ID] [P] [TDD] [REVIEW] [SUBAGENT] [Story] Description with file path (FR-..., SC-...)

[P] parallelizable · [TDD] RED-GREEN-REFACTOR · [REVIEW] operator gate · [SUBAGENT] delegable

Paths

Repository-relative from the umbrella root. workshop/ is the PRIVATE module; this umbrella and specs/ are PUBLIC. Counts, paths and identifiers only.

Phase 1: User Story 1 — a second engineer obtains the repository (P1)

Goal: a clone that completes on a machine that has never held this work. **Independent test:** clone into an empty directory on a fresh machine, run the documented bootstrap, read the status output.



T001 [US1] Untrack the 54 *.tar.gz.part-* archive chunks with `git rm --cached` and add the pattern to `workshop/.gitignore` (FR-002, FR-005, SC-002). **Evidence** (source): `git -C workshop ls-files '*.part*'` returns **0** tracked parts, down from 54 totalling 2,624.8 MB; `workshop/.gitignore:112` carries `*.tar.gz.part-*`, which is what stops the project's `git add .` commit wrapper putting them back; `commit c182e6a` shows all 54 deletions at 52,428,800 bytes each. Files stay on disk — nothing was deleted.



T002 [US1] Keep the three *.sha256 manifests TRACKED (≈ 9 KB) so a chunk fetched from elsewhere is still verifiable (FR-002, SC-003). **Evidence** (source): recorded in `workshop/scripts/verify-obtainability.sh` header — “without them a fetched file is

unverifiable”. The integrity record is the thing that makes moving the bytes out survivable.



T003 [US1] Record that the Git LFS remedy the spec’s own clarification chose was **prepared and WITHDRAWN before it was committed** (FR-002). **Evidence** (source): workshop/scripts/verify-obtainability.sh quotes the superseding operator instruction verbatim — remove archives from git entirely, fetch them from storage servers in a later iteration. **This is a divergence between spec.md’s Clarifications section and what was built, and it is stated here rather than left for a reader to trip over.** The spec still records LFS as the answer; the tree does not implement it.



T004 [US1] Put the documented clone command in the repository’s front-door documentation (FR-001, FR-003, SC-001). **Evidence** (source): workshop/README.md:57-60 carries `git clone --filter=blob:none ...` with the single flag justified in the same document at lines 105-160, including the REFUTED HTTP/2 hypothesis and the established no-resume mechanism.



T005 [US1] Document `--filter=blob:none` as **ROBUST, not merely fast** (FR-001, FR-003). **Evidence** (source): workshop/README.md:140 states it; the reasoning is that a 540 KB transfer that fails costs nothing to retry while a 2.6 GB one costs everything. This is the design consequence the spec’s root-cause section demanded be written down.



T006 [US1] State plainly, in the guard itself, that untracking **does not shrink the repository** (FR-005, SC-004). **Evidence** (source): workshop/scripts/verify-obtainability.sh — “`git count-objects -vH` does not fail, and a naive full-history `git clone` still transfers all of it”; workshop/README.md:159-160 repeats it. History rewriting is forbidden, so this limit is permanent and is printed rather than softened.



T007 [TDD] [US1] Ship workshop/scripts/verify-obtainability.sh — the obtainability-cost guard that warns **before** the failure point (FR-004, SC-002). **Evidence** (source): 3-valued exit contract documented in-file (0 below warning · 1 WARN/BREACH/half-finished-untracking · 2 could-not-ask); ceiling **adopted from SC-002’s 250 MB rather than selected**, `WARN_AT = CEILING - (largest single blob at HEAD)` capped at 80% of ceiling. **Both threshold rules are stated without reference to any count**, which is FR-010 applied to the guard’s own design.



T008 [US1] Make the guard measure the CHECKOUT PAYLOAD of the documented path, and report full history as a NOTE only (FR-004, SC-002). **Evidence** (source): recorded in-file — “a gate that failed on history would demand a remedy nobody is allowed to perform, and a gate whose remedy is forbidden is decoration.”



T009 [US1] Catch the half-finished untracking — a path .gitignore matches while git still tracks it (FR-002, FR-004). **Evidence** (source): that state is an explicit exit-1 finding in the guard’s contract. It is the failure mode where the rule was added and `git rm --cached` was not, so the rule does nothing and the bytes ship anyway.



T010 [TDD] [US1] **PARTIAL** — run `verify-obtainability.sh --prove-failure`, the §1.1 paired proof, and record its pass/mutation counts (FR-021, SC-014). The proof’s interface and mutation **M7** (a FINDING outranks an UNDETERMINED, never inverted) are declared in the guard’s header, but **no run of it was captured in this session**, so no count is claimed here.



T011 [US1] **OPEN** — perform SC-001’s actual trial: **3 of 3** clones by the documented command, on a machine that has never held the repository, with **0** transport failures (FR-001, SC-001). The refuting two-arm HTTPS test recorded in `spec.md` was run from a third host and proves the transport works; it is **not** the 3-of-3 acceptance run, and must not be reported as one.



T012 [US1] **BLOCKED** — **operator decision**. Publish the documented step by which recorded material is obtained from storage outside the repository, with its stated availability condition (FR-003, SC-003). The superseding instruction was explicit that the fetch instructions come in a later iteration. **Until it exists, SC-003’s “100% remains obtainable by a documented step” is NOT met** — the bytes are on one disk and the step is unwritten.



T013 [US1] **OPEN** — warn on the growth TREND, not only on the breach (Edge case: *a repository shrinks below the threshold but keeps growing*). The guard thresholds on a point-in-time size; nothing reads a series.

Checkpoint: the documented path is small and guarded; the acceptance trial (T011) and the retrieval step (T012) are the two things standing between this and SC-001/SC-003.

Phase 2: User Story 2 — a reader can tell a finding from a designed state (P1)

Goal: every red states whether it is a defect, a declared condition, or a check that could not run. **Independent test:** run every instrument; for each red, answer three questions from its output alone.



T014 [US2] Give workshop/scripts/verify-obtainability.sh a three-valued exit where **2 is never a pass** (FR-006, FR-008, SC-006). **Evidence** (source): contract stated in-file; precedence asserted by mutation **M7** so a moving tree cannot mask a breach already measured.



T015 [TDD] [US2] Make _tools/watch-deploy.sh read its exit codes and fold each cycle into a three-valued verdict (FR-006, FR-008, SC-006). **Evidence** (source, from scripts/check-registry.tsv and commit f824cd1): verdict is 0 clean / 2 a cycle could not run / 5 a cycle ran and failed, **failure outranking could-not-determine**. Before this, the watcher discarded \$? and could loop 240 times against a deploy failing every cycle and still exit 0.



T016 [TDD] [US2] Ship _tools/prove-watch-deploy-verdict.sh, the \$1.1 paired proof, driven by **data** rather than by editing its subject (FR-021, SC-014). **Evidence** (in_process, recorded in scripts/check-registry.tsv): **9 passed / 0 failed / 8 mutations** against the current watcher, and **3 passed / 6 failed** against a reconstruction of the prefix watcher — every one of the 6 reporting got 0 where a verdict was owed. Mutations are stub exit codes via WATCH_DEPLOY_CMD; the watcher's source sha256 is captured before the battery and re-asserted after, so a mutation that edited the subject would fail the proof.



T017 [US2] Rewrite the _tools/watch-deploy.sh registry exemption to carry its evidence and WITHDRAW the expired reason (FR-007). **Evidence** (source): scripts/check-registry.tsv — the prior reason “*it judges nothing*” is marked WITHDRAWN by name, with the date and the change that expired it. **An exemption whose justification expired still looks reviewed**; this is the fix for that class.



T018 [US2] Write read-only provider adapters for the two hosts that had none — gitflic.ru and gitverse.ru (FR-011c, SC-005). **Evidence** (source): scripts/verify-provider-ci.sh grew 585 lines in

commit f824cd1; both adapters speak HTTP directly because neither host ships a CLI, and each declares its own honest ceiling in-file (gitflic **cannot** reach a NONE verdict by construction and says so on every row; gitverse can, its route set established empirically).



T019 [TDD] [US2] Prove the two adapters with the M8..M13 battery, including the case that stops an adapter hardwired to refuse from passing (FR-011c, FR-021). **Evidence** (source): scripts/verify-provider-ci.sh:408-540 — M8 unreachable $\Rightarrow$ rc 2 with the row UNVERIFIED; **M9 a control that goes green** $\Rightarrow$ rc 0 with the row NONE. Without M9 the battery would be satisfied by an adapter that can only ever say no.



T020 [US2] **BLOCKED — operator credential**. Turn the 6 UNVERIFIED provider rows into verdicts (FR-011c, SC-005). **Re-measured 2026-09-08** (served, provider APIs): bash scripts/verify-provider-ci.sh exits **1** — 22 repositories / 40 upstream rows, **1 CONFIRMED**, and the 6 rows now report “*an adapter IS registered for this host but could not run*” rather than “*no adapter registered*”. <https://api.gitflic.ru> answers **HTTP 403** and <https://api.gitverse.ru> answers **HTTP 401**; both require a credential on every route and none is configured. **This is FR-011c working exactly as written** — “never asked” became “asked and could not determine”, which is progress, and it was NOT converted into a pass. Lifting it needs GITFLIC_TOKEN / GITVERSE_TOKEN, which only the operator can issue.



T021 [US2] Judge the content-boundary class A INWARD rows rather than allow-listing them (FR-007, FR-011, SC-007). **Evidence** (source): docs/findings/content-boundary-class-a-judgement.jsonl holds **1,933 rows, every one verdict: NOT_A_DISCLOSURE, 0 unjudged** — verified by counting the file’s own verdict field, not by transcription. docs/content-boundary-class-a-judgement.md groups them G1–G8 with a per-group reason. **Nothing was redacted, allow-listed, re-baselined, and scripts/verify-content-boundary.sh was not edited. The gate still exits 1, and it should.**



T022 [US2] Judge the **union** of two gate runs’ INWARD rows, not the more convenient one (FR-023, SC-007). **Evidence** (source): both runs are tabled in docs/content-boundary-class-a-judgement.md — totals 16,155 and 9,945, tied together with --expect-corpus; **51 rows appeared only in run 2** and are included. Judging the union is what stops a moving tree quietly dropping a row out of the reading queue.



T023 [US2] Attribute the total's fall rather than banking it (FR-011, FR-023). **Evidence** (source): the same document tests four candidate mechanisms for the 16,155 → 9,945 drop in 39 minutes and states “*nothing was silenced to achieve it*”. **An unexplained fall is worse than a stable red.**



T024 [US2] Report the direction split across five populations whose totals differ by more than 3,000 rows (FR-011, FR-018). **Evidence** (source): OUTWARD landed at 79.7 / 79.3 / 79.1 / 80.2 / 82.1 % every time. **The totals move; the direction does not** — that is the figure to act on, and it is what bounded the reading assignment.



T025 [TDD] [US2] Register scripts/verify-content-boundary-judgement.sh as a check with a --prove-failure paired proof (FR-020, FR-021, SC-014). **Evidence** (source): row added in scripts/check-registry.tsv in commit f824cd1; the gate is 450 lines.



T026 [US2] Install the pre-push gate hook, since .git/hooks/ is untracked and a fresh clone runs no gates (FR-020). **Evidence** (source): .git/hooks/pre-push is present, executable, 1,224 bytes, dated 2026-09-08 10:58. **Honest limit:** git push --no-verify still bypasses it with no record, and this hook protects only this checkout.



T027 [US2] Take the R3 hardcoded-path red to green by a real fix, not a re-baseline (FR-011, FR-011a, SC-007). **Re-measured 2026-09-08** (source): bash scripts/audit-hardcoded-paths.sh exits 0 — “*no machine-specific hardcoded paths*”, 16 file(s) explicitly allowed. The R3 baseline of 17 live occurrences is cleared.



T028 [US2] **PARTIAL** — work through the 387 baselined path occurrences, partitioned by owning repository (FR-011a, FR-011b). **Evidence so far** (source): commit f824cd1's own message records the partition and that **51% can only be fixed by an upstream commit returning as a gitlink bump**. The partition exists; the campaign does not. **FR-011b's re-derivation requirement is unmet:** no row has been re-derived to prune debt already fixed upstream, so the 387 is still an unpruned figure that overstates.



T029 [US2] Close the R2 continuation red and the R4 governance-pin red (FR-011). **Evidence** (source): commit f824cd1 updated CONTINUATION.md (+207 lines) and moved submodules/constitution three commits, the first move in six that was **not** corpus-neutral — the

constitution gained anchor 11.4.272 and every recorded figure in the four carriers was stale and was updated in lockstep. **BOTH REDS HAVE SINCE REOPENED, re-measured 2026-09-08 from the umbrella root** (source): `bash scripts/continuation-check.sh` exits **1** at *7 PASS · 1 DRIFT · 0 UNDET · 15 NOTE* — “CONTINUATION.md IS STALE”; `bash scripts/verify-submodule-remote-sync.sh` exits **1** at *12 CURRENT / 1 DRIFT / 0 UNDETERMINED*, the single DRIFT row being `submodules/constitution` (gitlink 71ac4373... vs remote 2db64ee 8..., **difference DETERMINED, direction NOT**, because the remote commit is absent from this object store). **Closing an instance is not closing the class, and this pin has now gone stale six times. Do not read T029’s tick as a green fleet.**



T030 [US2] Confirm the manifest and gitlinks moved together (FR-011). **Re-measured 2026-09-08** (source): `bash scripts/verify-manifest-pins.sh` exits **0** — **13 MATCH / 0 DRIFT / 0 UNDETERMINED** of 13 declared deps. C9 caught the manifest lagging its gitlinks **twice on 2026-09-08** (commits bfe2931, 3922e35) and both were repaired; that is the gate working.



T031 [US2] **OPEN** — build the fleet summary that reports **defect count and declared-condition count separately and never merges them** (FR-009, SC-008). **No such instrument exists.** Every gate reports its own verdict; nothing aggregates them into the two-number statement FR-009 requires, so SC-008’s “unexplained reds reaches 0, from a baseline of 11” is **not computable today**.



T032 [US2] **OPEN** — give every non-clean instrument an explicit KIND on its own output line: `defect / declared / could-not-run` (FR-006, SC-005). Several instruments do this well (T014, T015, T020); **it is not universal**, and SC-005’s “100% of non-clean results state their kind” has not been measured across the fleet.



T033 [US2] **OPEN** — make a declared condition name **who may lift it and what evidence would** (FR-007, SC-015). `scripts/check-registry.tsv` exemptions now carry evidence and reasons (T017), and T020’s row names the credential, but no instrument requires a lifting authority as a structural field.



T034 [TDD] [US2] **OPEN** — assert FR-010 mechanically: any threshold, bucket, allow-list or baseline change must carry a principle stated independently of the count it produces (FR-010, SC-007). Two guards do this in prose by construction — `verify-obtainability.sh`’s two threshold rules (T007) and the palette gate’s

fixed floor — but **nothing detects a weakening change**. This is the requirement most exposed to being satisfied by assertion.



T035 [REVIEW] [US2] **OPEN** — run SC-015’s acceptance: a reader answers “*is this a defect, what is the evidence, what happens next*” for every red from instrument output alone, without consulting a person. Not attempted; it depends on T031–T033.

Phase 3: User Story 3 — every learner-facing defect is repaired or disclosed (P2)

Goal: nothing silently broken; every omission stated. **Independent test:** walk one area end to end and confirm nothing is silently absent.



T036 [TDD] [US3] Ship the B6/B7 detector for questions citing withheld material in `workshop/pipeline/extract/verify_question_banks.py` (FR-012, SC-009). **Evidence** (source): 478 lines added in commit `c182e6a`; **B6** asserts every citation resolves in the registry AND is redacted: `false`, **B7** the same for every `lesson_sections` pid. The file records the mechanism in its own header: a question citing a redacted passage is withheld under `WithholdCitationRedacted`, **so a learner opening those areas was served a shorter test with nothing saying so.**



T037 [US3] **PARTIAL** — repair or record the removal of every question citing withheld material (FR-012, SC-009). The **detector** landed (T036); **whether the affected questions were repaired was NOT confirmed in this session**, because confirming it means running the bank verifier against the private corpus. **SC-009’s “0 questions cite material the server withholds, down from 5” is therefore UNVERIFIED here** — the instrument that would answer it exists and was not run. (population: `unstated`) — the R8 baseline of 5 was a source count over the authored banks; the served count was never separately established.



T038 [US3] **OPEN** — tell a learner, in the interface, that an area has no assessment and why (FR-013, SC-010). No task in the tree implements the learner-facing disclosure; the platform’s absence-honesty gates (`workshop/platform/gates/verify-absence-honesty.sh`) govern what the *server* asserts, not what the *learner reads*.



T039 [US3] Remove R7's structural obstacle — area documents never sectioned into the passage registry, so no citation could resolve and no question could be authored (FR-015, SC-011). **Evidence** (source): `git -C workshop ls-files 'docs/training/curriculum-areas/*.md'` returns **37** and `...*.sections.json` returns **37 — 37 of 37 sectioned**, against R7's baseline of 25 of 37 unsectioned. The obstacle is gone at source.



T040 [US3] **PARTIAL** — report the resulting coverage figure, and report **separately and honestly** the areas that carry no assessment because the corpus cannot support one (FR-015, SC-011). The structural blocker is cleared (T039); **the two populations have not been counted apart**, and SC-011 requires exactly that separation. Merging them would let a corpus limitation be reported as progress.



T041 [REVIEW] [US3] **BLOCKED — operator gate**. Read the 3 documents failing publication review on **62, 54 and 58** uncited claim blocks and decide publish-or-not (FR-014, SC-010). They remain unpublished, so a client sees fewer areas than are authored. This is the same gate carried as T060 in `../006-session-record-and-qa-readiness/tasks.md` — **one decision, recorded in two specs, not two decisions**.



T042 [US3] **OPEN** — assert FR-014 mechanically: a published document's primary references resolve to material a reader can actually reach. T041 is a human read of three documents; **no check enumerates published documents and resolves their primary references**.

Phase 4: User Story 4 — the interface looks designed on the surfaces people use (P2)

Goal: materially more hue variety on principal surfaces; contrast floors hold in both presentations. **Independent test:** sample what is painted on principal surfaces in both presentations; count hue variety and measure contrast within each pairing.



T043 [TDD] [US4] Make `workshop/platform/gates/verify-served-palette.sh` read the **SERVED bundle only**, through `lib/`

served_css.sh, and never src/styles/** (FR-016, FR-018, SC-012). **Evidence** (source): stated as the gate's entire reason for existing — measured defect **D6** was that the token source's hue span was widened, nothing regenerated the served stylesheets, and the token measurement was then reported as evidence about the interface. “**A check that read the token source to reach its verdict would have returned SUCCESS over exactly that defect.**”



T044 [US4] Family the colour an element **paints of its own** (paint), not only the ground behind it (bg) — defect **D-A** (FR-016, FR-018). **Evidence** (source): recorded in the same gate. For a non-text row the sampler sets bg to the ground *behind* the indicator, which is correct for contrast but meant **a coloured chip, wash, rule or fill could not contribute a hue family AT ALL, by construction**. The population excluded the very elements an interface uses to carry colour. A row with no paint still grades exactly as before, so an older sample is unaffected.



T045 [US4] Define a hue family as a 30° bucket (12 buckets), with imperceptible-chroma colours collapsing into ONE neutral family counted once (FR-016, FR-018, SC-012). **Evidence** (source): stated in-gate — “*counting each grey as its own hue would let a monochrome interface pass*”. Two colour spaces are used on purpose: HSL for the hue ANGLE (it need only partition colours that already have a hue), and a different measure for whether a colour has a hue at all, because HSL saturation divides by a span that vanishes near white and black. **The floor of 6 is fixed by the specification, not chosen by the gate** — FR-010 applied to the instrument.



T046 [US4] **PARTIAL** — re-measure the served hue-family count and record it against SC-012's floor of 6, up from a baseline of 2 (FR-016, SC-012). The corrected instrument exists (T043–T045) and workshop/platform/gates/sample-served-surfaces.mjs takes the browser sample. **No served run was captured in this session, so no post-fix count is claimed here.** CONTINUATION.md:646 still records the baseline “*the served interface paints 2 hue families against a floor of 6*”, and until a served run supersedes it, **that is the last measured figure in this tree**. Note the shape of the trap this gate was built against: when the instrument became honest the count went **down** from 4 to 2 — a corrected gate reporting a worse number than the broken one it replaced.



T047 [US4] **OPEN** — give every hue present a stated classification job, so no hue is added only to raise a count (FR-017, SC-012). The

design-system documents landed in commit f11a3c1 (workshop/docs/design-system/TOKENS.md, RATIONALE.md, COMPONENTS.md, wk-semantic-tokens.css), but **nothing asserts the mapping from every painted hue to a declared job**, which is what SC-012's second half requires.



T048 [TDD] [US4] Ship workshop/platform/gates/verify-served-contrast.sh with its paired prover, measuring text and non-text pairings separately in both presentations (FR-019, SC-013). **Evidence** (source): both verify-served-contrast.sh and prove-served-contrast.sh are tracked; workshop/docs/design-system/contrast.py (245 lines, commit f11a3c1) is the measurement, and its own commit records that **nine of the specification's own colours were short** — the spec was wrong and the instrument found it.



T049 [US4] **PARTIAL** — record the served contrast result: **0** failures across text and non-text pairings in both presentations (FR-019, SC-013). The gate and its prover exist; **no served run was captured in this session**.



T050 [US4] **OPEN** — justify the measured population independently of its effect on the count, and record that justification whenever the population changes (FR-018, SC-012). T044 **changed the population** (adding paint) for a stated defect reason, which is the right shape — but no rule requires the next change to do the same, and the edge case *“a red is fixed by making the instrument weaker”* lives exactly here.

Phase 5: Evidence discipline — cross-cutting



T051 [TDD] Drive every palette mutation by **data**, not by altering the check (FR-021, SC-014). **Evidence** (source): 8 fixtures tracked under workshop/platform/gates/fixtures/palette/ — empty, flat, good, low-contrast, near-grey-ramp, scheme-split, indicator-only, boundary-ulp. indicator-only is the fixture that pins defect D-A; boundary-ulp pins the bucket edge. **A gate green by construction is the inoperative-proof defect this project has already caught once**.



T052 Make the content-boundary gate carry evidence its corpus did not move during the measurement (FR-023). **Evidence** (source): both runs in docs/content-boundary-class-a-judgement.md report co

rpus MOVED with the changed-file count (66 of 14,891; 4 of 14,837) and are tied together by --expect-corpus. **The gate reports its own instability instead of absorbing it silently.**



T053 [TDD] **PARTIAL** — make every check whose subject is a set establish the set is **non-empty** before reporting the set contains no defects (FR-022, SC-014). Two instruments do this: verify-obtainability.sh returns **2** when nothing is tracked to measure, and workshop/platform/gates/verify-bank-reachability.sh was written precisely because the prior gate checked one fixture area while 89 questions were unreachable. **It is not universal, and no sweep asserts it across the fleet.**



T054 **OPEN** — satisfy SC-014: **100%** of requirements map to at least one automated check, and every such check ships a data-driven paired demonstration (FR-020, SC-014). **This file is the first half of that criterion and does not complete it.** Every FR and SC in spec.md is now cited by at least one task above — but a task is not a check, and 22 of the 54 tasks here are not ticked. **Requirements with no automated check at all: FR-009, FR-010, FR-013, FR-014, FR-017 and SC-015.** Those six are the real coverage gap.

Status

Counted from this file by script, not transcribed.

Phase	tasks	DONE	PARTIAL	OPEN	BLOCKED
1 US1 obtainability	13	9	1	2	1
2 US2 red-state discipline	22	15	1	5	1
3 US3 learner-facing	7	2	2	2	1
4 US4 presentation	8	4	2	2	0
5 Evidence	4	2	1	1	0
Total	54	32	7	12	3

```
grep -c '^-\ [' specs/005-clone-and-clear-red/tasks.md # 54
grep -c '^-\ [x\]' specs/005-clone-and-clear-red/tasks.md # 32
grep -c 'PARTIAL' specs/005-clone-and-clear-red/tasks.md # 7
task lines
grep -c 'BLOCKED' specs/005-clone-and-clear-red/tasks.md # 3
task lines
```

32 of 54 carry a captured run or an observed artefact. Of the 22 not ticked, **3 are BLOCKED on an operator** (T012 the storage-fetch instructions, T020 the two provider credentials, T041 the three unpublished documents), **7 are PARTIAL** — the instrument exists and the measurement or the campaign does not — and **12 are OPEN**, with no implementation found in the tree.

Read the PARTIAL column before the DONE column. Five of the seven are the same shape: *a corrected instrument exists and was never run against the served product*. T037, T046 and T049 each name a number this specification claims to have moved, and **none of the three post-fix figures was captured in this session.**

The three reds that reopened

T029's tick is the most misreadable line in this file, so it is restated here. R2 and R4 were both closed in commit f824cd1, and both were measured RED again on 2026-09-08 from the umbrella root:

```
bash scripts/continuation-check.sh          # rc 1 - 7 PASS · 1
DRIFT · 15 NOTE
bash scripts/verify-submodule-remote-sync.sh # rc 1 - 12
CURRENT / 1 DRIFT
bash scripts/verify-manifest-pins.sh         # rc 0 - 13 MATCH /
0 DRIFT
bash scripts/audit-hardcoded-paths.sh        # rc 0 - 16 file(s)
allowed
bash scripts/verify-provider-ci.sh           # rc 1 - 1
CONFIRMED, 6 could-not-run
```

The governance pin has now gone stale **six** times. **Treat it as a standing operator decision that recurs, not a task that completes** — and never quote a green from this document as current. Re-run the five commands.

Dependencies

T001 untrack → T007 guard → T010 proof
└──────────→ T011 the 3-of-3 trial ← the
SC-001 acceptance
T012 storage-fetch instructions ← BLOCKED, and SC-003 is unmet
without it

T014/T015/T018 three-valued instruments → T031 the fleet summary
→ T035 SC-015
└────────→ T032 kind

on every red

└─> T033

lifting authority

T039 sectioning → T040 the two coverage populations, counted apart

T043/T044/T045 corrected palette gate → T046 the served re-measurement

T031 blocks SC-008. Nothing today can state how many reds are defects and how many are declared, so the headline criterion of this specification — unexplained reds reaching 0 from a baseline of 11 — **is not computable**, however many individual gates are green.

Requirement coverage

Every FR and SC in spec.md is cited by at least one task above: **26 of 26 FR** and **15 of 15 SC**, verified by script rather than transcribed.

```
python3 - <<'PY'
import re
s=open('specs/005-clone-and-clear-red/spec.md').read()
t=open('specs/005-clone-and-clear-red/tasks.md').read()
ids=[]
for m in re.finditer(r'\*(FR-[0-9]+[a-z]?|SC-[0-9]+[a-z]?)\*.*',
                    s):
    if m.group(1) not in ids: ids.append(m.group(1))
un=[i for i in ids if not re.search(r'(?![0-9a-z])'+i+r'(?![0-9a-z])', t)]
print(len(ids), 'ids ·', len(un), 'uncited'); print(' '.join(un))
PY
```

A citation is not a check. The six requirements with no automated check anywhere in the tree are named in T054, and that list — not the coverage percentage — is the finding this file exists to surface.