

plan-coverage-fix-c

Plan-coverage proposal rule — the veto guard and Fix C

Feature: 006-session-record-and-qa-readiness · FR-007, FR-007a **Date:** 2026-09-09 **Status:** BOTH CHANGES APPLIED on operator decision. Not committed by the authoring session. **Subject:** pipeline/extract/propose_plan_coverage.py, platform/gates/verify-plan-coverage-proposal.sh, platform/gates/prove-plan-coverage-proposal.sh (private workshop submodule)

CONTENT BOUNDARY — BINDING. This file lives in a PUBLIC repository and its subject is inside the PRIVATE workshop submodule. Every planned point, passage and extracted row below is named **by identifier or path only**. No point text, no transcript text, no curriculum prose, no area title and no person’s name appears here, and none may be added. Matched vocabulary is reported as counts, never as words — a bag of discriminative terms is a paraphrase of the sentence it came from.

1. Why the rule was investigated at all

Three proposal rows that a human had **independently recorded covered**, each with an evidence locator, came back from the rule as cannot_determine (medium confidence):

plan_id	chapter
01M1H9DJ052XDSQQJVZNSN8KCB	02
01M1H9DJ052XDSQQJVZNSN8KCC	02
01M1SD6HJ78ZTKAXSW0185YKH8	02.01

The operator was offered “rewrite them”, “correct them to covered” and “withdraw them”, and **chose none of those: investigate the rule that produced them first**. What follows is that investigation and the two changes it led to. **No row was repaired by hand at any point**.

2. What the rule does

Two lanes per planned point, each scored as IDF-weighted recall of the point’s discriminative terms and each compared against a **matched null** of 300 synthetic points built from same-rarity terms drawn from the origin chapter’s vocabulary:

- **transcript lane** — a sliding 5-segment window over the chapter’s own transcript;
- **own-row lane** — the best match against a **single** extracted row (meeting note / TODO / open question) whose `origin_chapter` is this chapter.

An absence guard already existed for the negative branch: a lane may support `not_covered` only when `q_upper95 <= alpha`, where `q` is the probability that a same-rarity *unrelated* point leaves **no** trace in this chapter. That guard had replaced a `transcript_segments >= 500` constant under decision #14 of 2026-09-08, precisely because the constant made `not_covered` structurally unreachable for a short chapter.

3. Root cause — TWO causes, not one

The three rows do not share a cause, and forcing one story would have been wrong. All three were `ONE_LANE_ONLY`: the transcript lane beat its null, the own-row lane did not.

Cause 1 — 01M1SD6HJ78ZTKAXSW0185YKH8 (chapter 02.01). The own-row lane is *measurably blind* here (`q_upper95 = 0.8618`). The absence guard already computed that number, already printed it, and already refused to let this lane create a negative — and then let the same lane **veto a positive** with no guard at all. **Decision #14’s guard had been applied to one branch only.**

Cause 2 — ...KCB and ...KCC (chapter 02). Blindness does **not** explain these. Their own-row lane is measurably sighted (`q_upper95 = 0.0099`) and it genuinely dissented (`p = 0.1367` and `0.1867`, no ties, 299 distinct

null values). The rule’s local reasoning was valid. The defect is one level up: a lane with **no measurable relationship to the outcome** was a mandatory conjunct.

The measurement that carried both decisions. Against the 7 points a human has recorded covered on this tree:

lane	fires	captures human- recorded covered	expected by chance	one-sided hypergeometric
transcript	7/49	4/7	1.00	p = 0.0049
own-row	5/49	1/7	0.71	p = 0.55

Supporting structural measurements. The median own row carries 5 content tokens after stoplisting (chapter 02) against a point of 6–8 discriminative terms; chapter 02.01’s entire own-row evidence base is **3 rows for 30 planned points**, so 27 of those 30 have zero overlap with any own row and score exactly 0.0 with $p = 1.0$ by arithmetic. This is the single-unit design the transcript lane’s 5-segment window exists to avoid — *“a single segment cannot hold a whole point and a per-segment match would be unfair by construction”* — and that argument was never applied to the own-row lane. **Widening it does not help and was tried:** a union over *all* own rows fires on 1/49 instead of 5/49.

Excluded by measurement, not by assumption. Threshold and seed instability are not the cause: across four seeds only 2–3 of 49 rows change any lane’s significance and the covered count stayed 1 every time. A mid-p correction for ties changes nothing.

Withdrawn during the investigation. An intermediate finding that covered was “physically unreachable for 44 of 49 rows” **is withdrawn** — it restated the observed lane failure rather than proving unreachability. Checked properly, a single own row carrying 1–2 of the point’s rarest terms would suffice everywhere ($k_{\text{needed}} = 1$ or 2 against a max own-row length of 10–30). There is no unreachability theorem; there is an extreme power deficit.

4. The two changes

4.1 The veto guard (applied first)

A lane whose $q_{upper95} > \alpha$ is **disqualified as a witness** on both branches: it may neither create a negative nor veto a positive. Same α , same already-computed number, **no new parameter and no new draw from the RNG**.

Blast radius, predicted before and confirmed after: exactly 1 row of 49. 01M1SD6HJ78ZTKAXSW0185YKH8, cannot_determine → covered; covered 1 → 2. It fixed cause 1 and, correctly, did **not** fix cause 2.

4.2 Fix C (applied second, separate operator decision)

covered now requires the **transcript lane and nothing else**. The own-row lane is reported on every row and decides nothing — it may neither create a verdict nor veto one. This mirrors what the rule has always done for not_covered: *a coverage claim is a claim about the SESSION, and the transcript is the session*.

Three positive codes, differing **only** in what the demoted lane did, so that demoting it hides nothing:

code	meaning	rows
TWO_LANES_BEAT_MATCHED_NULL	both lanes agreed	1
ONE_LANE_BEAT_MATCHED_NULL_THE_OTHER_IS_MEASURABLY_BLIND	the own-row lane could not see	1
TRANSCRIPT_LANE_DECISIVE_OWN_ROW_LANE_SIGHTED_AND_DISSENTING	the own-row lane could see, dissented , and was overridden	5

ONE_LANE_ONLY is **renamed** to OWN_ROW_LANE_ONLY_THE_DECISIVE_LANE_DID_NOT_AGREE (4 rows) — the old name asserted a symmetry between the lanes that no longer exists.

Removed by Fix C: the veto guard’s *symmetric* arm — a significant own-row lane while the transcript lane was blind. Under Fix C the own-row lane never creates a verdict. That arm **never fired** on this tree, so the removal is measurable at zero rows; it is recorded here rather than left to be discovered.

5. Blast radius of Fix C — predicted 6 rows, observed 6 rows

covered 1 → 7 against the original artefact (2 → 7 against the veto-guard artefact). **No recorded human decision is contradicted, and none could be: all 7 recorded dispositions are covered; no human has recorded a negative anywhere in this tree.**

chapter	plan_id	→ code	human decision?
02	01M1H9DJ052XDSQQJVZNSN8KCB	sighted-dissent overridden	covered, with evidence locator — refuted row 1
02	01M1H9DJ052XDSQQJVZNSN8KCC	sighted-dissent overridden	covered, with evidence locator — refuted row 2
02.01	01M1SD6HJ78ZTKAXSW0185YKH8	overridden blind lane	covered, with evidence locator — refuted row 3
02	01M1H9DJ052XDSQQJVZNSN8KC6	sighted-dissent overridden	none — surfaced for reading
02	01M1H9DJ052XDSQQJVZNSN8KC8	sighted-dissent overridden	none — surfaced for reading
02	01M1H9DJ052XDSQQJVZNSN8KCT		

chapter	plan_id	→ code	human decision?
		sighted-dissent overridden	none — surfaced for reading

6. HONEST BOUNDARY — what Fix C costs and what it does not establish

These four points are carried in the generator’s own docstring and in the gate’s header, not only here.

1. **It discards a measured dissent by design.** On ...KC6 and ...KC8 the own-row lane is measurably able to see and dissented strongly ($p = 0.8433$ and 0.5733). Those are disagreements, not silences, and this rule overrides them. The third positive code exists so every such row says so about itself; `evidence.own_row_lane.p_value` is the size of what was discarded. **The pre-Fix-C behaviour is recoverable exactly by ignoring every row carrying that code.**
2. **Specificity is UNMEASURED and the evidence is one-sided.** $p = 0.0049$ and $p = 0.55$ are agreement on **positives only**. **No human has recorded a not_covered or a dropped anywhere in this tree**, so there is no negative ground truth and nothing establishes either lane’s specificity. A lane that fired on every row would score as well. The ground truth is 7 rows. Read the change as “*the own-row lane adds nothing measurable to the positives*”, never as “*the transcript lane is right*”.
3. **Three of the six moved rows are flagged near_decision_threshold** — ...KC6, ...KCB, ...KCT. At least one of their lanes has a p-value inside $0.02..0.12$, so a different Monte-Carlo seed could move them back. A verdict that moves with a seed is not a verdict.
4. **The three undecided rows are NOT measured false positives.** They are rows no human has read, which this change now puts in front of one. Nothing here says they are covered.

And the standing limit both changes inherit: a planned point is the output of a cue-phrase heuristic over ASR text and is *what was said* about the next session, never a ratified agenda. Every verdict inherits that and cannot be stronger than it. A proposal remains a **draft for a human**; it is not the decision record, and the served figure still moves only when a human decides.

7. Evidence

- `verify-plan-coverage-proposal.sh` — **rc 0, zero findings** (P1–P11).
- `prove-plan-coverage-proposal.sh` — **71 passed / 0 failed / 61 mutations / 10 controls**, rc 0.
- **RNG invariant.** The null draws are one shared sequential stream, so an added or removed draw would shift every downstream row. Measured after both changes: **0 of 49 rows show any change to evidence** — every p-value, null p95, `zero_trace_trials` and `q_upper95` byte-identical to the original artefact. Both changes touch the DECISION only.

P11, the new assertion, is the mirror of P7 and enforces two things: the **decisive lane must have carried the positive** (the demoted lane may not decide), and the **code must match what the own-row lane actually did**, so a row that overrode a real sighted dissent can never be filed as corroboration or as blindness. P11 checks that a positive is well-formed and honestly labelled. **It does not check, and cannot check, that a positive is true.**

Controls, one per admissible positive shape (C11, C11a, C11b), keep covered from being made unreachable by tightening the gate. M29 — a covered created by the own-row lane alone — and M35 — an overridden sighted dissent mislabelled as an overridden blind lane — are the two arms that hold Fix C to its stated shape.

P11 was rewritten the day it was written. Its first form forbade exactly what Fix C then adopted, and its paired mutation M29 existed to prevent that adoption. Both were replaced rather than deleted when the operator adopted the larger rule on measured grounds. **A proof left pointing at a withdrawn rule is a gate green over something nobody believes.**

8. Population reconciliation, 66 → 49 — traced per row

Recorded here because an exact total can hide two offsetting errors, and here it did.

Every one of the 17 removed rows is accounted for. 0 unreconciled.

Traced two ways: from each ledger entry's own `affected_plan_points` field, and by **counterfactual re-derivation** (a throwaway copy of curriculum/ with all dispositions removed, re-run through `cmd/ session-record`).

The 6 ledger retractions remove 11 ROWS, not 6 — five origin-01 points cost two rows each (once in chapter 02, once carried into 02.01); one origin-02 point costs one:

retracted point	minted in	rows removed
01M1H9DJ052XDSQQJVZNSN8KC4	01	02/...8KC4, 02.01/...8KC4
01M1H9DJ052XDSQQJVZNSN8KCG	01	02/...8KCG, 02.01/...8KCG
01M1H9DJ052XDSQQJVZNSN8KCK	01	02/...8KCK, 02.01/...8KCK
01M1H9DJ052XDSQQJVZNSN8KCQ	01	02/...8KCQ, 02.01/...8KCQ
01M1H9DJ052XDSQQJVZNSN8KCF	01	02/...8KCF, 02.01/...8KCF
01M1SD6HJ78ZTKAXSW0185YKHB	02	02.01/...YKHB

The other 6 rows are decision #12's human covered decisions — six chapter 02 points resolved and therefore stopped carrying into 02.01. Measured, not inferred: emptying the dispositions restores exactly these six rows and no others — 02.01/...8KCB, ...8KCC, ...8KCH, ...8KCN, ...8KCP, ...8KCV. (The 7th decision, ...YKH8 in 02.01, removes no row: that chapter has no successor.)

So $66 - 11 - 6 = 49$ is right and the attribution was inverted. The reconciliation as previously written credited the 11 to a separate population of “extraction defects” and the 6 to the ledger. In fact **the ledger accounts for the 11 and the coverage decisions account for the 6**. There is no separate population: decision #13's “11” and the ledger's “6” are **the same retirement counted two ways — 11 rows, 6 distinct points**.

Two supporting measurements: the 6 retracted ids are **absent from every chapter's next-meeting-points.jsonl** — the extractor fix, not the ledger, is the mechanism, and the ledger is documentation of it — and the counterfactual derivation with an emptied ledger returns 55, not 60, which is what proves that.

A label that survives because two wrong numbers cancel is a false-green in prose. That is why the per-row trace was required, and it is the transferable lesson of this section.