

quickstart

Quickstart — validating this feature end to end

Feature: Session record + QA readiness · **Date:** 2026-09-08 **Build measured:** index generation 7, 39 served areas, served CSS sha256 c6d03995...

Every command below was executed. Where a result is not yet obtainable, that is stated rather than written as an expectation.

Exit contract: 0 clean · 1 a real finding · 2 could-not-determine. A 2 is never a pass.

0 — Confirm you are measuring the served product, not the source

Do this first, every time. The most-repeated defect in this project is a source measurement reported as a served one.

```
curl -s 127.0.0.1:8087/api/areas | python3 -c 'import
    json,sys;d=json.load(sys.stdin);print(d["generation"],
    len(d["areas"]))'
podman inspect workshop-curriculum_platform_1 --format
    '{{.State.StartedAt}}'
ls -l workshop/platform/bin/workshop-server
```

If the binary is newer than the container start time, every figure below is about a stale process. That state persisted four hours undetected today.

1 — The session record (US1, US2)

```
curl -s 127.0.0.1:8087/api/chapters/02 | python3 -c '
import json,sys; sr=json.load(sys.stdin)["session_record"]
for n in sr["section_order"]:
    s=sr["sections"][n]
    print(f"{n:16} {s[\"own_state\"]:9} own={s[\"own_count\"]:>3}
        carried={s[\"carried_count\"]:>3}")'
```

Expected — measured:

meeting_notes	authored	own= 50	carried= 0
open_questions	authored	own= 45	carried= 11
todo	authored	own= 35	carried= 27
next_meeting	authored	own= 18	carried= 24

Check the split, not just the totals. own_count=0 with an authored reason is a real state; without the split it would read as a productive session.

```
bash workshop/platform/gates/verify-session-record.sh # rc 0
bash workshop/platform/gates/prove-session-record.sh # rc 0,
14 passed
```

Vacuity: the prover's empty-tree case must exit 2. A record check over no chapters has verified nothing.

2 — Carry-forward, and that nothing vanished

```
curl -s 127.0.0.1:8087/api/chapters/02 | grep -o 'CARRIED-
FORWARD' | wc -l # 126
curl -s 127.0.0.1:8087/api/chapters/02.01 | python3 -c '
import json,sys; cf=json.load(sys.stdin)["session_record"]
["carry_forward"]
print("outbox:",cf["is_outbox"],"points:",cf["count"],"max_slip:",cf["max_slip_cou
```

Expected: outbox: True points: 42. **The latest chapter has no successor, so its points must still land somewhere findable** (FR-011).

planned=66 covered=0 carried=66 dropped=0 vanished=0. **covered=0 is honest, not a stub** — no coverage decision has been recorded by anybody, and coverage_state: undecided says so.

3 — The learner journey (US4)

```
AID=$(curl -s 127.0.0.1:8087/api/areas | python3 -c 'import
    json,sys;print(json.load(sys.stdin)["areas"][0]["id"])')
S='X-Session: quickstart'
curl -s -H "$S" 127.0.0.1:8087/api/areas/$AID/lessons | python3
    -c '
import json,sys; ls=json.load(sys.stdin)["lessons"]
print("lessons:",len(ls),"with a body:",sum(1 for l in ls if
    (l.get("body") or "").strip()))'
```

The header is mandatory — without it these routes answer 400, and the browser client omitting it made the entire surface unreachable.

Then: fetch the assessment while locked and grep the body for prompt, choices, correct, explanation — **all must be absent**. Complete the required lessons, re-fetch, confirm questions appear **without** answer fields, submit, and confirm the result is stable on re-read.

4 — The gates, against the running build

```
for g in lesson-content bank-reachability session-record sc009-
    citation-span \
    answer-key-disclosure submission-validation server-unity
    crossref-currency; do
    printf '%-28s ' "$g"; bash workshop/platform/gates/verify-$g.sh
    >/dev/null 2>&1; echo "rc=$?"
done
```

Measured after rebuild and restart: **all rc 0**, with server-unity at PASS=43 FAIL=0 UNDET=0 DEBT=8.

`verify-lesson-content.sh --live` was rc 1 before the restart — 318 of 318 EMPTY — and that red arm is the point. A gate with only the arm that passes proves nothing.

5 — Cross-references, and a long job that is not a defect

```
bash workshop/platform/gates/verify-crossref-currency.sh
```

Measured **rc 1** immediately after restart (generation 7 has NO crossref run), then **rc 0** 45 minutes later: 498580 edge(s) over 24929 source(s).

The red was correct while derivation ran. Do not treat it as a defect without checking the container log for deriving cross-references.

6 — The QA and client documents (US3)

Regenerating. The versions on disk were written at 12:30, before the rebuild, and describe 14 areas with tests and 0 lessons with content against a live 27 and 318 of 318. **A stale expected-result column manufactures false defects.** Both must carry the build id they were written against.

7 — Evidence discipline (US5)

Pick any claim in contracts/evidence.md and re-run its producer. It must reproduce, and it must state its **population** — source, in_process or served.

```
bash workshop/platform/gates/prove-answer-key-disclosure.sh
# rc 0, 16 mutations
bash workshop/platform/gates/prove-submission-validation.sh
# rc 0, 11 mutations
```

Check each battery has a **control that goes green** and a **vacuity refusal exiting 2**. A battery of only refuse-to-pass cases is satisfied by a gate hardwired to fail.

What this quickstart can and cannot prove today

Scenario	Status
0 Population check	PASS
1 Session record served	PASS — 12/12 sections, 0 empty-without-reason
2 Carry-forward, vanished=0	PASS
3 Learner journey, real content	PASS — 318/318 bodies, key withheld while locked
4 Gate sweep	PASS — 8 of 8 rc 0
5 Cross-references	PASS — after a 45-minute derivation
6 QA + client documents	

Scenario	Status
	IN FLIGHT — regenerating against this build
7 Evidence discipline	PASS
Three standalone section routes	IN FLIGHT
/meeting-notes route	BLOCKED — withholding rule is an operator decision
Coverage verdicts	UNDECIDED by design — a proposal may not move the figure