

tasks

Tasks: A Complete Session Record per Chapter, and a Platform a QA Team Can Test

Input: [spec.md](#) · [plan.md](#) · [research.md](#) · [data-model.md](#) · [contracts/](#) · [quickstart.md](#)

How to read this file

[x] means a command was run and its exit code observed, with the evidence on the task line. Tasks are kept rather than deleted because eleven parallel streams built this during the session that specified it, and a list hiding completed work would misrepresent both what remains and how the remainder must be verified.

A checked box is not a claim that the capability is live. Where a task was verified in-process only, the line says so — that distinction is the subject of constitution principle *Source Is Not Served*, which this feature's own spec violated before it was written.

Task Format

[ID] [P] [TDD] [REVIEW] [SUBAGENT] [Story] Description with file path

[P] parallelizable · [TDD] RED-GREEN-REFACTOR · [REVIEW] human gate before consumers · [SUBAGENT] delegable

Paths

Repository-relative from the umbrella root. `workshop/` is the PRIVATE module; `submodules/curriculum-kit/` is the public reusable library.

Phase 1: Setup — the shape the content is authored against



T001 [REVIEW] Publish the reusable content library as `vasic-digital/curriculum-kit` (PUBLIC), mount as a gitlink, declare in `helix-deps.yaml` with a real ref (FR-030, SC-013). **Evidence** (source): 82 files; content-boundary scan 0 prose / 0 name rows; verify-`standalone-clone.sh` 3 gaps → 2.



T002 Add `Lesson.Body` to `submodules/curriculum-kit/pkg/curriculum/model.go` as opaque marked-up text the library neither parses nor renders (FR-013). **Evidence** (in_process): `go build/go test/go vet rc 0`; commit `07af6fb`.



T003 [P] Keep `Body` **optional** in the library and record why in `TestValidateDoesNotRequireALessonBody` (FR-013). **Evidence**: a materials-only lesson is legitimate, so “every lesson has a body” is a consumer policy the library may not impose; a future rule must change a test that explains its absence.



T004 [P] Assert `Body` and `Summary` have not collapsed into one value in `TestLessonBodyIsCarriedAndIsDistinctFromSummary`, including that an absent body decodes as `""` rather than a copy of the summary (FR-013).

Execution note: `DecodeDocument` calls `DisallowUnknownFields()`, so a body key the library does not declare fails the **entire directory load**. Prove one file decodes before writing 42.

Phase 2: Foundational — one derivation per fact

No user-story work may begin until T005 and T008 are complete.
Every story's correctness derives from them.



T005 [TDD] [REVIEW] Compute the publication decision once in workshop/platform/backend/internal/api/area_publication.go and consume it from all three surfaces (FR-023, FR-025). **Evidence** (population: unstated): verify-area-publication-consistency.sh rc 0 (6 passed); prover rc 0 (7 mutations); TestAreaListAndDetailAgree asserts the **pair**, which is what was missing while both handlers' own tests passed.



T006 Make catalog staleness four-valued in workshop/platform/backend/pkg/learning/catalog.go — Fresh/Stale/StalenessUnknown/NoSource (FR-026a). **Evidence**: NoSource is separate so an in-memory test catalog, having no directory to have drifted from, keeps its legitimate negatives.



T007 Guard every determined negative about the catalog in workshop/platform/backend/internal/api/lessons.go via catalogNegativeRefused() (FR-024, FR-026a). **Evidence** (population: unstated): stale ⇒ 503 with catalog_snapshot: "stale" and the changed filenames; polarity proven — guard off returns the false 404, guard on returns 503.



T008 Build the session-record derivation in workshop/platform/backend/pkg/sessionrecord/sessionrecord.go — one place the rules live (FR-007, FR-012, SC-003). **Evidence** (source): Resolve has **no fourth branch**, so vanished=0 holds by construction; a drop without a reason is refused and carries instead.



T009 [P] Order chapters by the existing curriculum.ChapterID.OrdinalPath rather than a new scheme (FR-002, FR-008). **Evidence**: a second ordering is a second thing to disagree.

Checkpoint: areas[] is exactly the set whose detail answers 200 (817 unopenable → 0), and no planned point can vanish.

Phase 3: User Story 1 — a participant reads what a session produced (P1)

Goal: four sections per chapter, cumulative, each with content specific to that session. **Independent test:** open each chapter, read all four sections, confirm none is empty or generic.



T010 [US1] Reuse the existing producer workshop/pipeline/extract/meeting_notes.py rather than forking it (FR-001, SC-001). **Evidence** (source): it already emitted all four sections into curriculum/chapter-<slug>/knowledge/*.jsonl; the brief's "0 matches in chapters/" was accurate and misleading — the material lives under curriculum/.



T011 [US1] Serve session_record on GET /api/chapters/{id} in workshop/platform/backend/internal/api/chapters.go, degrading on that key alone and never 503-ing the chapter (FR-005).



T012 [US1] Split own_count from carried_count per section (FR-004, FR-009). **Evidence** (population: unstated): two sections carry own_count=0 with an authored reason; without the split they would read as productive sessions.



T013 [US1] Require a reason when own_state is empty (FR-004). **Evidence** (population: unstated): 12 of 12 sections present, **0 empty-without-reason** — this is the property that lets a tester tell a defect from a blank.



T014 [P] [US1] Carry origin_chapter on every item (FR-003, SC-002).



T015 [US1] Build GET /api/chapters/{id}/open-questions, preserving the three-valued resolution the client already types (answered_in_chapter/answered_by_research/unanswered) (FR-005). **The wire must not narrow it** — the consumer already knows what it needs.



T016 [P] [US1] Build GET /api/chapters/{id}/todo reading pkg/sessionrecord, never re-deriving (FR-005).



T017 [P] [US1] Build GET /api/chapters/{id}/next-meeting reading pkg/sessionrecord (FR-005).



T018 [TDD] [US1] Assert the routes and the embedded key serve the **same items** in `workshop/platform/gates/verify-session-record.sh` (FR-005, FR-023). **This is the pair assertion** — a route and a key that agree today are exactly what drift, and that class produced the 819-vs-817 catalogue bug.



T019 [REVIEW] [US1] **BLOCKED — operator decision.** `GET /api/chapters/{id}/meeting-notes` is NOT built. Its withholding rule **may not be guessed** (FR-005a): one note is a content judgement left undecided. Update its DEBT row in `workshop/platform/gates/route-manifest.tsv` to name that specific blocker rather than the general one (FR-028).

Phase 4: User Story 2 — a planned point that slipped is visibly carried forward (P1)

Goal: nothing planned disappears; slippage is visible. **Independent test:** take a point planned and not covered; confirm it appears in the next chapter, marked, with its origin.



T020 [US2] Derive
`plan(N) = unresolved(plan(N-1))` with `slip+1 + next-meeting` points of `N-1` in `pkg/sessionrecord` (FR-006, FR-008, FR-010, FR-012). **Evidence** (population: unstated): 66 planned · 0 covered · 66 carried · 0 dropped · **0 vanished**.



T021 [US2] Resolve absence of a coverage decision to carried, never to a silent drop (FR-007, SC-003). **Evidence:** Resolve has no fourth branch — the invariant holds by construction, not by diligence.



T022 [P] [US2] Emit marker: `CARRIED-FORWARD` and `marker_label` with the count (FR-009, FR-010, SC-004). **Evidence** (population: unstated): 126 markers, 281 slip references on chapter 02.



T023 [US2] Serve the latest chapter's carry-forward as an **outbox** with `is_outbox: true` and a stated reason (FR-011). **Evidence** (population: unstated): 42 points, `max_slip=2`.



T024 [US2] Record basis: undetermined for chapter 01 and **compute no coverage percentage for it** (FR-006, FR-024). **Evidence:** no

predecessor and no recorded intent; a percentage over an invented plan is a fiction with a denominator.



T025 [TDD] [US2] Prove the derivation with 16 corpus mutations in workshop/platform/gates/prove-session-record.sh (FR-025, SC-011).

Evidence (population: unstated): rc 0, 14 assertions — control green, vacuity refusal exits 2, duplicate planned point rc 1, reasonless drop refused.



T026 [SUBAGENT] [US2] Propose a coverage verdict per point with **per-point evidence and confidence**, marked PROPOSED, written to a separate artefact — never into plan-coverage.json, which is the decision record (FR-007a).



T027 [TDD] [US2] Assert coverage_state is still undecided and the served counts have not moved after a proposal exists (FR-007a). **This is what makes “a proposal is not a decision” mechanical rather than aspirational.**



T028 [REVIEW] [US2] **Operator gate**: confirm or overturn the proposed verdicts (FR-007a). Until then the served figure stays undecided.

Phase 5: User Story 3 — a QA team can test the platform manually (P1)

Goal: a tester unfamiliar with the platform completes the document without asking a question. **Independent test**: hand it to someone who has never seen this and watch.



T029 [US3] Write workshop/docs/qa/MANUAL-TEST-PLAN.md with every scenario's expected result stated **before** the action (FR-018, FR-019, FR-020, FR-022). **Evidence** (source): 82 scenarios across 11 surfaces; \$0 is an “already known” table read first; \$3 specifies seven mandatory defect-report fields.



T030 [US3] Write workshop/docs/qa/CLIENT-WALKTHROUGH.md, every step performed (population: unstated) (SC-015).



T031 [SUBAGENT] [US3] **Regenerate both against the current build** (FR-020a). They describe 14 areas with tests and 0 lessons

with content; live is 27 and 318 of 318. A stale expected-result column manufactures false defects and spends a tester's day disproving our own fixes.



T032 [US3] Record the build identifier in both — served-CSS sha256, served area count, index generation — and state they must be regenerated when it moves (FR-020a, FR-026).



T033 [US3] Disclose up front (FR-020, FR-020b): 39 served against 42 authored; 12 areas carrying no test with the reason; every score indeterminate while banks mix machine-marked and free-text questions.



T034 [P] [US3] Disclose that /meeting-notes is deliberately unbuilt and that three sibling routes may or may not be live — **check, do not assume** (FR-020, FR-028, SC-008).

Phase 6: User Story 4 — every area detail shows real teaching and a real test (P2)

Goal: lessons with content; a test where one exists; a stated reason where none does. **Independent test:** open every published area.



T035 [US4] Emit "body" verbatim from workshop/pipeline/extract/build_learning_catalog.py (FR-013, SC-005). **Evidence** (source): the builder already parsed every body, used it only for estimatedMinutes, and discarded it — 339 of 339 now carry one, median 328 words, longest 4,563.



T036 [US4] Emit body unconditionally from workshop/platform/backend/pkg/learning/wire.go (FR-013). **Evidence:** "" is a determined "unauthored"; an absent key is indistinguishable from a transport fault.



T037 [US4] Render the body through the app's existing markdown renderer in workshop/platform/frontend/src/app/features/areas/area-lessons.component.ts — **not** a second renderer and **not** [innerHTMLHTML] (FR-013).



T038 [TDD] [US4] Ship verify-lesson-content.sh with **two arms** (FR-013, FR-023, FR-026). **Evidence:** in-process 339/339 rc 0; **live rc**

1 at 318 of 318 EMPTY before the restart — the red baseline on the broken artefact. A gate with only the passing arm proves nothing.



T039 [US4] Send X-Session from the learning client via the **one** exported header builder in workshop/platform/frontend/src/app/core/api.ts (FR-013, FR-014). **Evidence** (served): lessons and assessment 400 → 200; the entire surface was unreachable in a browser.



T040 [US4] Fix the assessment envelope read one level too high, so a healthy 200 no longer renders as malformed_response (FR-014).



T041 [US4] Stop treating an empty questions array behind a **shut gate** as “no test exists” — consult question_count and the required-lesson gate first (FR-014, FR-015).



T042 [TDD] [US4] Ship verify-bank-reachability.sh — the **populati on** check that was missing (FR-014, FR-023, FR-026). **Evidence** (population: unstated): the existing gate checked route mountedness against one fixture area, so 89 questions were unreachable while it passed.



T043 [US4] Make the no-test copy state only what the server established (FR-015, FR-026a, SC-006). **Evidence** (population: unstated): it asserted “*there is genuinely none*” for 13 areas that had a bank — **confident prose asserting the opposite of the truth**.



T044 [P] [SUBAGENT] [US4] Author text alternatives for the 205 video materials (FR-016, FR-028). **Evidence needed**: 205 of 213 carry alt: ""; the interface names the absence rather than hiding it, so this is a disclosed gap, not a hidden one.

Phase 7: User Story 5 — a reviewer can confirm no claim is unearned (P2)

Goal: every claim resolves to machine-produced evidence or is marked unverified. **Independent test**: sample claims; each resolves or is explicitly unverified.



T045 [US5] Withhold answer keys from the open practice route in workshop/platform/backend/internal/api/questions_graded.go

(FR-?). **Evidence** (population: unstated): the decks are the same rows — 224 of 224 match by prompt, 0 of 224 by id — so gating practice would delete a feature; the defect was the disclosure.



T046 [US5] Fail closed on a nil, unreadable **or stale** catalog (FR-024, FR-026a). **Evidence**: a graded prompt the snapshot never saw would look “graded nowhere” to a literal index and be published.



T047 [US5] Refuse a malformed submission in workshop/platform/backend/internal/api/assessment_submit_body.go, naming the offending key, before the store is touched (FR-014). **Evidence** (population: unstated): a misspelled key returned 200 and permanently recorded a 0% attempt; responses: [] still grades.



T048 [US5] Assert word-span containment on **both** endpoints in workshop/platform/backend/pkg/knowledge/wordjoin.go (FR-?). **Evidence** (population: unstated): containment was checked on word *starts* while the published span was derived from word *ends*; 7 of 2158 joined segments could publish a span their own passage does not contain.



T049 [P] [US5] Fix the gate-server leak in workshop/platform/gates/lib/learning_stack.sh (FR-?). **Evidence** (population: unstated): `BASE=$(ls_start ...)` runs in a **subshell**, so `LS_PID=$!` never reached the caller and every trap read an empty variable and returned 0 — 21 orphans reaped from this host.



T050 [TDD] [SUBAGENT] [US5] Add the build identifier to /api/health (FR-026). **Evidence needed**: nothing noticed the container was four hours older than the code, because no reading could be tied to a revision.



T051 [SUBAGENT] [US5] Commission a **second** independent adversarial review against the rebuilt container (FR-029, SC-014). The first found four false headline claims; a review that confirms is worth little.

Phase 8: Polish & cross-cutting



T052 Untrack the 54 video-archive parts with `git rm --cached` and `gitignore` the pattern (FR-028). **Evidence** (source): tracked 54 → 0,

on-disk bytes unchanged at 2,752,266,473, fresh checkout 2,711.7 MB → 87.0 MB. **size-pack did not fall and could not** — history rewriting is forbidden; the gate prints that limit itself.



T053 [P] Return **rc 2** with `TOD0(storage-fetch)` when archive parts are absent, and keep **rc 1** for a present-but-corrupt part (FR-024, SC-010). **Evidence**: a missing download is not a corrupt archive, and conflating them sent a previous investigation at the wrong thing.



T054 [P] Make `_tools/watch-deploy.sh` read its exit codes (FR-023, FR-025). **Evidence** (`in_process`): it could loop 240 times against a deploy failing every cycle and exit 0; the proof fails **6 of 8** against a reconstruction of the pre-fix watcher.



T055 [P] Fix the hardcoded checkout path in `workshop/docs/training/curriculum-areas/16-*.md` and regenerate its four derived artefacts (FR-?). **Evidence** (`source`): the document teaches pointer discipline and broke it; audit 22 occurrences → `rc 0`; 597 citations still resolve; 18 of 18 sidecar pids present in the registry.



T056 Restart once when every writer has stopped, verified by two `git status` readings 30 s apart (FR-027). **A restart triggers a 45-minute cross-reference derivation** over 24,929 passages, during which `verify-crossref-currency.sh` correctly reports `rc 1`.



T057 Re-run the full gate sweep against the rebuilt container and record each figure with its **population** — `source`, `in_process` or `served` (FR-023, FR-026).



T058 [P] [SUBAGENT] Register every new gate in `workshop/platform/gates/check-registry-006.tsv` (FR-?). **Evidence needed**: an unregistered gate fails the registry's own R5 anti-drift rule, which caught one of mine today.



T059 [P] [SUBAGENT] Re-validate every exempt row whose subject changed this session (FR-028). **Evidence needed**: one said “*it judges nothing*” of a script that had gained a three-valued verdict — an exemption whose justification expired still looks reviewed.



T060 [REVIEW] **Operator gate**: read the 3 documents failing publication review on 62, 54 and 58 uncited claim blocks (FR-020b, FR-023). They stay unpublished until read; a client sees 39 areas, not 42.

Dependencies

Phase 1 (kit + Body)

└> Phase 2 (T005 publication · T008 session derivation) ← the checkpoint

└> Phase 3 US1	└	independent of one another
└> Phase 4 US2	└	
└> Phase 5 US3	└	
└> Phase 6 US4	└	
└> Phase 7 US5	└	

Phase 8 – T056 requires EVERY writer stopped; T057 requires T056

T005 and T008 block every story. Nothing else does — the five stories touch disjoint file sets and were built concurrently.

T019, T028 and T060 are operator gates, not work items. T031 blocks the QA session — handing over stale documents is worse than handing over none.

Parallel Execution

Stream	Owns	Tasks
Content	docs/training/, curriculum/, pipeline/	T010, T020–T026, T035, T055
Backend	platform/backend/, platform/gates/	T005–T009, T011–T018, T036, T042, T045–T050
Frontend	platform/frontend/ src/app/	T037, T039–T041, T043, T044
QA docs	docs/qa/	T029–T034
Umbrella	_tools/, scripts/	T052–T054

Cross-stream rule: any stream measuring another's output must **rebuild immediately before measuring and state which build it read**. The partition held today; the timing did not — a stylesheet was read mid-rewrite and a build id was hashed over an orphaned file a cp -a left behind.

Implementation Strategy

MVP = Phases 1, 2 and 3. A chapter that carries its four sections honestly.

Increment 2 = Phase 4 — carry-forward, which is what makes the record more than a snapshot: nothing planned disappears.

Increment 3 = Phases 5–7. QA readiness, learner completeness, evidence.

Phase 8 is not optional polish. Until T056 and T057 run, every fix in Phases 6–7 is verified **in process** and unverified **as served** — the exact distinction constitution principle *Source Is Not Served* exists to enforce, and the one this feature's own specification violated before it was written.

Status

Counted from this file by script, not transcribed.

Phase	tasks	done	open
1 Setup	4	4	0
2 Foundational	5	5	0
3 US1 session record	10	5	5
4 US2 carry-forward	9	6	3
5 US3 QA readiness	6	2	4
6 US4 area completeness	10	9	1
7 US5 evidence	7	5	2
8 Polish	9	4	5
Total	60	40	20

```
grep -c '^- \[' specs/006-session-record-and-qa-readiness/
tasks.md      # 60
grep -c '^- \[x\]' specs/006-session-record-and-qa-readiness/
tasks.md      # 40
```

40 of 60 carry a captured run. Of the 20 open, **3 are operator gates** (T019, T028, T060) and **2 require the single restart** (T056, T057). The remaining 15 are engineering work, and T031 is the one that blocks the QA session.

Superspec bridge — where the bite-sized steps live

/speckit-superspec-tasks ran superpowers:writing-plans over the remainder. Output:

workshop/docs/superpowers/plans/2026-09-08-qa-handover-remainder.md (313 lines)

Scope check: of the 20 open tasks, **9 are in flight** with live agents (T015–T018 routes, T026 coverage, T031–T034 QA documents) and **3 are operator gates** (T019, T028, T060). Writing bite-sized steps for work in flight would duplicate an agent mid-task and create two authorities for one file. The blueprint therefore covers the **8 unheld** tasks.

Blueprint	tasks.md	Produces
A1	T044	verify-material-alt.sh + prover; 205 video alts authored
A2	T050	build identity on /api/health, dirty-tree marked
A3	T059	expired exemptions withdrawn, not overwritten
A4	T058	every new gate registered; R5 green
B1	T056	the single restart, gated on two stable readings
B2	T057	full re-measurement, every figure carrying its population
C1	T051	the second adversarial review, against the rebuilt container

Two properties worth naming.

A1 writes the gate BEFORE authoring the content, and requires it to fail first. If it passes on today's tree its population is wrong, because the 205 missing alternatives are real. Its prover includes an overreach case — *an alt identical to its caption fails* — which a gate written only against emptiness would have certified as correct.

The blueprint's own self-review found a sequencing gap and states it rather than hiding it: A2's build stamp is invisible until B1 restarts, so A2's *live* verification belongs to Track B. That is written into A2's step 4.

Execution order:

A1 A2 A3 A4 (independent, now) A2's live check defers to B2
[every writer stops] -> B1 -> (45 min derivation) -> B2 -> C1